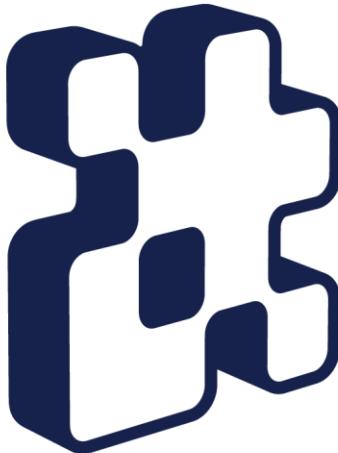


# Introduction to GraphBLAS

A linear algebraic approach for concise, portable,  
and high-performance graph algorithms

Gábor Szárnyas (CWI)

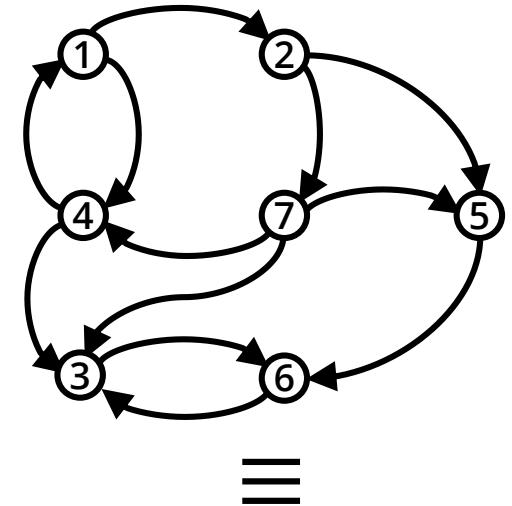


**GRAPHBLAS**

# Motivation

# GRAPHBLAS

- Graph algorithms are challenging to program
  - irregular access patterns → poor locality
  - caching and parallelization are difficult
- Optimizations often limit portability
- The **GraphBLAS** introduces an abstraction layer using the language of linear algebra
  - graph  $\equiv$  sparse matrix
  - traversal step  $\equiv$  vector-matrix multiplication

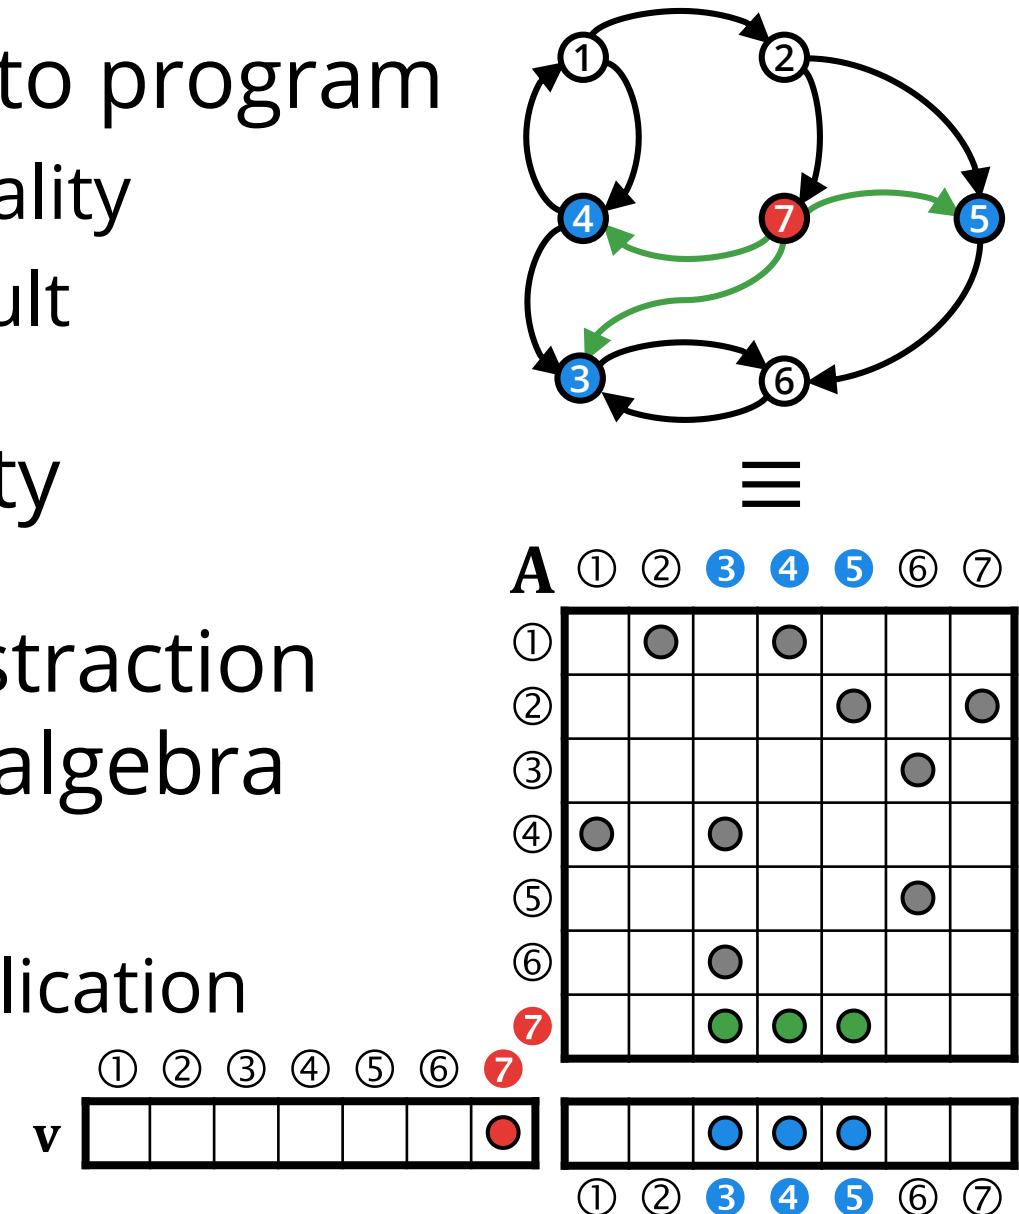


$\equiv$

A	①	②	③	④	⑤	⑥	⑦
①		●		●			
②					●		●
③							●
④	●			●			
⑤							●
⑥			●				
⑦		●	●	●	●		

# GRAPHBLAS

- Graph algorithms are challenging to program
  - irregular access patterns → poor locality
  - caching and parallelization are difficult
- Optimizations often limit portability
- The **GraphBLAS** introduces an abstraction layer using the language of linear algebra
  - graph  $\equiv$  sparse matrix
  - traversal step  $\equiv$  vector-matrix multiplication



# What makes graph computations difficult?

# GRAPH PROCESSING CHALLENGES

connectedness

the “curse of connectedness”

computer  
architectures

contemporary computer architectures are good at processing linear and hierarchical data structures, such as *lists*, *stacks*, or *trees*

caching and  
parallelization

a massive amount of random data access is required, CPU has frequent cache misses, and implementing parallelism is difficult

B. Shao, Y. Li, H. Wang, H. Xia (Microsoft Research),  
*Trinity Graph Engine and its Applications*,  
IEEE Data Engineering Bulletin 2017



# GRAPH PROCESSING CHALLENGES

Graph algorithm have a *high communication-to-computation ratio*.

Speedup with a CPU that has better arithmetic performance:

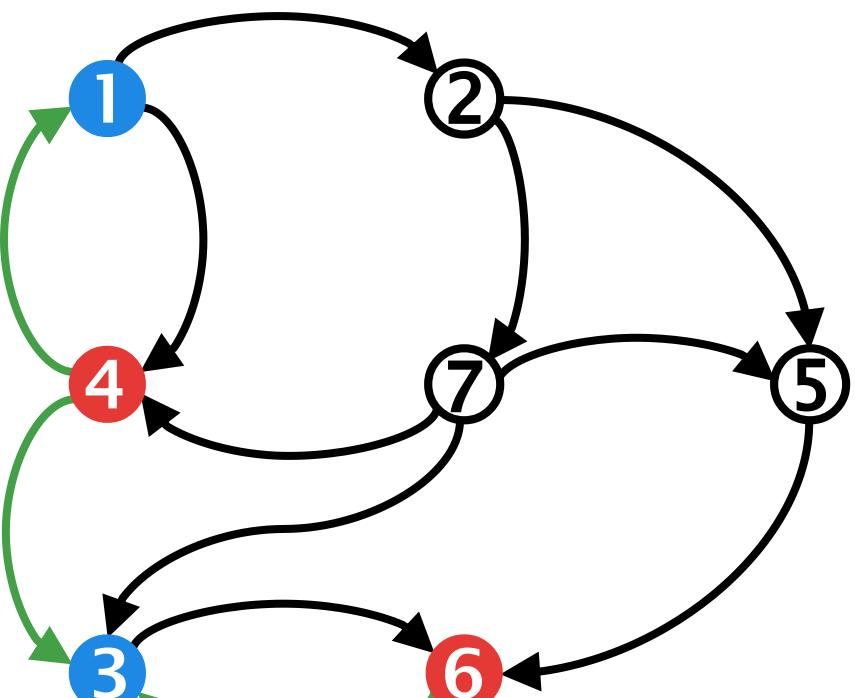
- machine learning → **a lot**
- relational queries → **some**
- graph processing → **very little**

Standard latency hiding techniques break down, e.g. pre-fetching and branch prediction provide little benefit.

Use a data representation and computation model which are expressive, machine-friendly, and portable.

# LINEAR ALGEBRA-BASED GRAPH PROCESSING

- Graphs are encoded as sparse adjacency matrices.
- Use vector/matrix operations to express graph algorithms.



A	1	2	3	4	5	6	7
1		1	1				
2				1	1		1
3							1
4	1		1				
5							1
6			1				
7				1	1	1	

f	1	2	3	4	5	6	7
frontier				1		1	

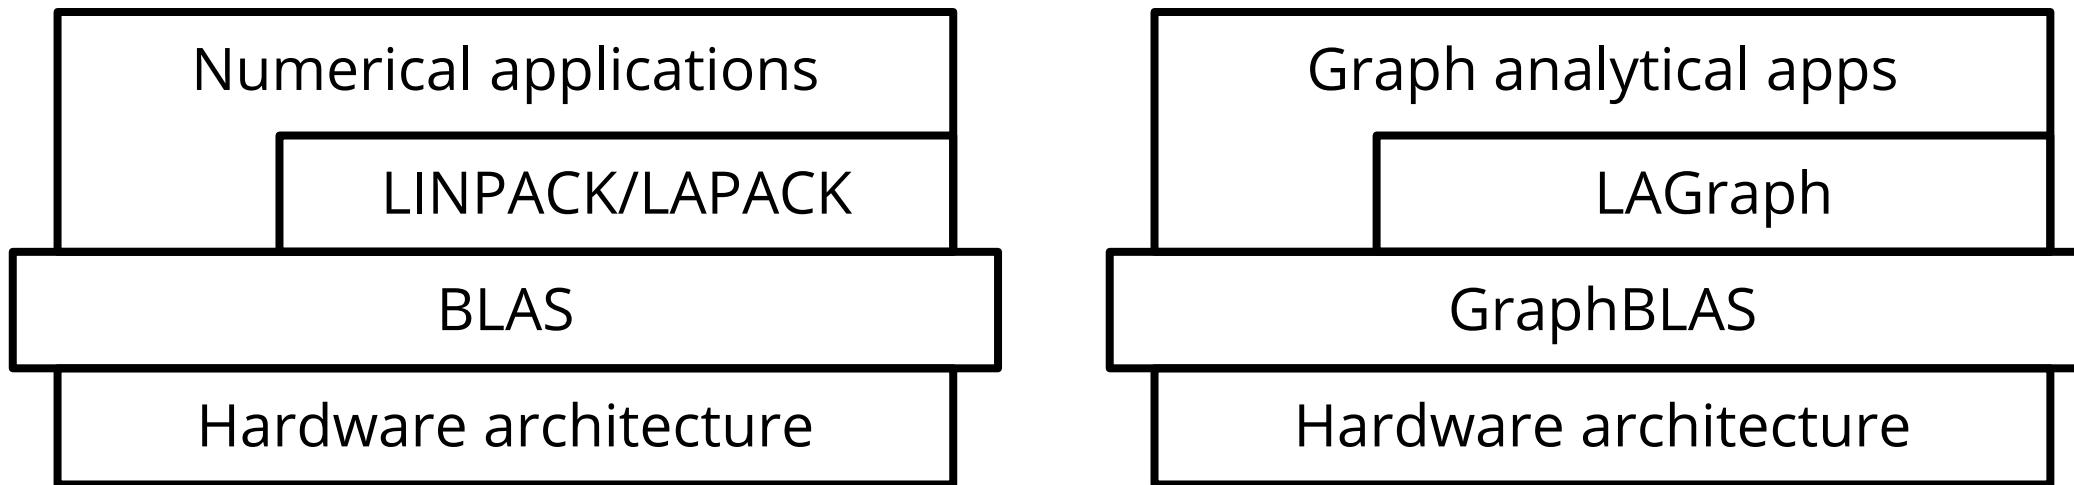
fA	1	2	3	4	5	6	7
	1	2					

# The GraphBLAS standard

# THE GRAPHBLAS STANDARD

Goal: separate the concerns of the hardware, library, and application designers.

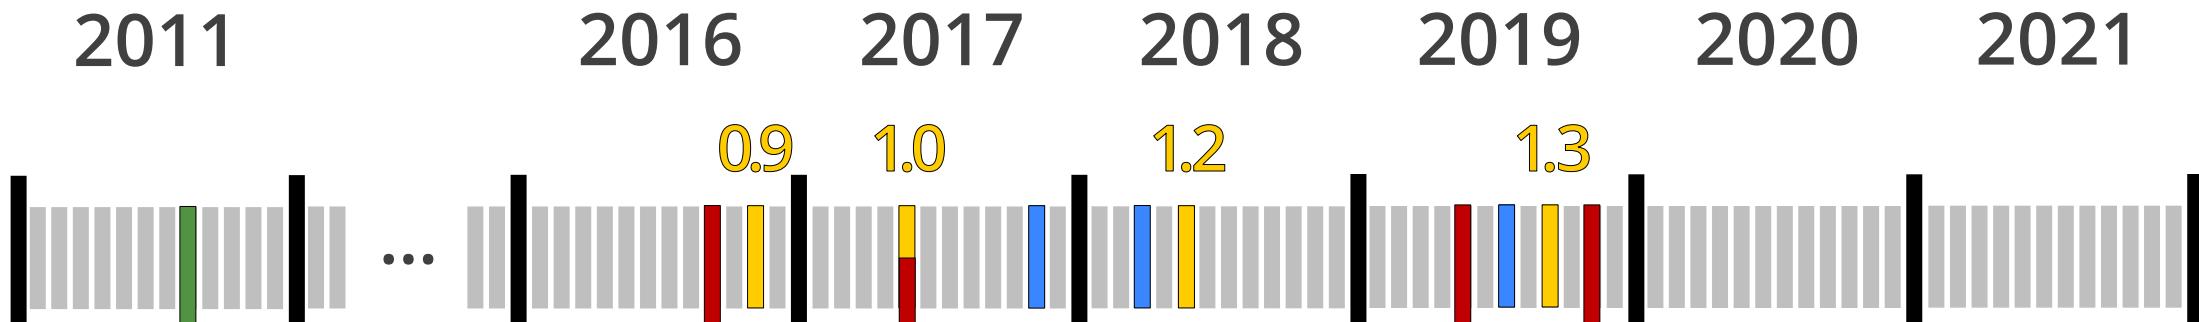
- 1979: BLAS                      Basic Linear Algebra Subprograms
- 2001: Sparse BLAS              an extension to BLAS (little uptake)
- 2013: GraphBLAS                an effort to define standard building blocks  
for graph algorithms in the language of linear algebra



S. McMillan @ SEI Research Review (Carnegie Mellon University, 2015):  
*Graph algorithms on future architectures*

# GRAPHBLAS TIMELINE

Book – Papers – GraphBLAS standards – SuiteSparse:GraphBLAS releases



**Graph Algorithms  
in the Language  
of Linear Algebra**

**Mathematical  
foundations,  
HPEC**

**C API,  
GABB@  
IPDPS**

**LAGraph, SuiteSparse:  
GrAPL @IPDPS**  
**SuiteSparse:  
GraphBLAS,  
TOMS**

- 2020:
- SuiteSparse:GraphBLAS v4.0.1draft
  - C++ API proposal
  - Distributed API proposal

# GRAPH ALGORITHMS IN LINEAR ALGEBRA

Notation:  $n = |V|, m = |E|$ . The complexity cells contain asymptotic bounds.

Takeaway: The majority of common graph algorithms can be expressed efficiently in LA.

problem	algorithm	canonical complexity $\Theta$	LA-based complexity $\Theta$
breadth-first search		$m$	$m$
single-source shortest paths	Dijkstra	$m + n \log n$	$n^2$
	Bellman-Ford	$mn$	$mn$
all-pairs shortest paths	Floyd-Warshall	$n^3$	$n^3$
minimum spanning tree	Prim	$m + n \log n$	$n^2$
	Borůvka	$m \log n$	$m \log n$
maximum flow	Edmonds-Karp	$m^2n$	$m^2n$
maximal independent set	greedy	$m + n \log n$	$mn + n^2$
	Luby	$m + n \log n$	$m \log n$

Based on the table in J. Kepner:

*Analytic Theory of Power Law Graphs*,  
SIAM Workshop for HPC on Large Graphs, 2008



See also L. Dhulipala, G.E. Blelloch, J. Shun:  
*Theoretically Efficient Parallel Graph Algorithms  
Can Be Fast and Scalable*, SPAA 2018

# KEY FEATURES OF GRAPHBLAS

- **Portable:** supports x86; GPUs (WIP), Arm (planned)
- **Efficient:** within one order-of-magnitude compared to hand-tuned code
- **Concise:** most textbook algorithms can be expressed with a few operations
- **Composable:** the output of an algorithm can be used as an input of a subsequent algorithm
- **Flexible:** can express algorithms on typed graphs and property graphs

# Theoretical foundations of the GraphBLAS

# DENSE MATRIX MULTIPLICATION

Definition:

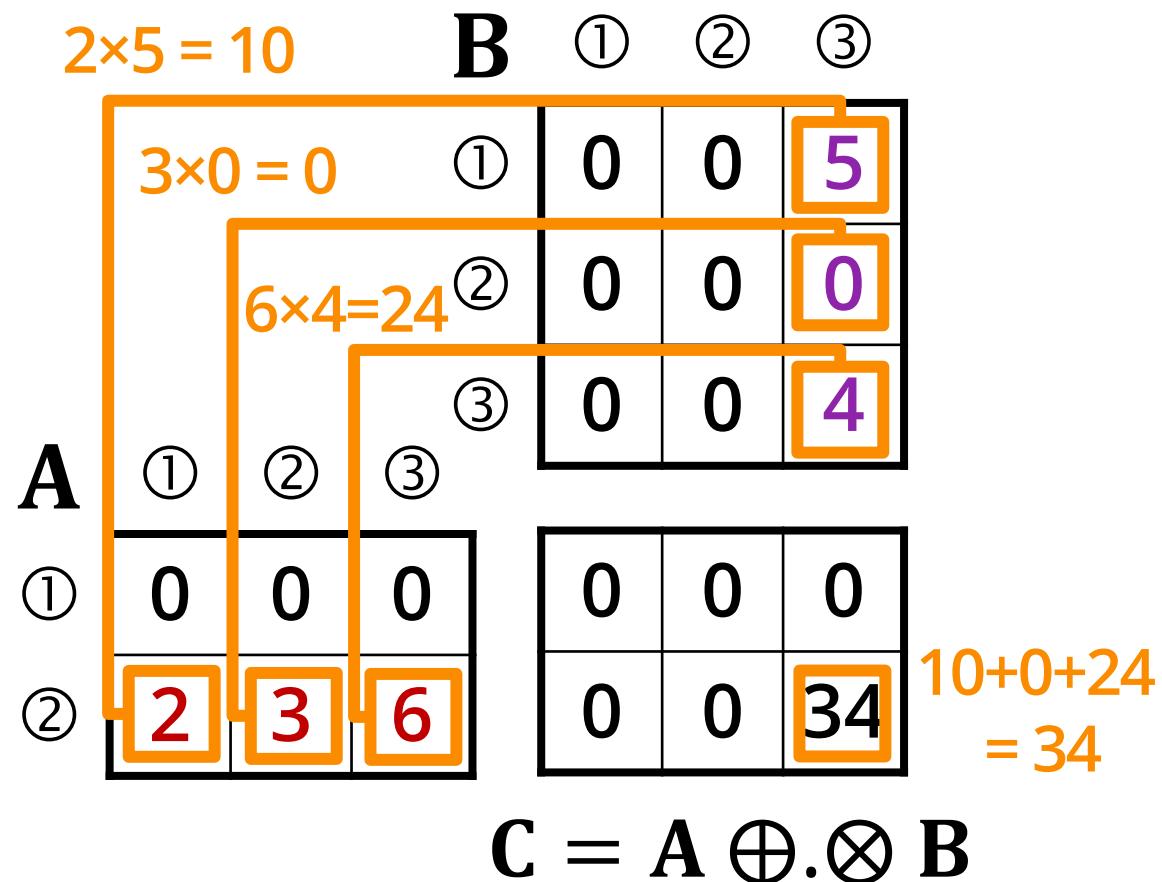
$$\mathbf{C} = \mathbf{AB}$$

$$\mathbf{C}(i, j) = \sum_k \mathbf{A}(i, k) \times \mathbf{B}(k, j)$$

Example:

$$\begin{aligned}\mathbf{C}(2,3) &= \mathbf{A}(2,1) \times \mathbf{B}(1,3) + \\ &\quad \mathbf{A}(2,2) \times \mathbf{B}(2,3) + \\ &\quad \mathbf{A}(2,3) \times \mathbf{B}(3,3)\end{aligned}$$

$$= 2 \times 5 + 3 \times 0 + 6 \times 4 = 34$$



# SPARSE MATRIX MULTIPLICATION

Definition:

$$\mathbf{C} = \mathbf{AB} = \mathbf{A} \oplus_{\cdot} \otimes \mathbf{B}$$

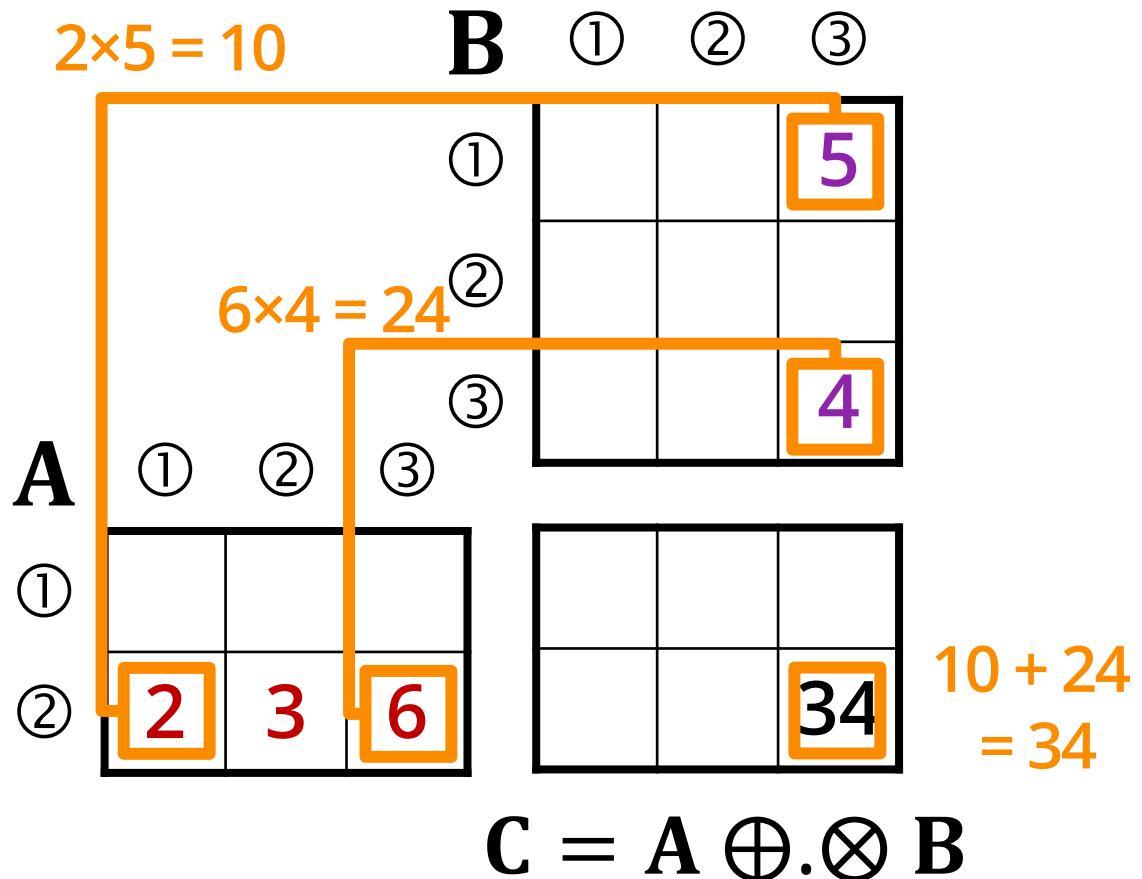
$$\mathbf{C}(i,j) = \bigoplus_{k \in \text{ind}(\mathbf{A}(i,:)) \cap \text{ind}(\mathbf{B}(:,j))} \mathbf{A}(i,k) \otimes \mathbf{B}(k,j)$$

Sparse matrix multiplication only evaluates the multiplication operator  $\otimes$  for positions where there is a non-zero element in both  $\mathbf{A}(i,k)$  and  $\mathbf{B}(k,j)$ .

Example:

$$\begin{aligned}\mathbf{C}(2,3) &= \mathbf{A}(2,1) \times \mathbf{B}(1,3) + \\ &\quad \mathbf{A}(2,3) \times \mathbf{B}(3,3)\end{aligned}$$

$$= 2 \times 5 + 6 \times 4 = 34$$



# MATRIX MULTIPLICATION $C = A \oplus.\otimes B$

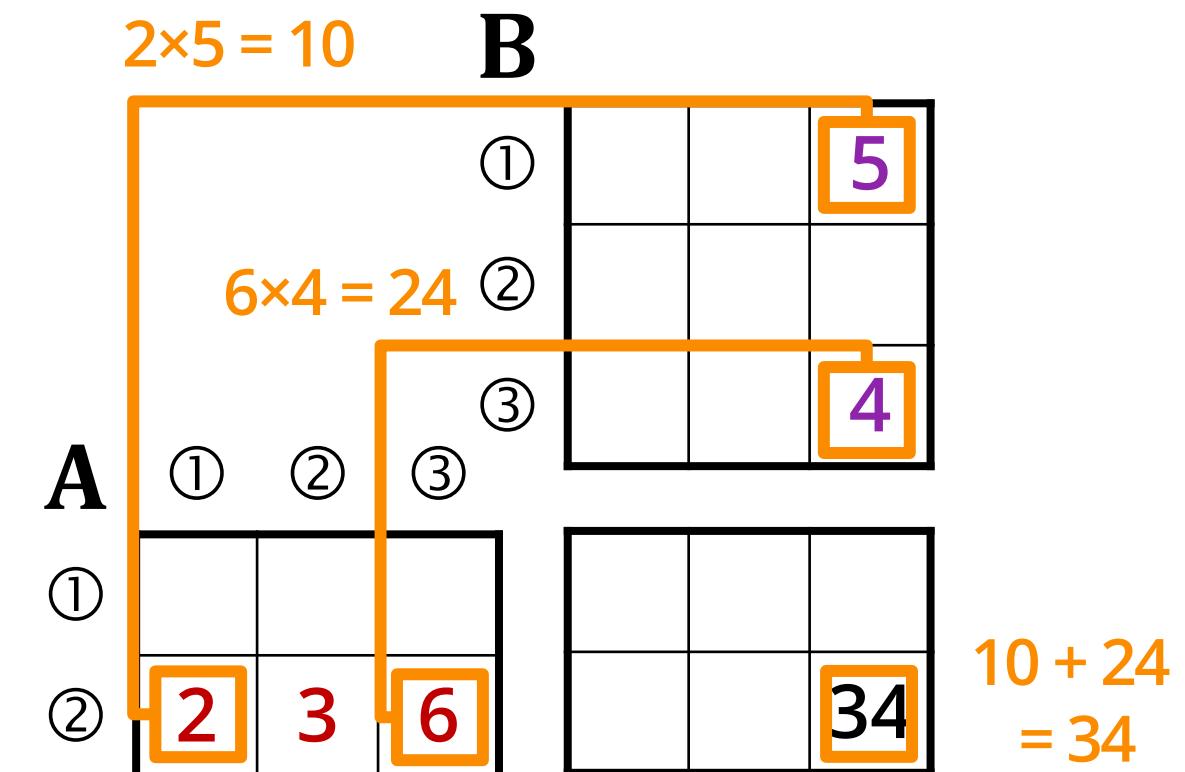
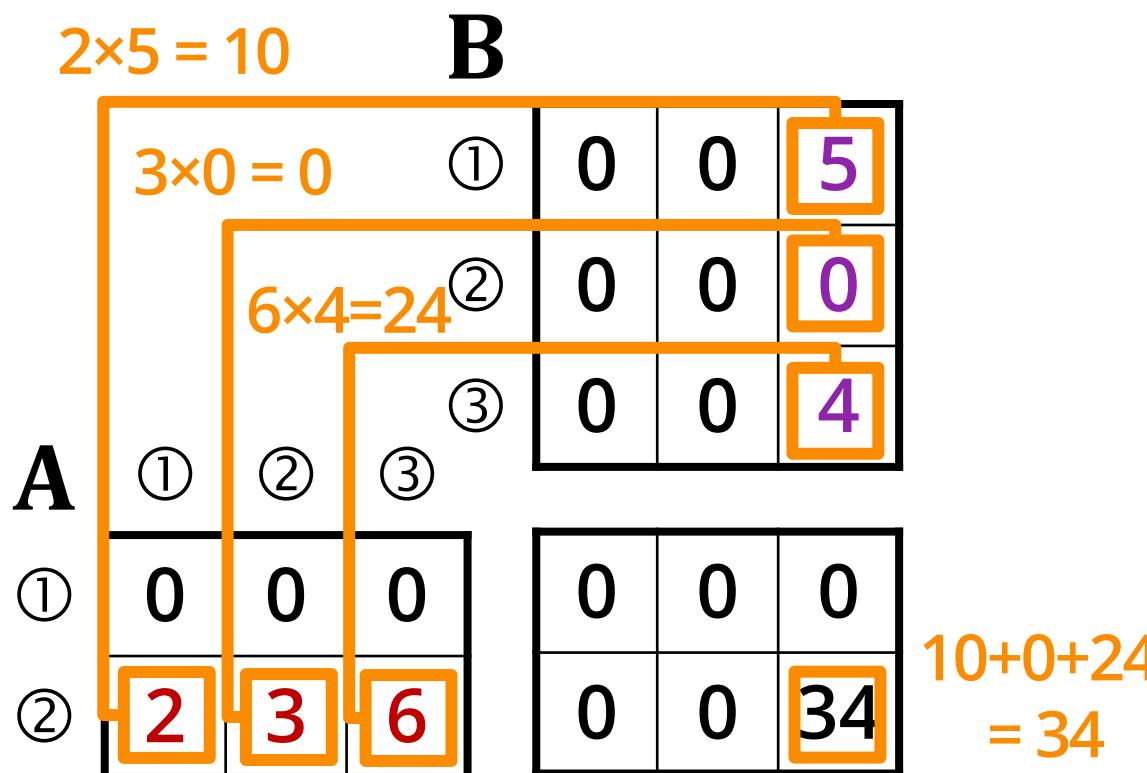
Multiplication on dense matrices

$$C(i, j) = \bigoplus_j A(i, k) \otimes B(k, j)$$

Multiplication on sparse matrices

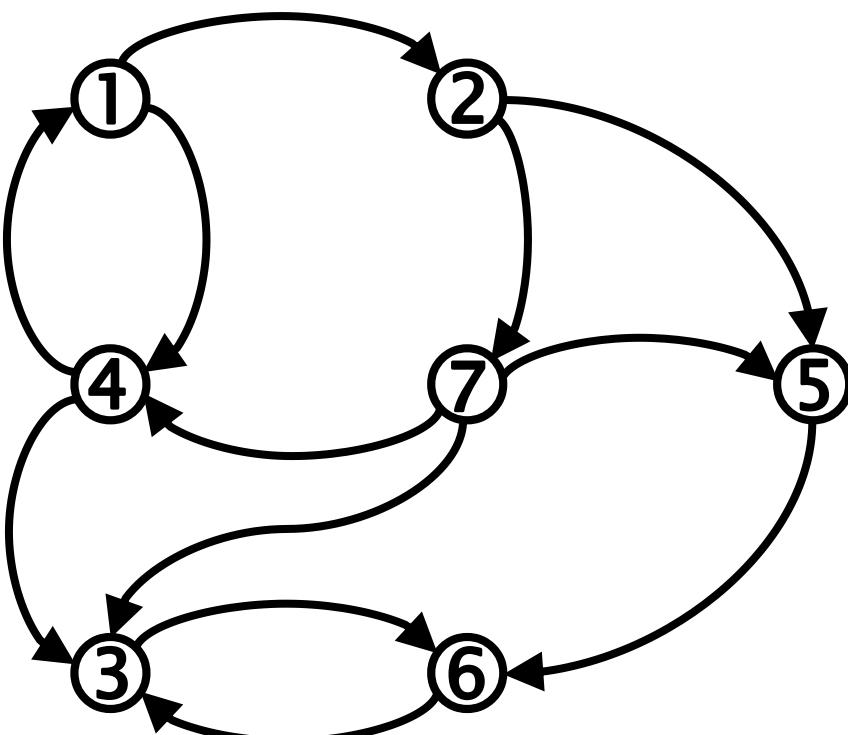
$$C(i, j) = \bigoplus_{k \in \text{ind}(A(i,:)) \cap \text{ind}(B(:,j))} A(i, k) \otimes B(k, j)$$

Example:  $C = A \oplus.\otimes B$



# ADJACENCY MATRIX

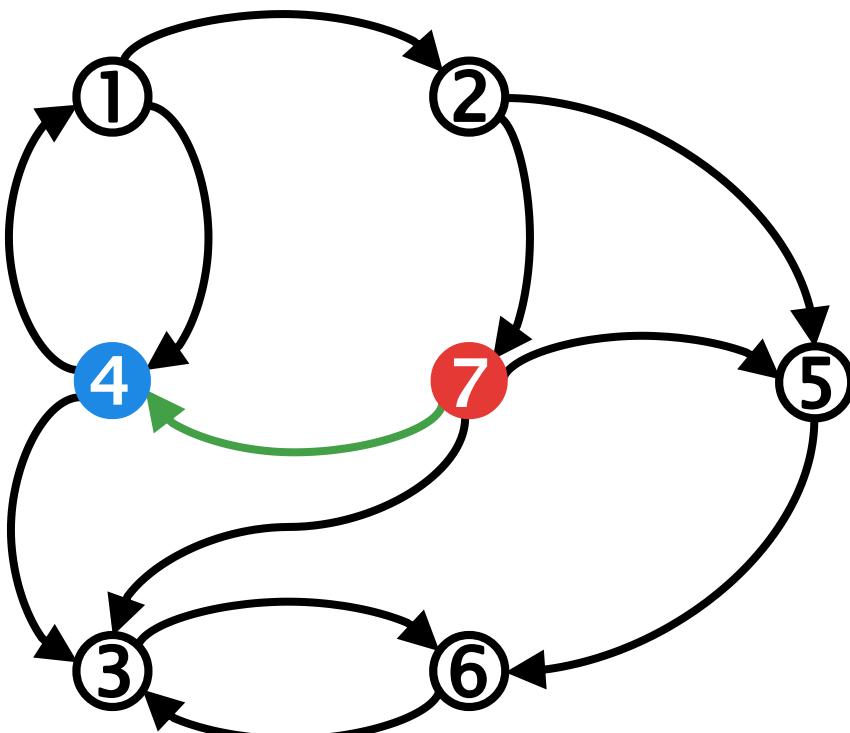
$$A_{ij} = \begin{cases} 1 & \text{if } (v_i, v_j) \in E \\ 0 & \text{if } (v_i, v_j) \notin E \end{cases}$$



A	①	②	③	④	⑤	⑥	⑦
①		1	1				
②				1	1		
③						1	
④	1		1				
⑤						1	
⑥				1			
⑦		1	1	1			

# ADJACENCY MATRIX

$$A_{ij} = \begin{cases} 1 & \text{if } (v_i, v_j) \in E \\ 0 & \text{if } (v_i, v_j) \notin E \end{cases}$$



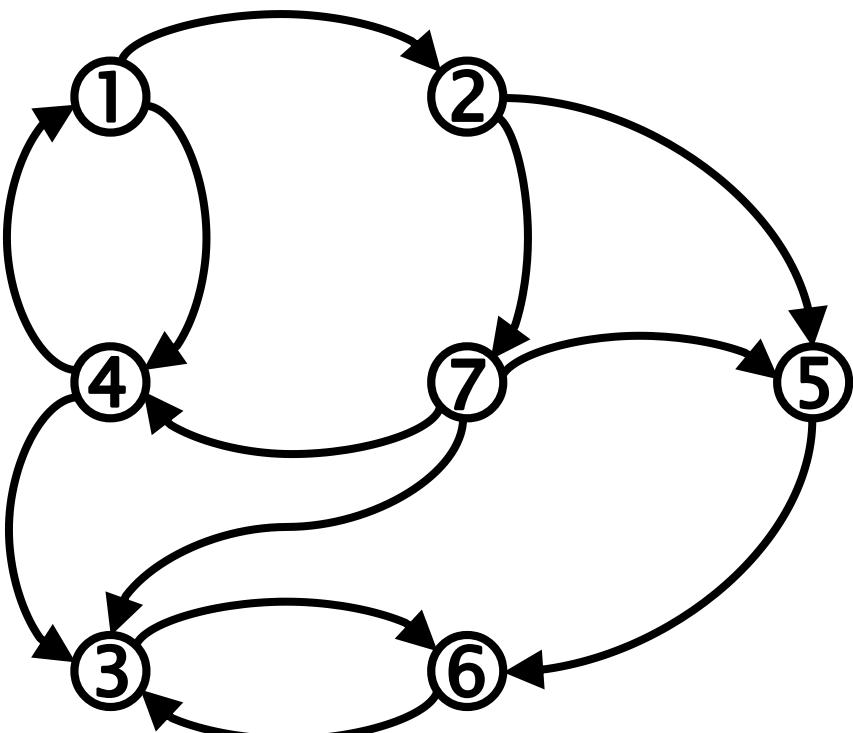
target

source 7

A	①	②	③	④	⑤	⑥	⑦
①		1	1				
②					1	1	
③							1
④	1		1				
⑤							1
⑥				1			
⑦				1	1	1	

# ADJACENCY MATRIX

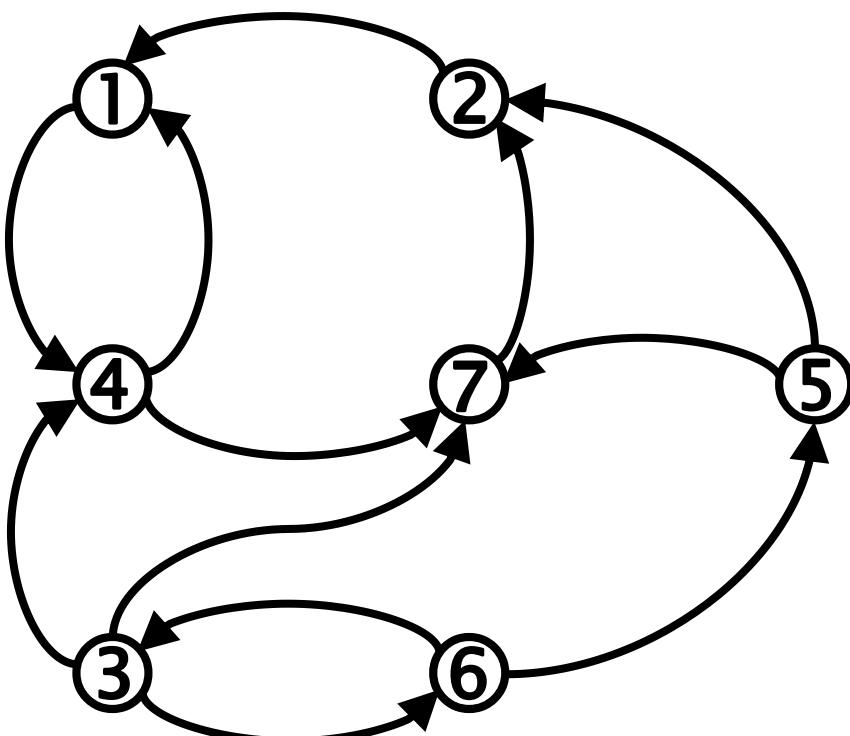
$$A_{ij} = \begin{cases} 1 & \text{if } (v_i, v_j) \in E \\ 0 & \text{if } (v_i, v_j) \notin E \end{cases}$$



A	①	②	③	④	⑤	⑥	⑦
①		1	1				
②				1	1		
③						1	
④	1		1				
⑤						1	
⑥			1				
⑦		1	1	1			

# ADJACENCY MATRIX TRANSPOSED

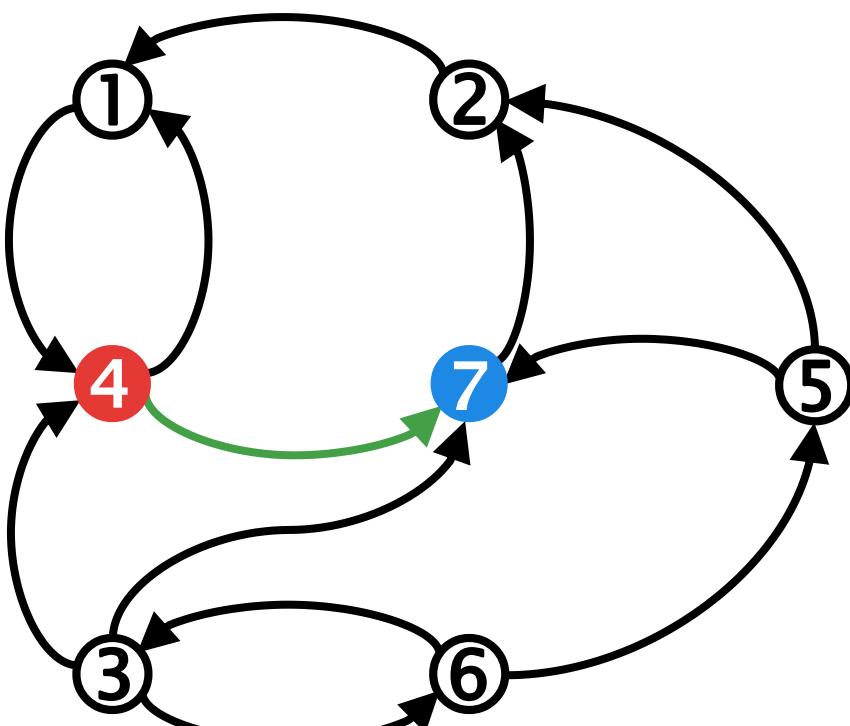
$$A_{ij}^T = \begin{cases} 1 & \text{if } (v_j, v_i) \in E \\ 0 & \text{if } (v_j, v_i) \notin E \end{cases}$$



A <sup>T</sup>	①	②	③	④	⑤	⑥	⑦
①					1		
②	1						
③			1	1	1		
④	1					1	
⑤		1					1
⑥			1		1		
⑦	1						

# ADJACENCY MATRIX TRANSPOSED

$$A_{ij}^T = \begin{cases} 1 & \text{if } (v_j, v_i) \in E \\ 0 & \text{if } (v_j, v_i) \notin E \end{cases}$$

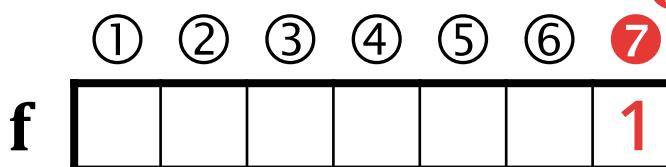
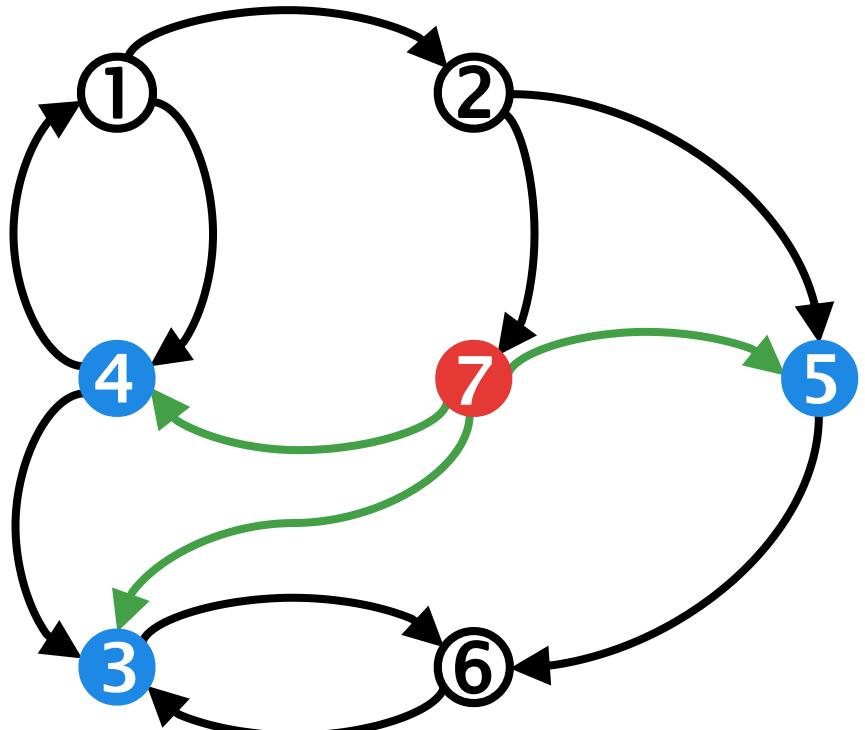


source

target

$A^T$	①	②	③	④	⑤	⑥	⑦
①				1			
②	1						
③			1	1	1		
④	1						1
⑤		1					
⑥			1		1		
⑦	1						

# GRAPH TRAVERSAL WITH MATRIX MULTIPLICATION



$fA^k$  means  $k$  hops in the graph

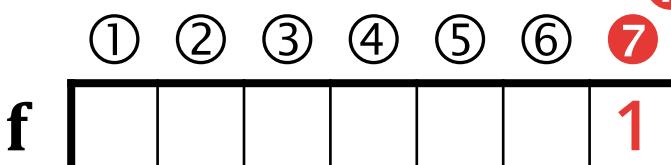
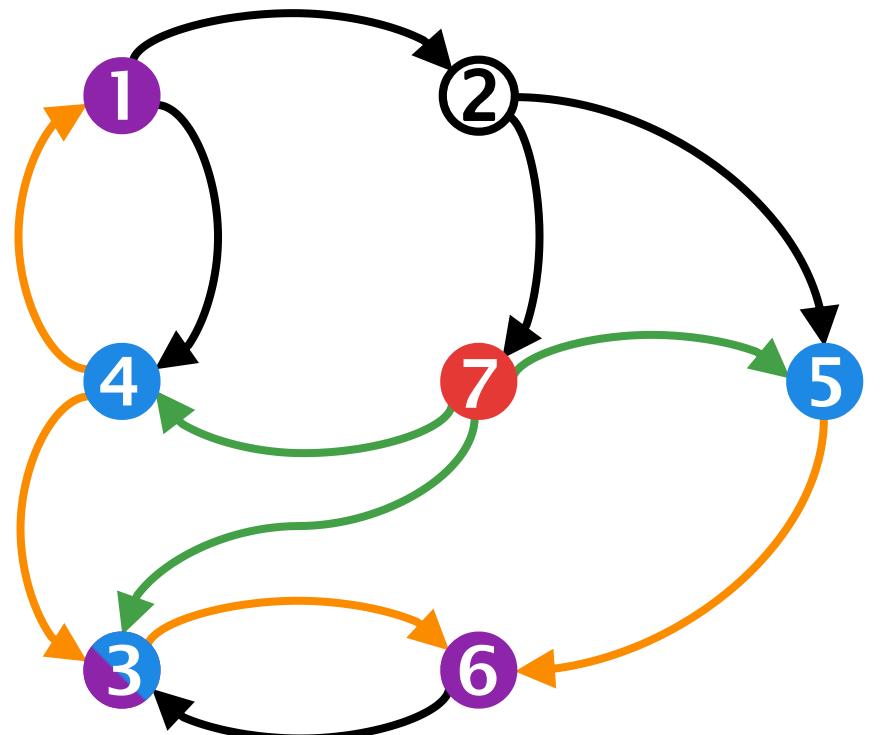
$A$

	1	2	3	4	5	6	7
1	1		1				
2				1		1	
3							1
4	1		1				
5							1
6			1				
7			1	1	1		

		1	1	1		
1	2	3	4	5	6	7

one hop:  $fA$

# GRAPH TRAVERSAL WITH MATRIX MULTIPLICATION



one hop:  $fA$

		1	1	1		
1	2	3	4	5	6	7

$fA^k$  means  $k$  hops in the graph

$A$

1	2	3	4	5	6	7
1		1	1			
2				1	1	1
3					1	
4			1	1		
5					1	
6					1	
7					1	1

$A$

1	2	3	4	5	6	7
1		1	1		1	1
2				1	1	1
3					1	
4			1	1	1	
5					1	
6					1	1
7					1	1

one hop:  $fA$

two hops:  $fA^2$

1	2	3	4	5	6	7
1		1		1	2	

# GRAPHBLAS SEMIRINGS\*

The  $\langle D, \oplus, \otimes, 0 \rangle$  algebraic structure is a GraphBLAS semiring if

- $\langle D, \oplus, 0 \rangle$  is a commutative monoid using the addition operation  $\oplus: D \times D \rightarrow D$ , where  $\forall a, b, c \in D$ , if the following hold:
  - Commutative       $a \oplus b = b \oplus a$
  - Associative       $(a \oplus b) \oplus c = a \oplus (b \oplus c)$
  - Identity           $a \oplus 0 = a$
- The multiplication operator is a closed binary operator  $\otimes: D \times D \rightarrow D$ .

The mathematical definition of a semiring requires that  $\otimes$  is a monoid and distributes over  $\oplus$ . GraphBLAS omits these requirements.

# SEmirings

semiring	domain	$\oplus$	$\otimes$	0	graph semantics
any-pair	$\{T, F\}$	any pair	F	traversal step	
integer arithmetic	$\mathbb{N}$	+	$\times$	0	number of paths
min-plus	$\mathbb{R} \cup \{+\infty\}$	min	+	$+\infty$	shortest path

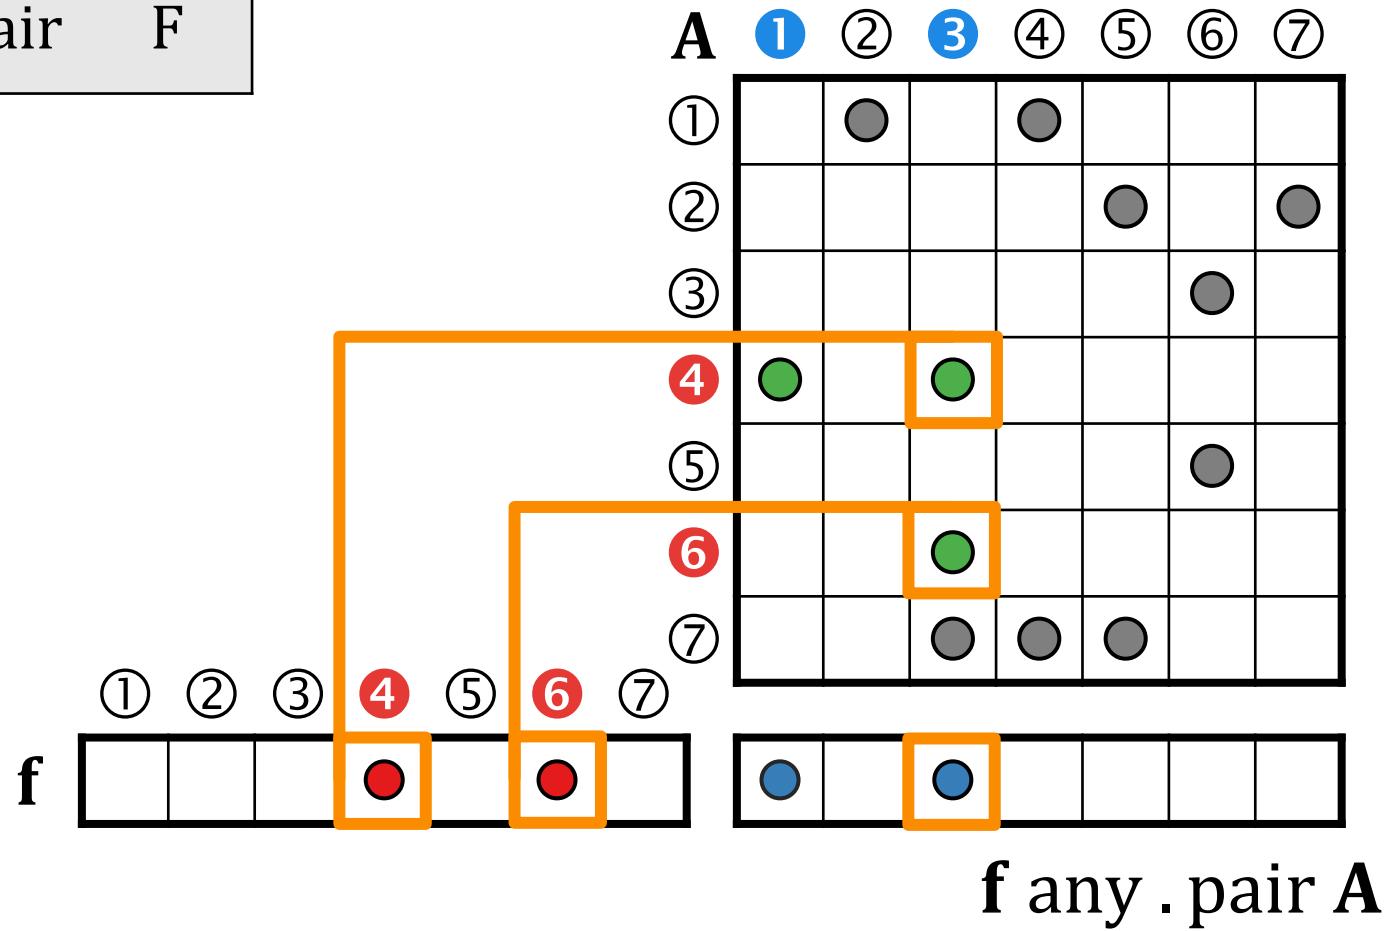
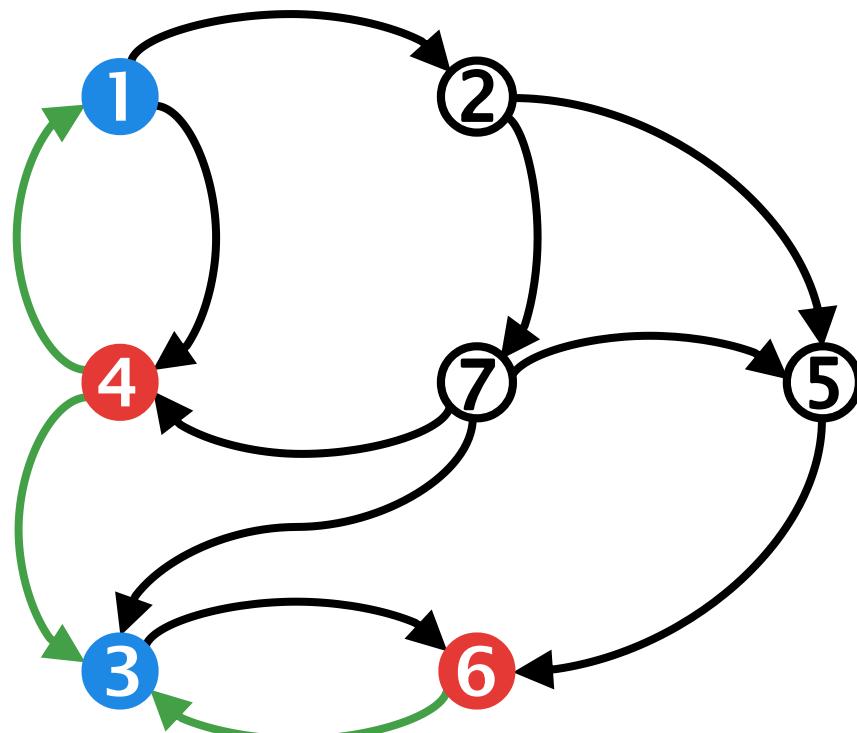
The default semiring is the conventional one:

- $\otimes$  defaults to the arithmetic multiplication operator.
- $\oplus$  defaults to the arithmetic addition operator.

# MATRIX-VECTOR MULTIPLICATION SEMANTICS

semiring	domain	$\oplus$	$\otimes$	0
any-pair	{T, F}	any	pair	F

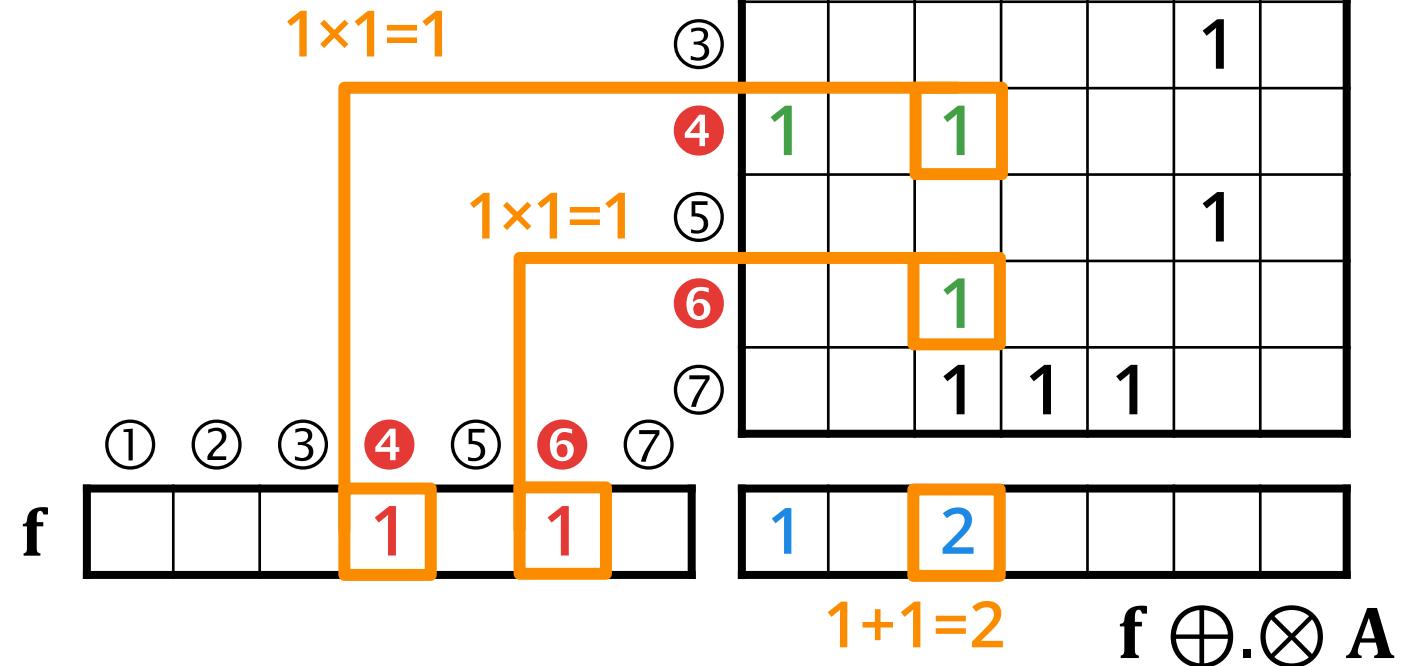
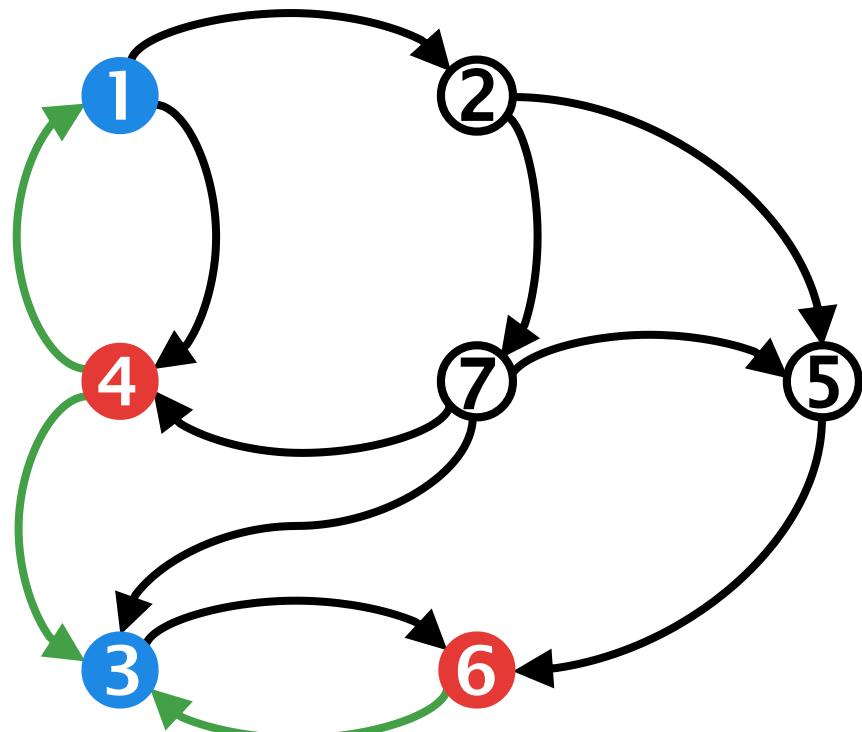
Semantics: traversal step



# MATRIX-VECTOR MULTIPLICATION SEMANTICS

semiring	domain	$\oplus$	$\otimes$	0
integer arithmetic	$\mathbb{N}$	$+$	$\times$	0

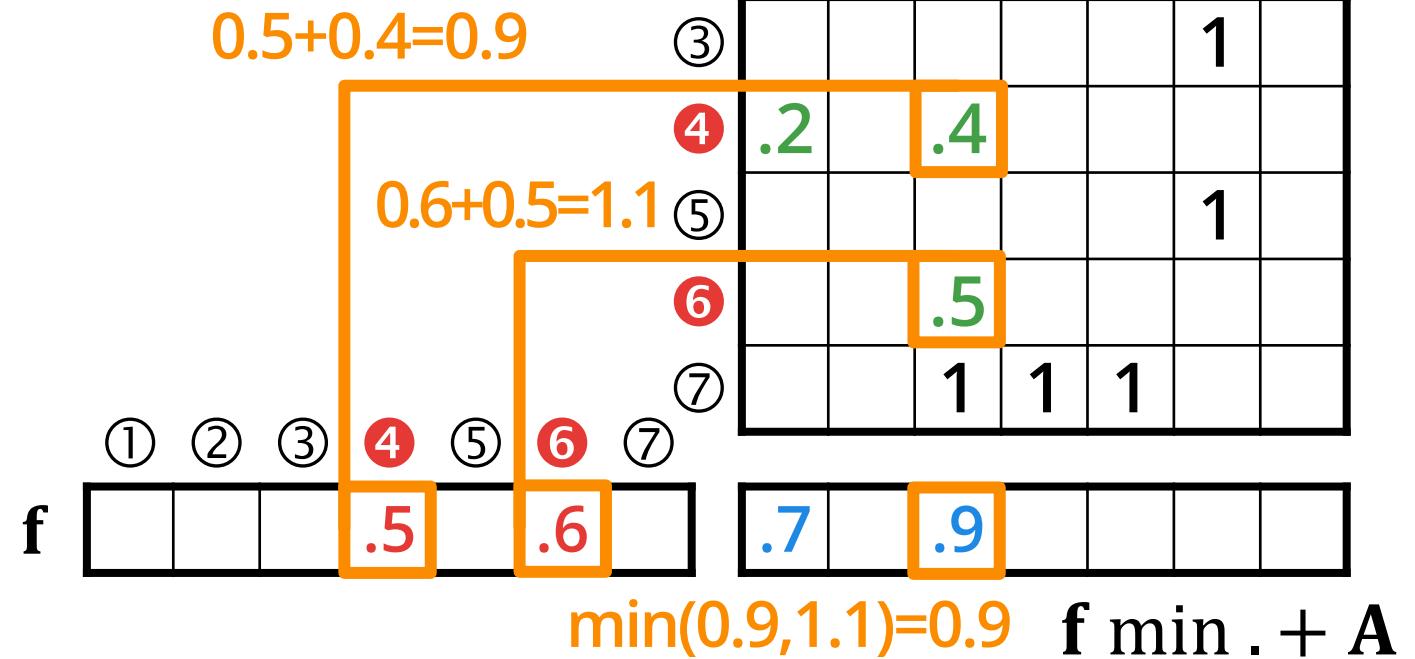
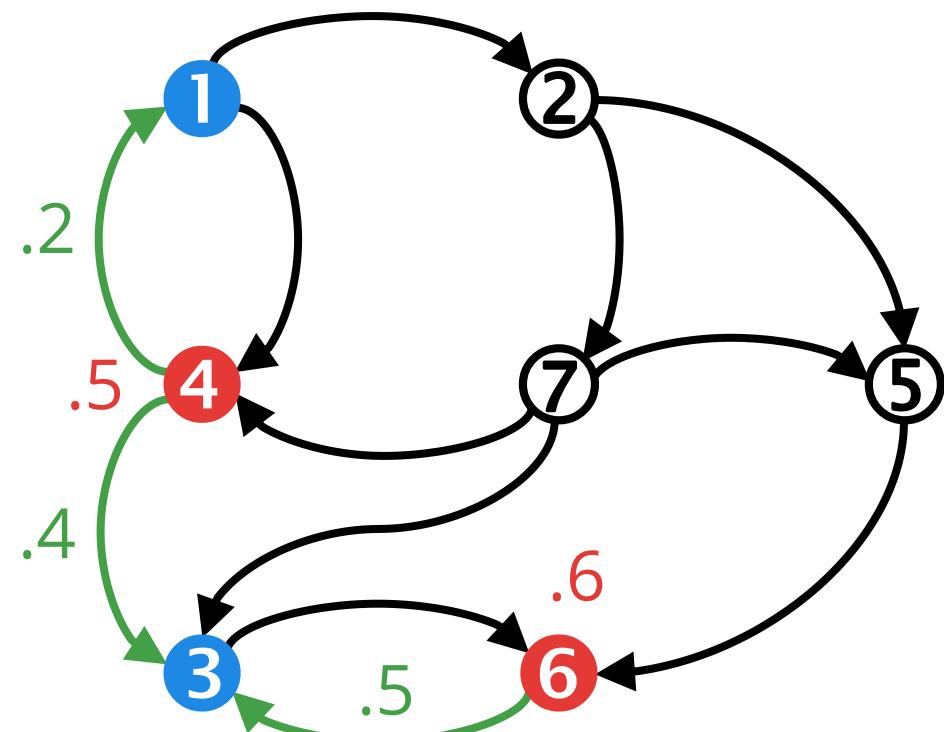
Semantics: number of paths



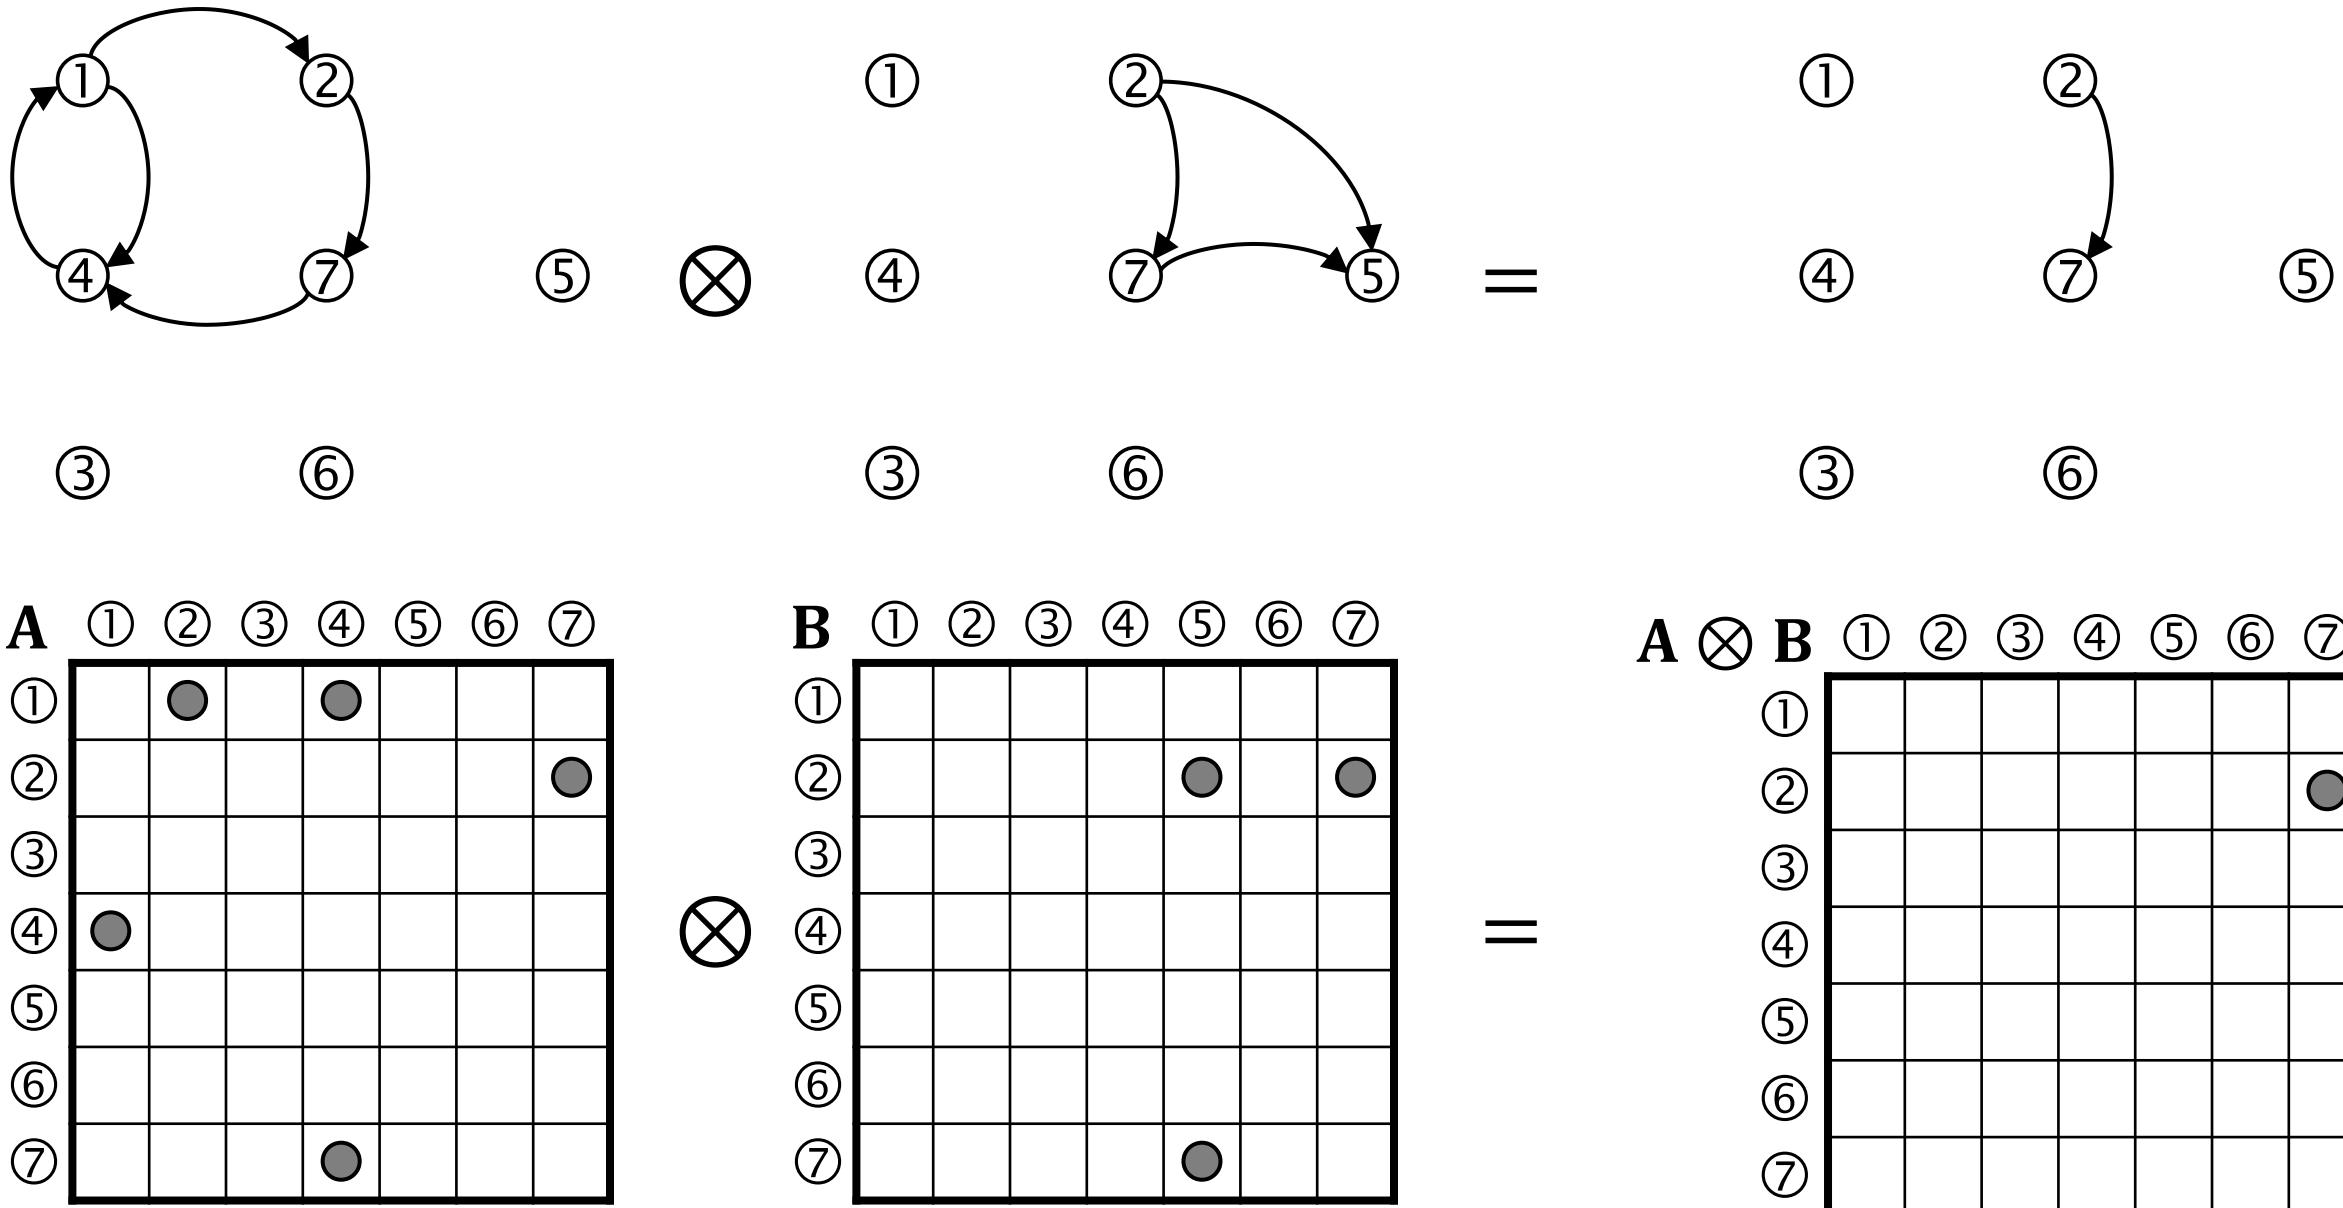
# MATRIX-VECTOR MULTIPLICATION SEMANTICS

semiring	domain	$\oplus$	$\otimes$	0
min-plus	$\mathbb{R} \cup \{+\infty\}$	min	+	$+\infty$

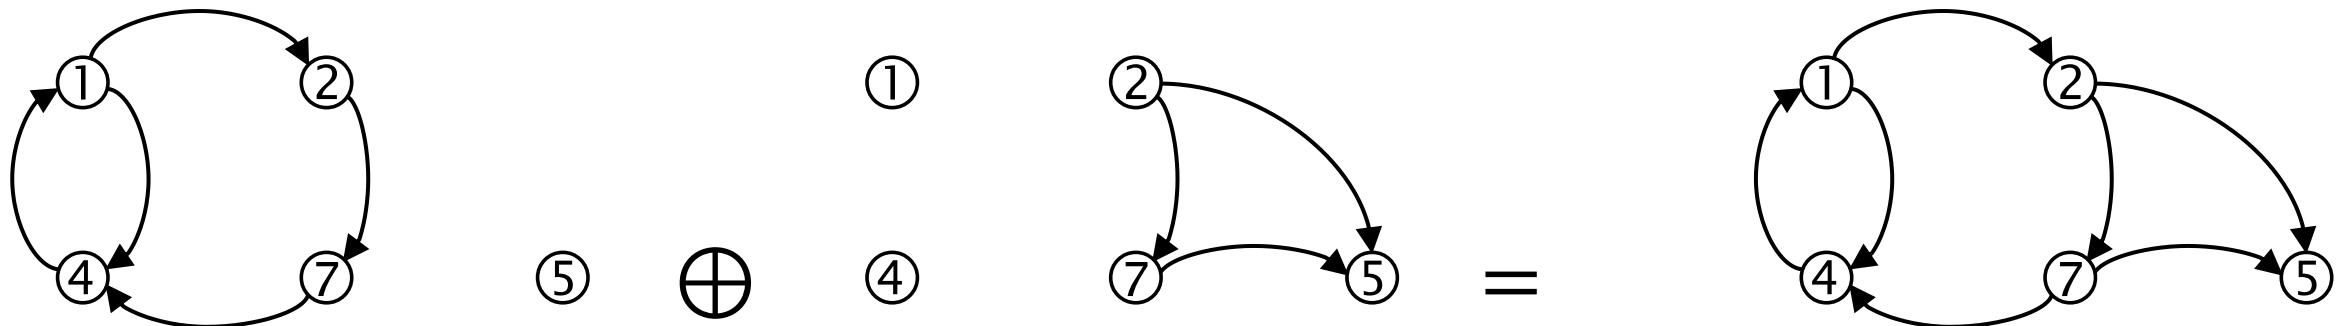
Semantics: shortest path



# ELEMENT-WISE MULTIPLICATION: $A \otimes B$



# ELEMENT-WISE ADDITION: $A \oplus B$



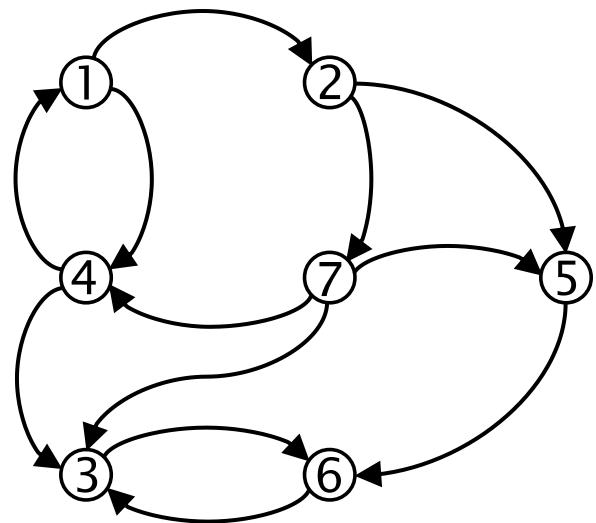
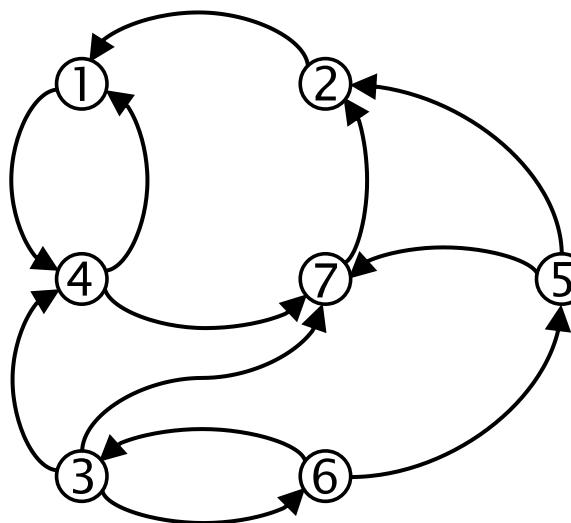
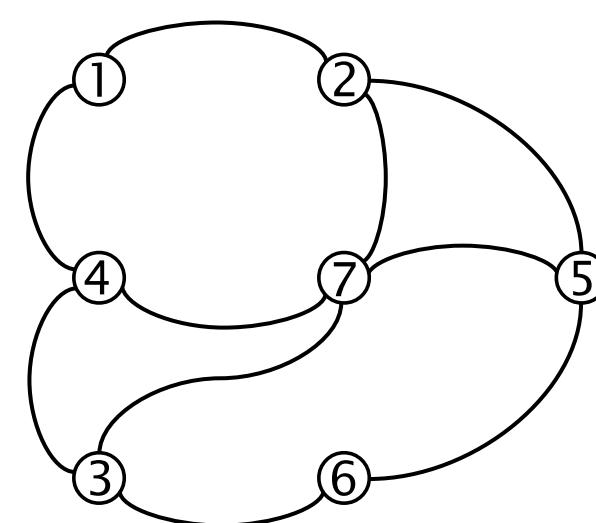
③            ⑥            ③            ⑥            ③            ⑥

<b>A</b>	1	2	3	4	5	6	7
1		●		●			
2							●
3							
4	●						
5							
6							
7				●			

<b>B</b>	1	2	3	4	5	6	7
1							
2					●		●
3							
4							
5							
6							
7					●		

<b><math>A \oplus B</math></b>	1	2	3	4	5	6	7
1		●		●			
2						●	●
3							
4	●						
5							
6							
7					●	●	

# TURNING A GRAPH INTO UNDIRECTED: $A \oplus A^T$

 $\oplus$  $=$ 

<b>A</b>	1	2	3	4	5	6	7
1		●		●			
2				●		●	
3							●
4	●		●				
5						●	
6			●		●		
7		●	●	●	●		

 $\oplus$ 

<b><math>A^T</math></b>	1	2	3	4	5	6	7
1				●			
2	●						
3				●		●	●
4	●					●	●
5		●					●
6			●		●		
7		●					

 $=$ 

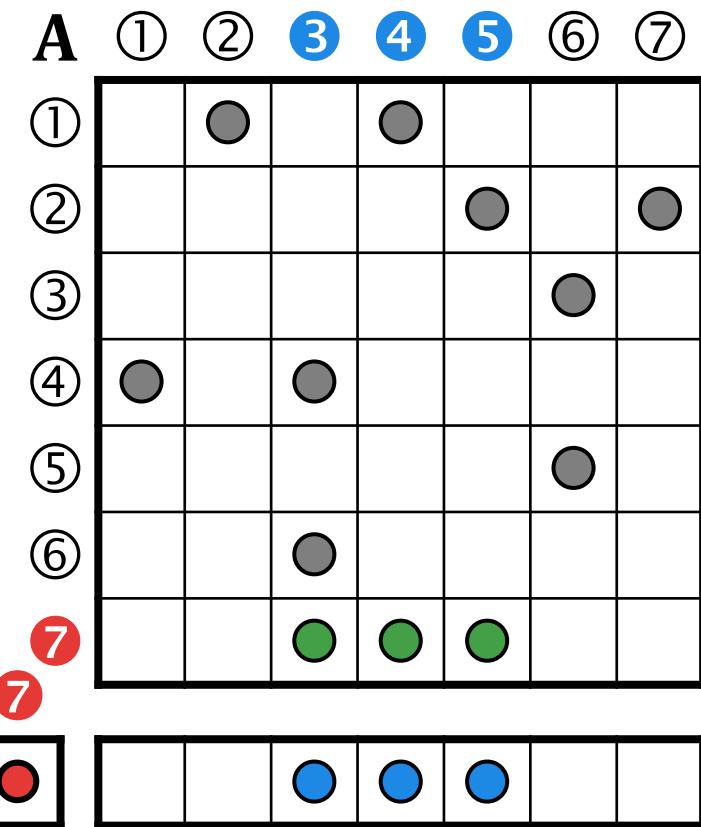
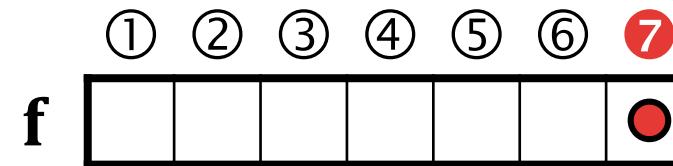
<b><math>A \oplus A^T</math></b>	1	2	3	4	5	6	7
1		●		●			
2	●			●		●	●
3				●		●	●
4	●					●	●
5		●					●
6			●		●		
7		●	●	●	●	●	●

# MASKING

Prevent redundant computations by reducing the scope of an operation

Operations can be executed

- **without a mask**
- with a regular mask
- with a complemented mask



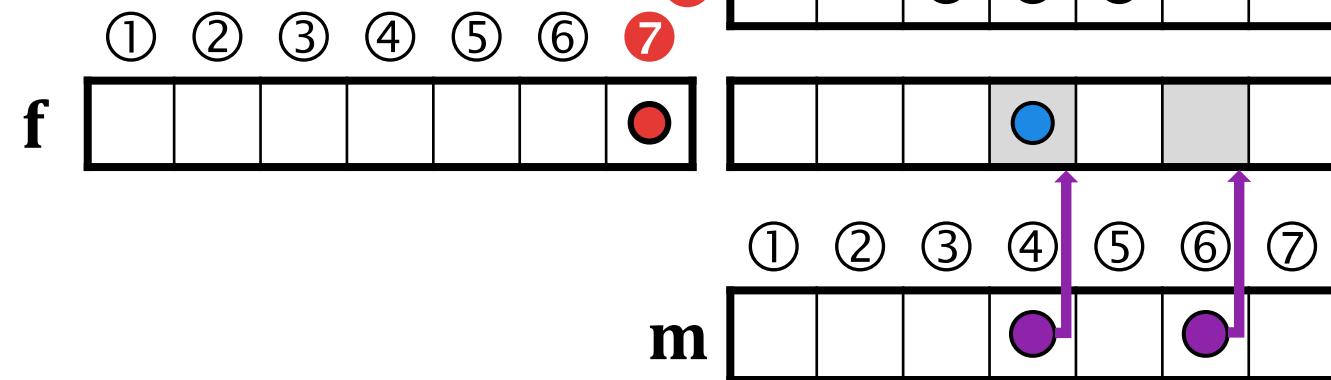
$$w = f \oplus . \otimes A$$

# MASKING

Prevent redundant computations by reducing the scope of an operation

Operations can be executed

- without a mask
- **with a regular mask**
- with a complemented mask



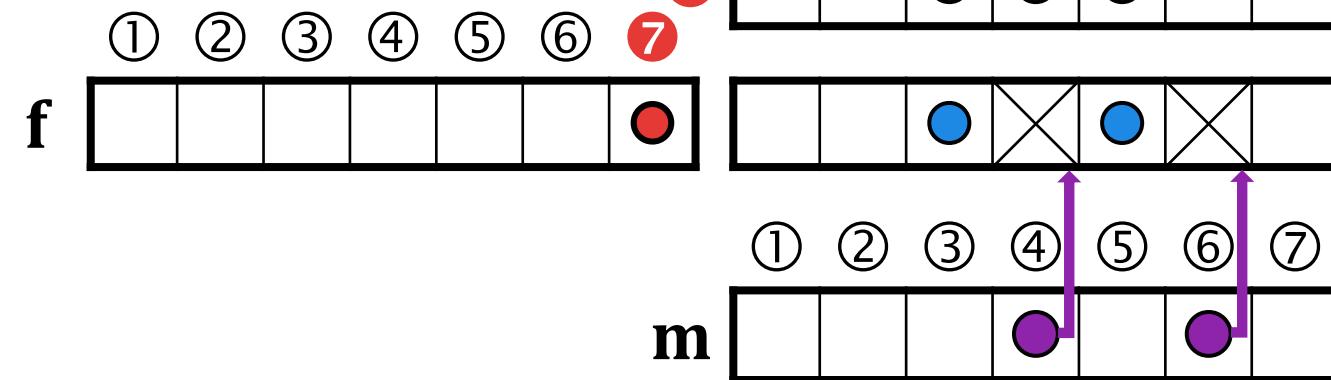
$$w\langle m \rangle = f \oplus_{.} \otimes A$$

# MASKING

Prevent redundant computations by reducing the scope of an operation

Operations can be executed

- without a mask
- with a regular mask
- **with a complemented mask**



$$w \langle \neg m \rangle = f \oplus . \otimes A$$

# NOTATION\*

- Symbols:

- **A, B, C, M** – matrices
- **u, v, w, m** – vectors
- $s$  – scalar
- $i, j$  – indices
- $\langle \mathbf{M} \rangle, \langle \mathbf{m} \rangle$  – masks

- Operators:

- $\oplus$  – addition
- $\otimes$  – multiplication
- $\oslash$  – division
- $^T$  – transpose

symbol	operation	notation
$\oplus \cdot \otimes$	matrix-matrix multiplication	$\mathbf{C} \langle \mathbf{M} \rangle = \mathbf{A} \oplus \cdot \otimes \mathbf{B}$
	vector-matrix multiplication	$\mathbf{w} \langle \mathbf{m} \rangle = \mathbf{v} \oplus \cdot \otimes \mathbf{A}$
	matrix-vector multiplication	$\mathbf{w} \langle \mathbf{m} \rangle = \mathbf{A} \oplus \cdot \otimes \mathbf{v}$
$\otimes$	element-wise multiplication (set intersection of patterns)	$\mathbf{C} \langle \mathbf{M} \rangle = \mathbf{A} \otimes \mathbf{B}$ $\mathbf{w} \langle \mathbf{m} \rangle = \mathbf{u} \otimes \mathbf{v}$
	element-wise addition (set union of patterns)	$\mathbf{C} \langle \mathbf{M} \rangle = \mathbf{A} \oplus \mathbf{B}$ $\mathbf{w} \langle \mathbf{m} \rangle = \mathbf{u} \oplus \mathbf{v}$
$f$	apply unary operator	$\mathbf{C} \langle \mathbf{M} \rangle = f(\mathbf{A})$ $\mathbf{w} \langle \mathbf{m} \rangle = f(\mathbf{v})$
	reduce to vector	$\mathbf{w} \langle \mathbf{m} \rangle = [\oplus_j \mathbf{A}(:,j)]$
$[\oplus \dots]$	reduce to scalar	$s = [\oplus_{ij} \mathbf{A}(i,j)]$
	transpose matrix	$\mathbf{C} \langle \mathbf{M} \rangle = \mathbf{A}^T$

Vectors can act as both column and row vectors.

(Notation omitted for accumulator, selection, extraction, assignment...)

# LINEAR ALGEBRAIC PRIMITIVES FOR GRAPHS #1

Element-wise addition:  
union of non-zero elements

The diagram shows two 4x4 sparse matrices with gray dots representing non-zero elements. A plus sign ( $\oplus$ ) is placed between them. The resulting matrix is shown on the right, where the non-zero elements are the union of the non-zero elements from both input matrices.

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$\oplus$

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

Element-wise multiplication:  
intersection of non-zero elements

The diagram shows two 4x4 sparse matrices with gray dots representing non-zero elements. A multiplication sign ( $\otimes$ ) is placed between them. The resulting matrix is shown on the right, where the non-zero elements are the intersection of the non-zero elements from both input matrices.

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$\otimes$

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

Sparse matrix times sparse vector:  
process incoming edges

The diagram shows a 4x4 sparse matrix with gray dots representing non-zero elements and a vertical sparse vector with gray dots representing non-zero elements. A sequence of operators ( $\oplus$ ,  $\cdot$ ,  $\otimes$ ) is placed between them. The resulting vector is shown on the right, where the non-zero elements are the result of processing incoming edges from the matrix.

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$\oplus \cdot \otimes$

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

Sparse vector times sparse matrix:  
process outgoing edges

The diagram shows a horizontal sparse vector with gray dots representing non-zero elements and a 4x4 sparse matrix with gray dots representing non-zero elements. A sequence of operators ( $\oplus$ ,  $\cdot$ ,  $\otimes$ ) is placed between them. The resulting vector is shown on the right, where the non-zero elements are the result of processing outgoing edges from the matrix.

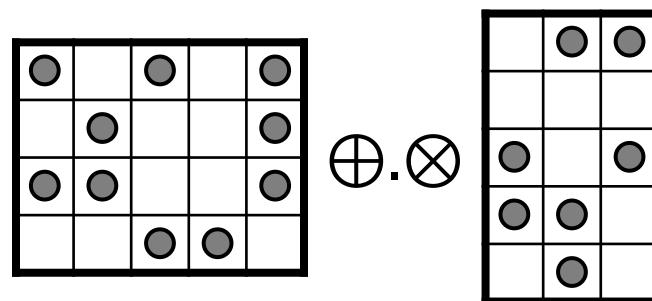
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$\oplus \cdot \otimes$

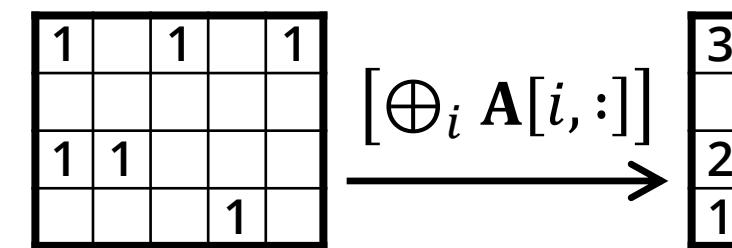
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

# LINEAR ALGEBRAIC PRIMITIVES FOR GRAPHS #2

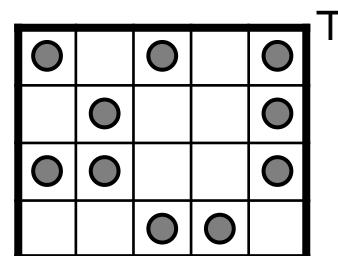
Sparse matrix times sparse matrix:  
process connecting outgoing edges



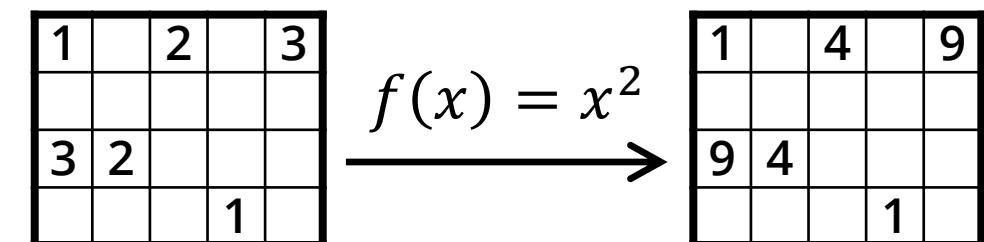
Reduction:  
aggregate values in each row



Matrix transpose:  
reverse edges



Apply:  
apply unary operator on all values



# Graph algorithms in GraphBLAS

## Breadth-first search

# BFS: BREADTH-FIRST SEARCH

- Algorithm:
  - Start from a given vertex
  - “Explore all neighbour vertices at the present level prior to moving on to the vertices at the next level” [Wikipedia]
- Variants:
  - **Levels** compute traversal level for each vertex
  - **Parents** compute parent for each vertex
  - **MSBFS** start traversal from multiple source vertices

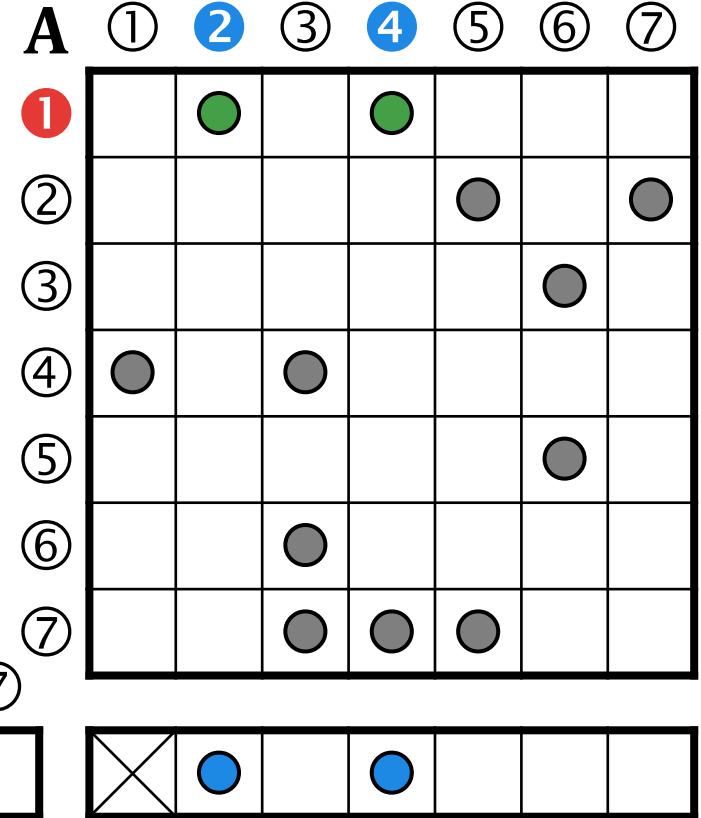
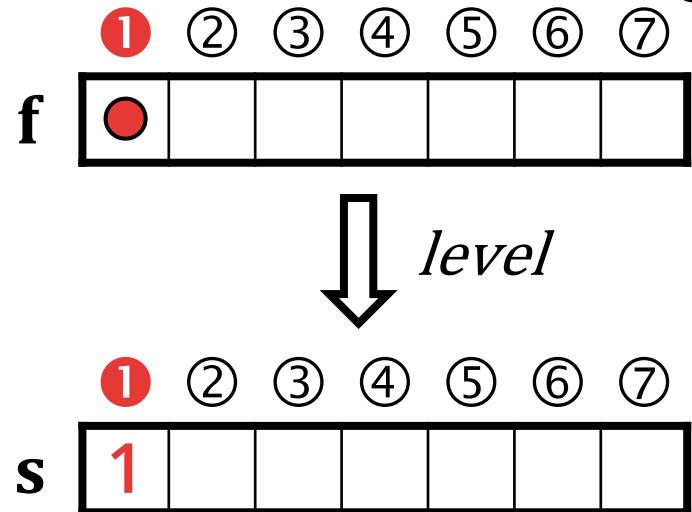
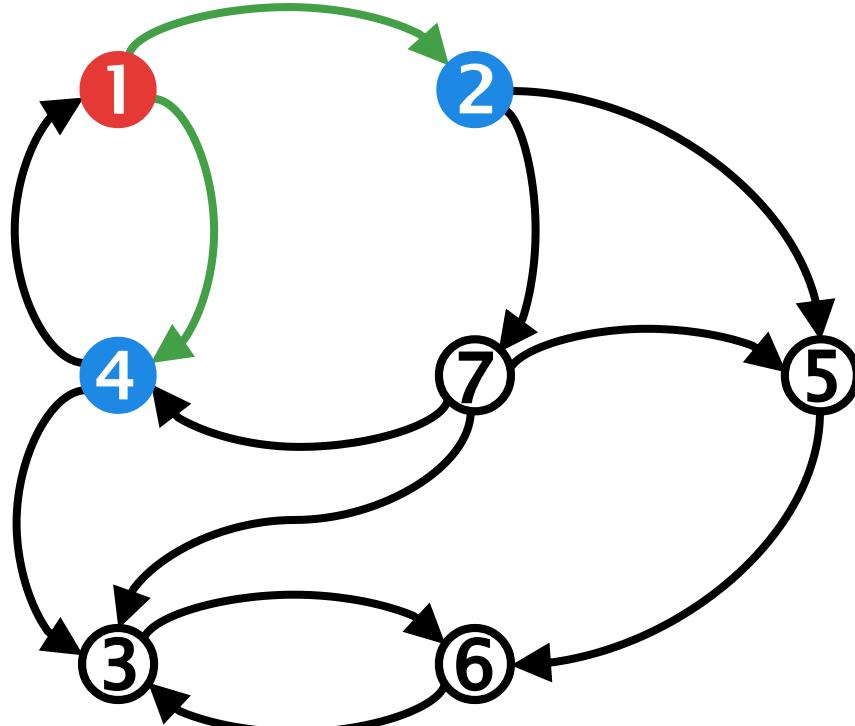
# Graph algorithms in GraphBLAS

Breadth-first search / Levels

# BFS - LEVELS

semiring	domain	$\oplus$	$\otimes$	0
any-pair	{T, F}	any	pair	F

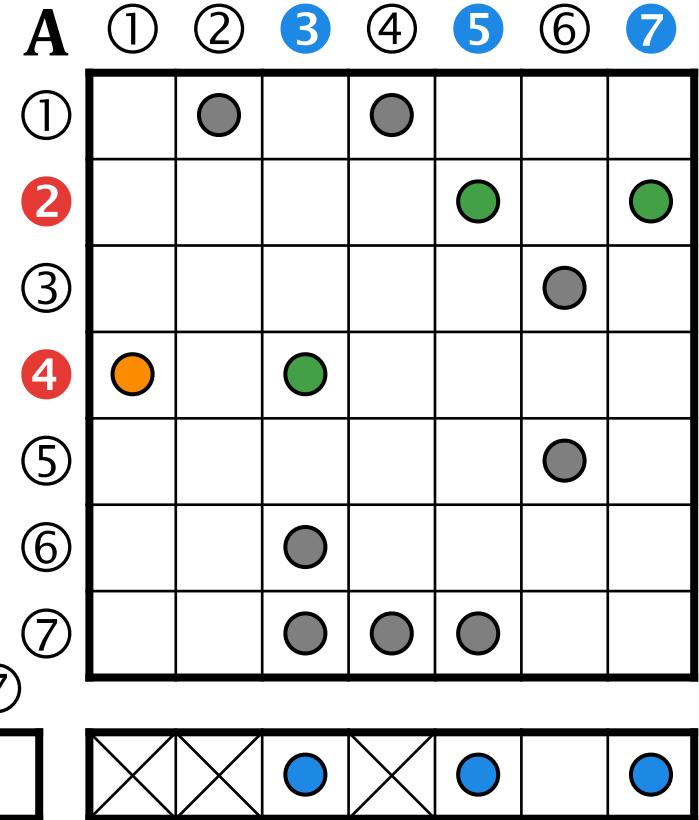
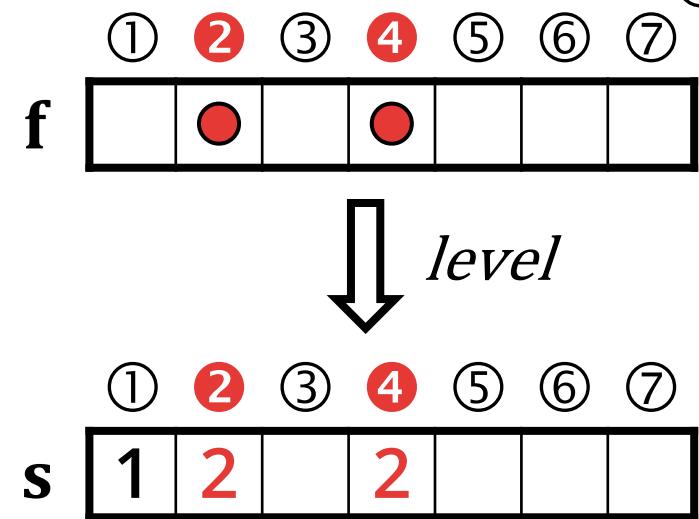
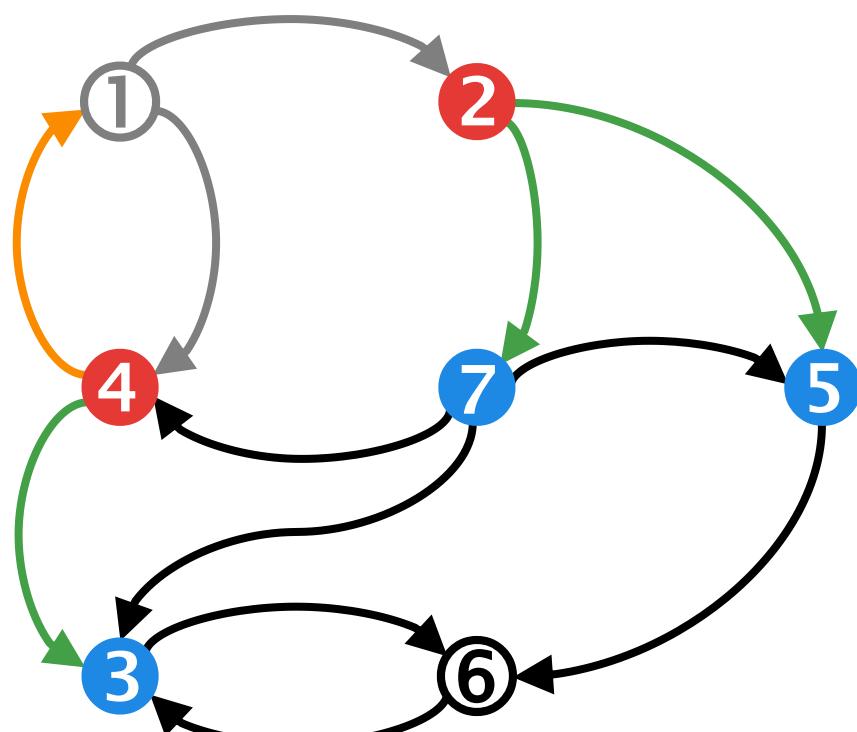
*level* = 1



# BFS - LEVELS

semiring	domain	$\oplus$	$\otimes$	0
any-pair	{T, F}	any	pair	F

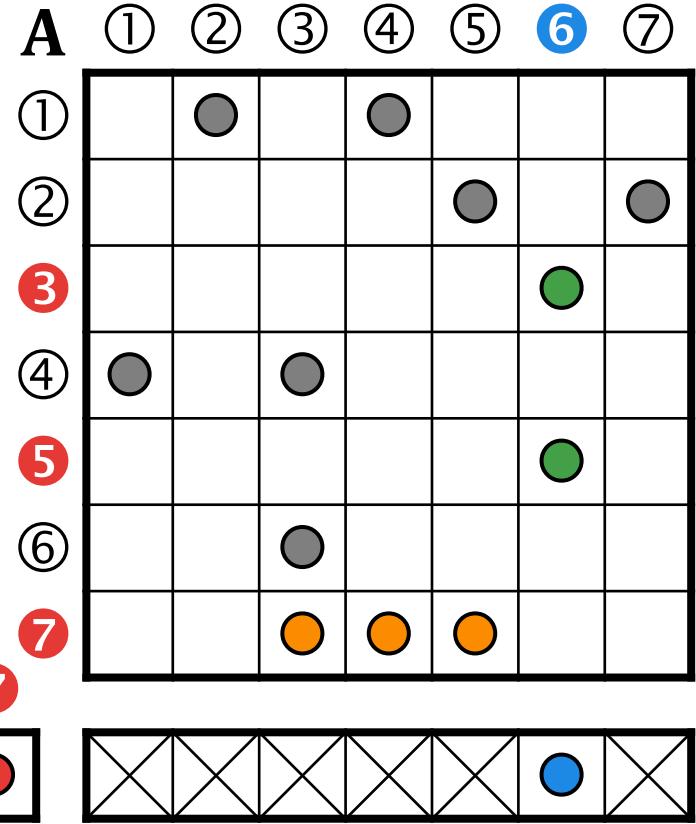
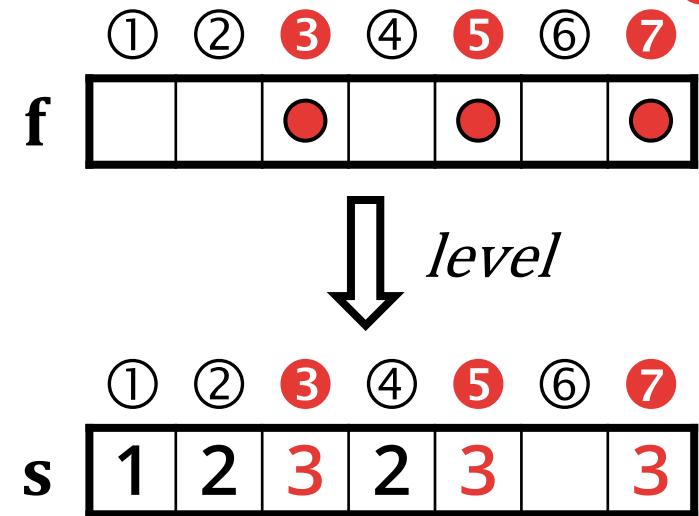
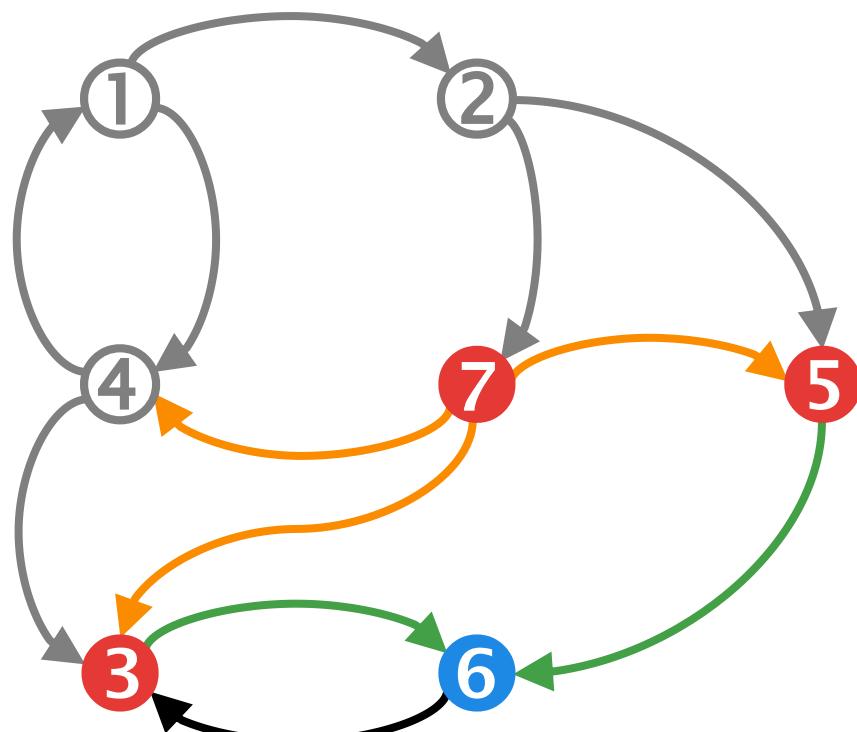
*level* = 2



# BFS - LEVELS

semiring	domain	$\oplus$	$\otimes$	0
any-pair	{T, F}	any	pair	F

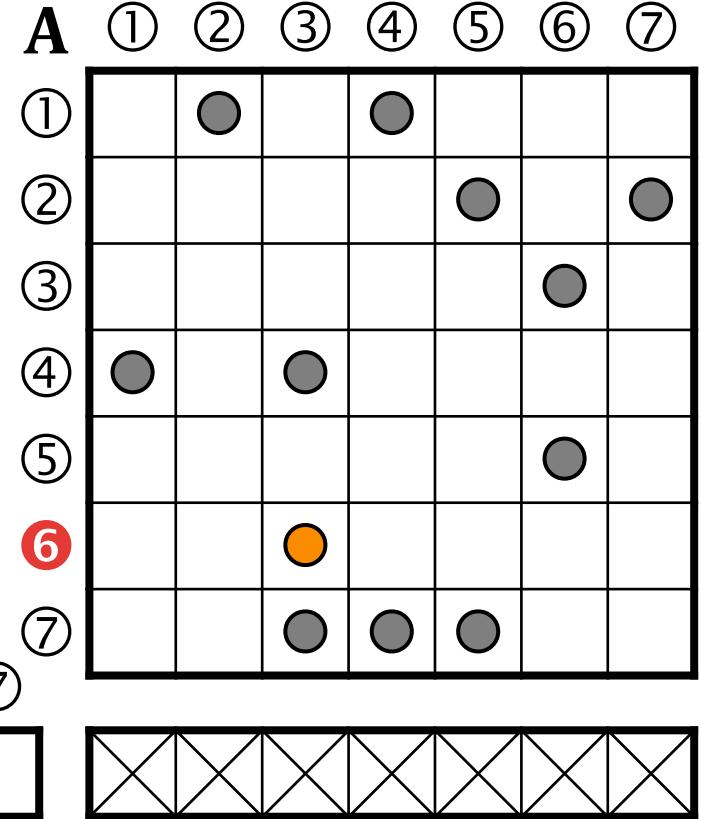
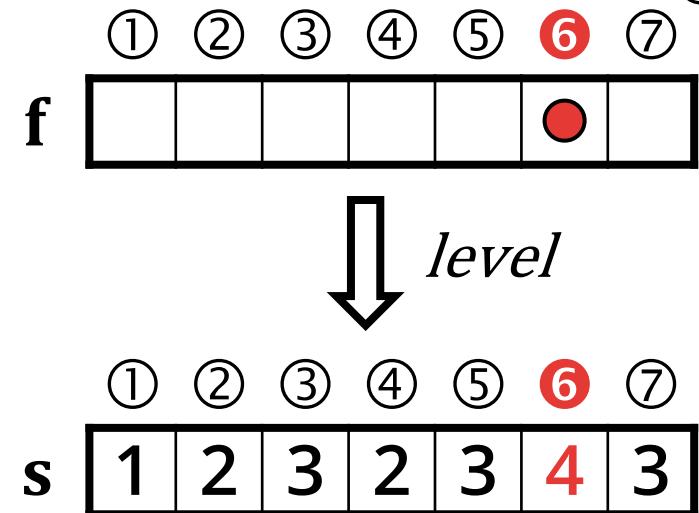
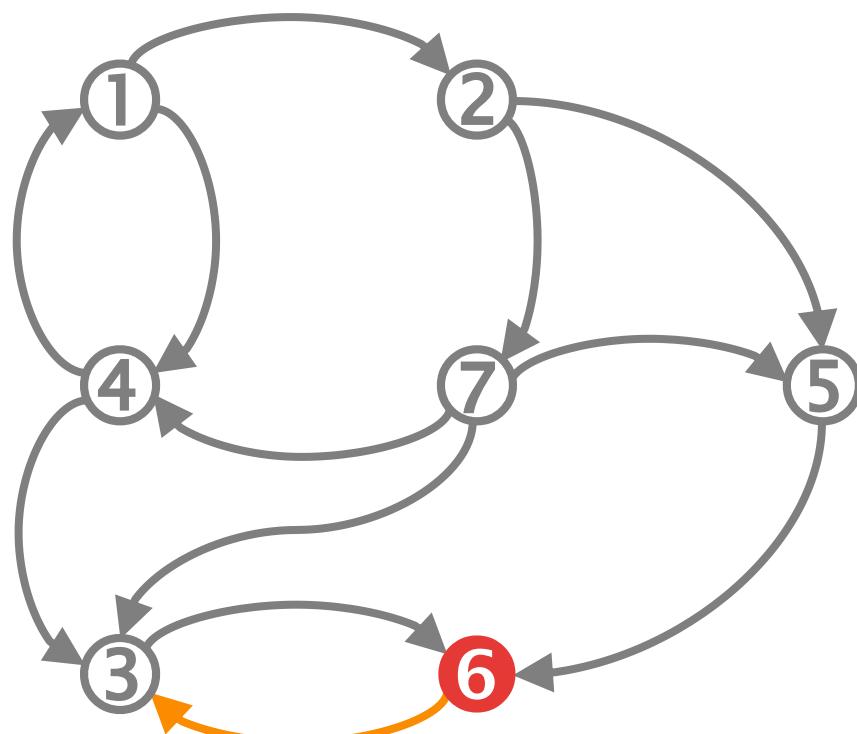
*level* = 3



# BFS - LEVELS

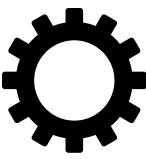
semiring	domain	$\oplus$	$\otimes$	0
any-pair	{T, F}	any	pair	F

*level* = 4



$f \langle \neg s \rangle = f$  any. pair A

$f$  is empty  
→ terminate



# BFS – LEVELS: ALGORITHM

- **Input:** adjacency matrix  $A$ , source vertex  $s$ , #vertices  $n$
- **Output:** vector of visited vertices  $v$  (integer)
- **Workspace:** frontier vector  $f$  (Boolean)

1.  $f(s) = T$
2. for  $level = 1$  to  $n - 1$  \*terminate earlier if  $f$  is empty
3.    $s\langle f \rangle = level$  assign the level value to the vertices in the frontier
4.   clear( $f$ ) clear the frontier  $f$
5.    $f\langle \neg s \rangle = f$  any. pair  $A$  advance the frontier

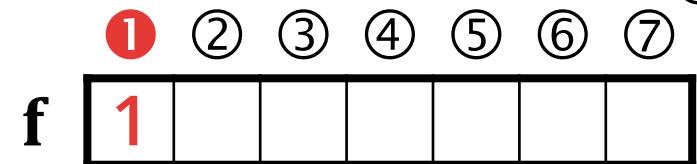
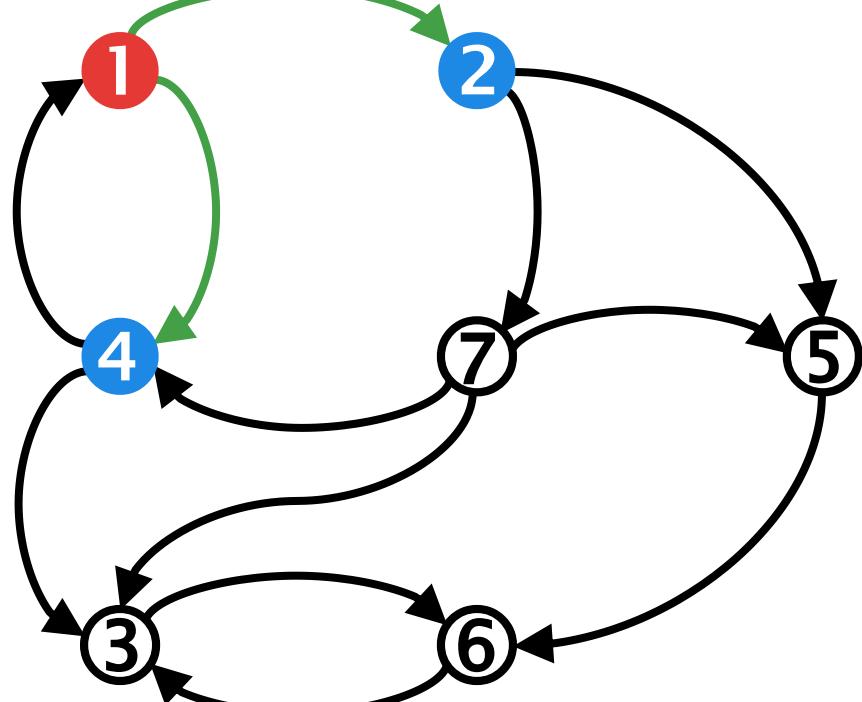
# Graph algorithms in GraphBLAS

Breadth-first search / Parents

# BFS – PARENTS

semiring	domain	$\oplus$	$\otimes$	0
min-first	$\mathbb{N}$	min	first	0

$$\text{first}(x, y) = x$$



A

1	2	3	4	5	6	7
1		green		green		
2					grey	grey
3					grey	grey
4	grey		grey			
5						grey
6			grey			
7		grey	grey	grey		

$f(\neg p)$

X	1	1				
---	---	---	--	--	--	--

$$f(\neg p) = f \text{ min . first } A$$

$\swarrow p(f) = f$

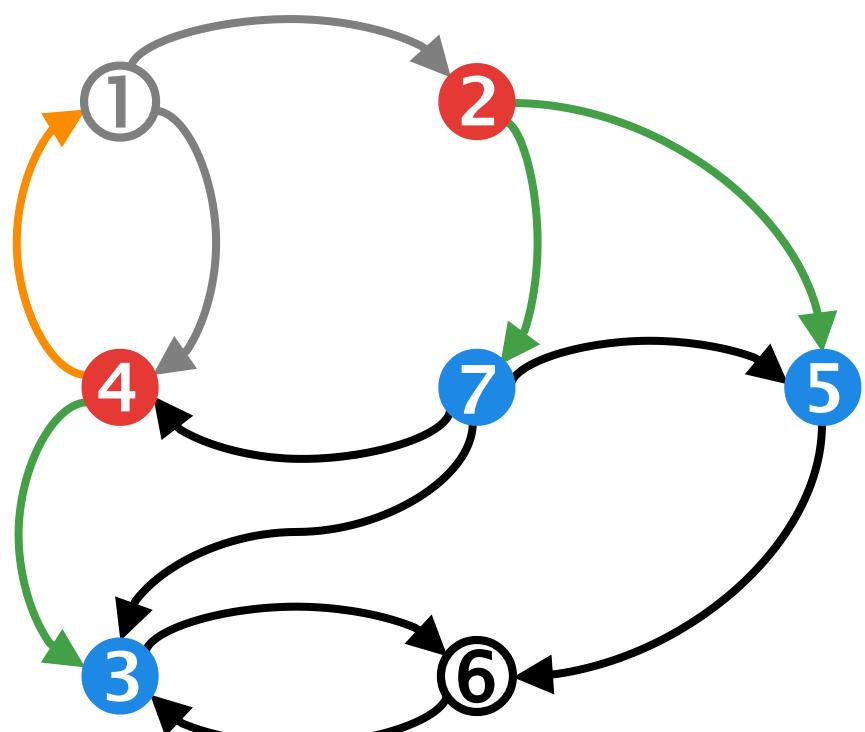
$\downarrow f(f) = \text{id}$

1	2	3	4	5	6	7
	2		4			

# BFS – PARENTS

semiring	domain	$\oplus$	$\otimes$	0
min-first	$\mathbb{N}$	min	first	0

$$\text{first}(x, y) = x$$

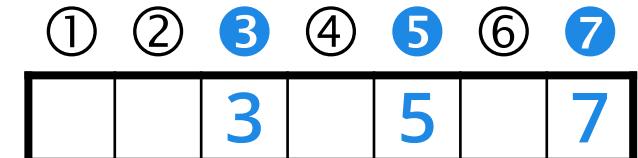


A	1	2	3	4	5	6	7
1		●		●			
2					●	●	●
3						●	
4	●			●			
5						●	
6				●			
7		●	●	●			



$\swarrow p(f) = f$

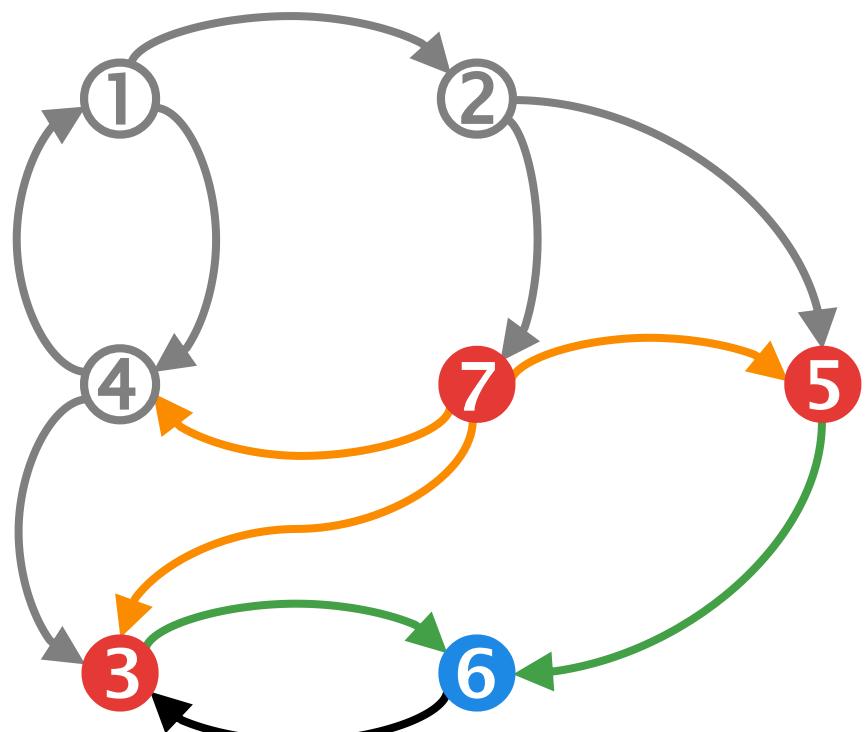
$\downarrow f(f) = \text{id}$



# BFS – PARENTS

semiring	domain	$\oplus$	$\otimes$	0
min-first	$\mathbb{N}$	min	first	0

$$\text{first}(x, y) = x$$



f

1	2	3	4	5	6	7
		3		5		7

p

0	1	4	1	2	3	2
					3	2

A

1	2	3	4	5	6	7
1		●		●		
2					●	●
3						●
4	●		●			
5						●
6			●			
7			●	●	●	

$f(\neg p)$

✗	✗	✗	✗	✗	✗	3	✗
---	---	---	---	---	---	---	---

$$f(\neg p) = f \text{ min . first } A$$

$\swarrow$   $p(f) = f$

$\downarrow$   $f(f) = \text{id}$

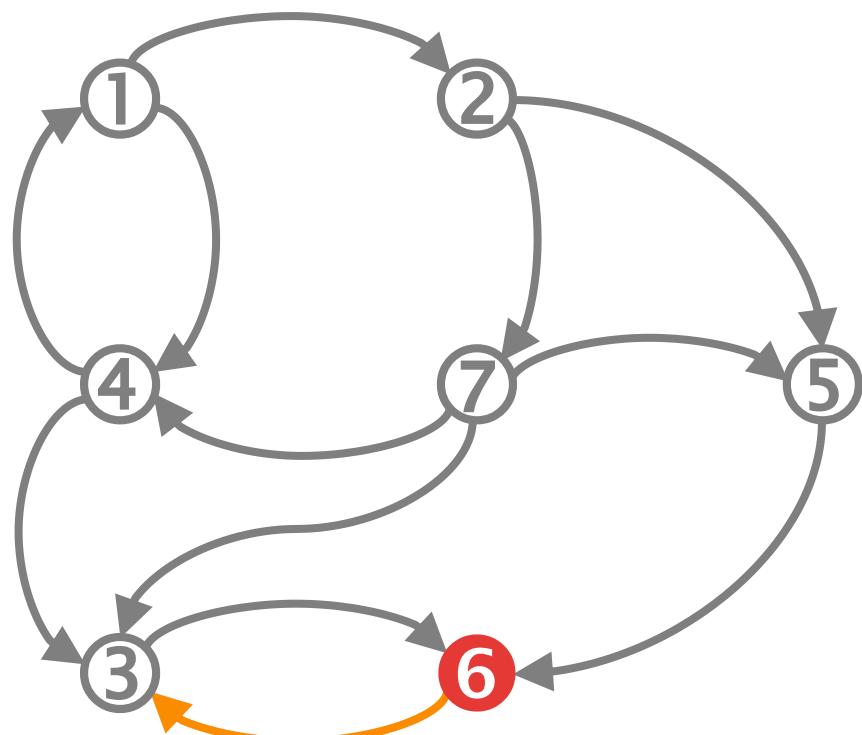
$f(f)$

1	2	3	4	5	6	7
					6	

# BFS – PARENTS

semiring	domain	$\oplus$	$\otimes$	0
min-first	$\mathbb{N}$	min	first	0

$$\text{first}(x, y) = x$$



f

1	2	3	4	5	6	7
					6	

p

0	1	4	1	2	3	2
---	---	---	---	---	---	---

A

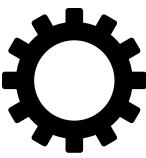
1	2	3	4	5	6	7
1						
2						
3						
4						
5						
6						
7						

X	X	X	X	X	X	X
---	---	---	---	---	---	---

$f \langle \neg p \rangle = f \text{ min . first } A$

**f** is empty  
 $\rightarrow$  terminate



# BFS – PARENTS: ALGORITHM

- **Input:** adjacency matrix  $A$ , source vertex  $s$ , #vertices  $n$
- **Output:** parent vertices vector  $\mathbf{p}$  (integer)
- **Workspace:** vertex index vector  $\mathbf{idx}$  (integer), frontier vector  $\mathbf{f}$  (integer)

1.  $\mathbf{idx} = [1 \ 2 \ \dots \ n]$  we assume 1-based indexing here
2.  $\mathbf{f}(s) = s$
3.  $\mathbf{p}(s) = 0$
4. for  $l = 1$  to  $n - 1$  \*terminate earlier if the frontier is empty
5.    $\mathbf{f}(\neg\mathbf{p}) = \mathbf{f} \text{ min .first } A$  advance the frontier
6.    $\mathbf{p}(\mathbf{f}) = \mathbf{f}$  assign parent ids to the frontier's vertices
7.    $\mathbf{f}(\mathbf{f}) = \mathbf{idx}$  assign vertex ids  $\mathbf{f}(i) = i$

# BFS – PARENTS: OPTIMIZATIONS

- If getting deterministic results is not a requirement (i.e. any parent vertex can be returned), instead of min. sel1st, one can use the any. first semiring.
- This optimization is allowed by the GAP Benchmark Suite.
- Direction-optimizing traversal (push/pull) can be exploited.
- The secondi (note the “i”) semiring can be used to express the BFS Parent algorithm. When using this semiring, one does not even have to look at the values in either  $A$  or  $f$ .

**This algorithm is described in:**

*Evaluation of Graph Analytics Frameworks Using the GAP Benchmark Suite*, IISWC 2020

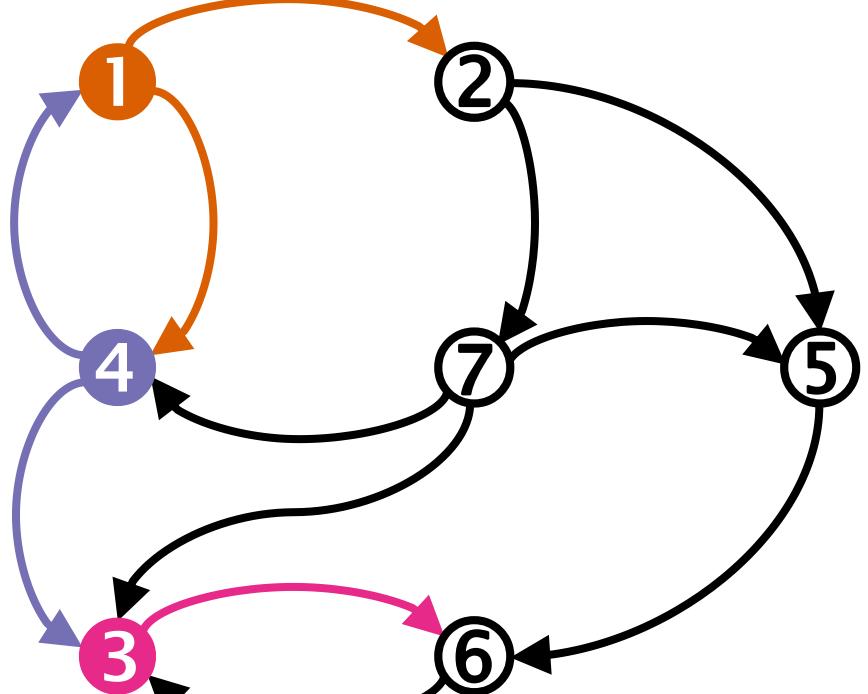


# Graph algorithms in GraphBLAS

Multi-source BFS

# MULTI-SOURCE BFS - LEVELS

semiring	domain	$\oplus$	$\otimes$	0
any-pair	{T, F}	any	pair	F



F	1	2	3	4	5	6	7
t1	●						
t2			●				
t3				●			

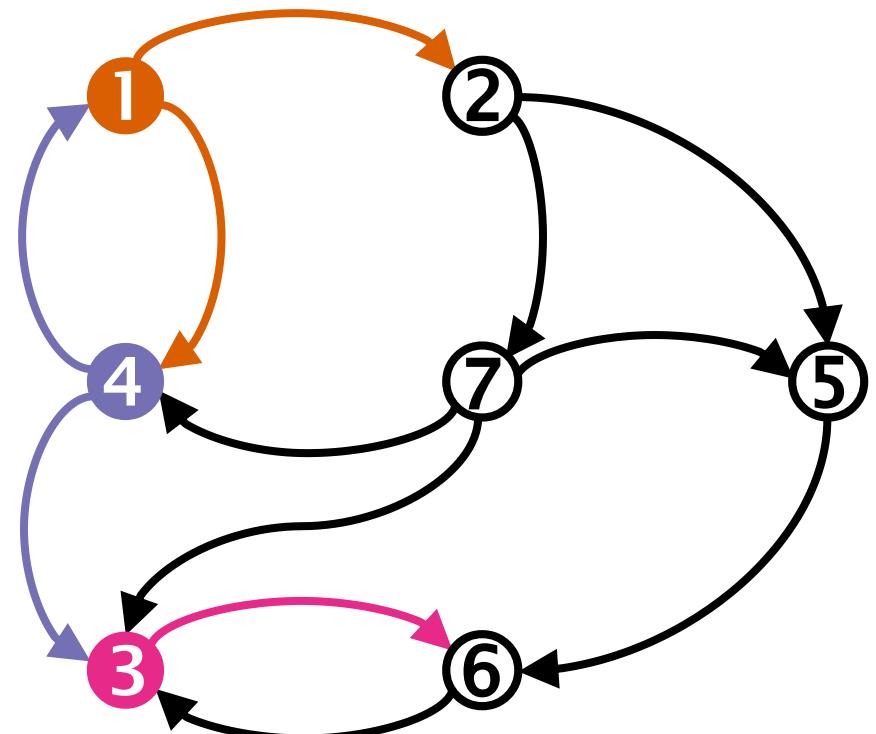
A	1	2	3	4	5	6	7
1		●			●		
2						●	●
3						●	
4	●			●			
5							●
6					●	●	
7						●	

		●	●				

$F \langle \neg S \rangle = F \text{ any. pair } A$

# MULTI-SOURCE BFS – PARENTS

semiring	domain	$\oplus$	$\otimes$	0
min-first	$\mathbb{N} \cup \{+\infty\}$	min	first	$+\infty$



F

	1	2	3	4	5	6	7
t1	1						
t2			3				
t3				4			

A

	1	2	3	4	5	6	7
1		●			●		
2						●	●
3						●	
4	●		●				
5							●
6			●				
7			●	●	●		

	1	1					
			X	X	X	X	
	4	4					

$$F \langle \neg P \rangle = F \text{ min . first } A$$

# BFS – PERFORMANCE

- Naïve BFS impls can be slow on real graphs with skewed distributions – further optimizations are needed.
- Direction-optimizing BFS was published in 2012.
  - Switches between push ( $vA$ ) and pull ( $A^T v$ ) during execution:
    - Use the push direction when the frontier is small
    - Use the pull direction when the frontier becomes large
  - Adopted to GraphBLAS in 2018 (Yang et al.'s ICPP paper)



S. Beamer, K. Asanovic, D. Patterson:  
*Direction-Optimizing Breadth-First Search*, SC 2012



C. Yang: *High-performance linear algebra-based graph framework on the GPU*, PhD thesis, UC Davis, 2019



C. Yang, A. Buluç, J.D. Owens: *Implementing Push-Pull Efficiently in GraphBLAS*, ICPP 2018



A. Buluç: *GraphBLAS: Concepts, algorithms, and applications*, Scheduling Workshop, 2019

# Graph algorithms in GraphBLAS

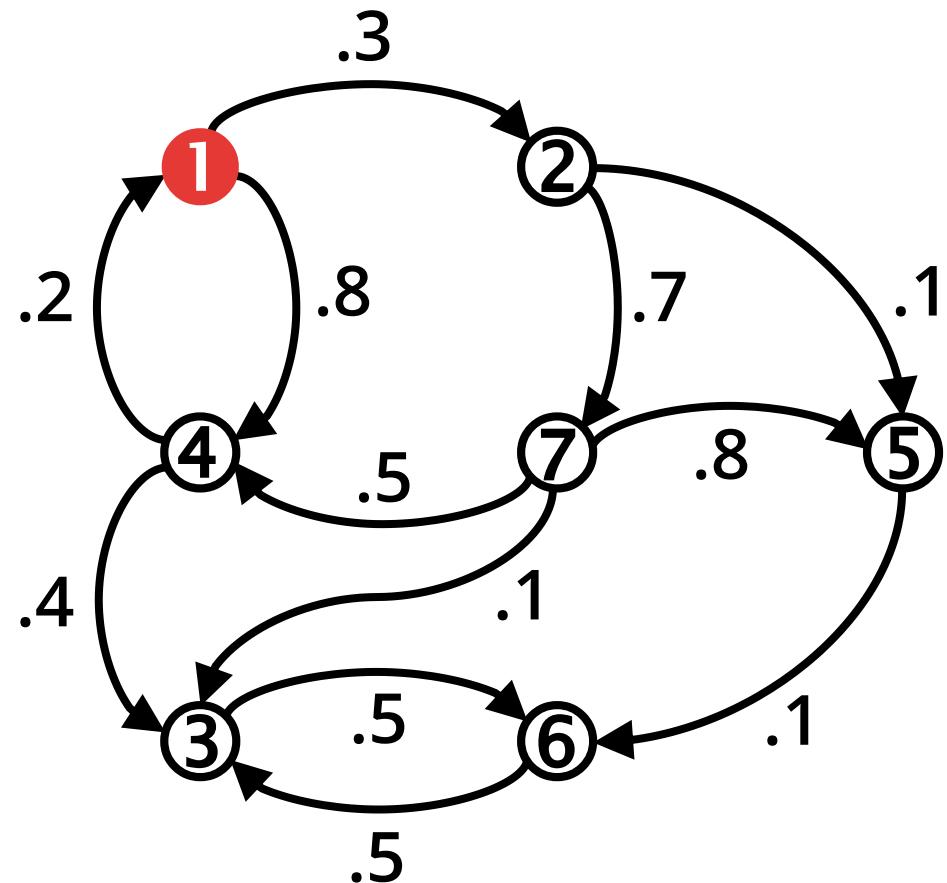
Single-source shortest paths

# SSSP – SINGLE-SOURCE SHORTEST PATHS

- Problem:
  - From a given start vertex  $s$ , find the shortest paths to every other (reachable) vertex in the graph
- Bellman-Ford algorithm:
  - Relaxes all edges in each step
  - Guaranteed to find the shortest path using at most  $n - 1$  steps
- Observation:
  - The relaxation step can be captured using a VM multiplication
  - Unlike in BFS, there is no masking here, as revisiting edges that have been visited previously can be useful.

# SSSP – ALGEBRAIC BELLMAN-FORD

semiring	domain	$\oplus$	$\otimes$	0
min-plus	$\mathbb{R} \cup \{+\infty\}$	min	+	$+\infty$



A	①	②	③	④	⑤	⑥	⑦
①	0	.3		.8			
②		0			.1		.7
③			0				.5
④	.2		.4	0			
⑤					0	.1	
⑥			.5			0	
⑦				.1	.5	.9	0

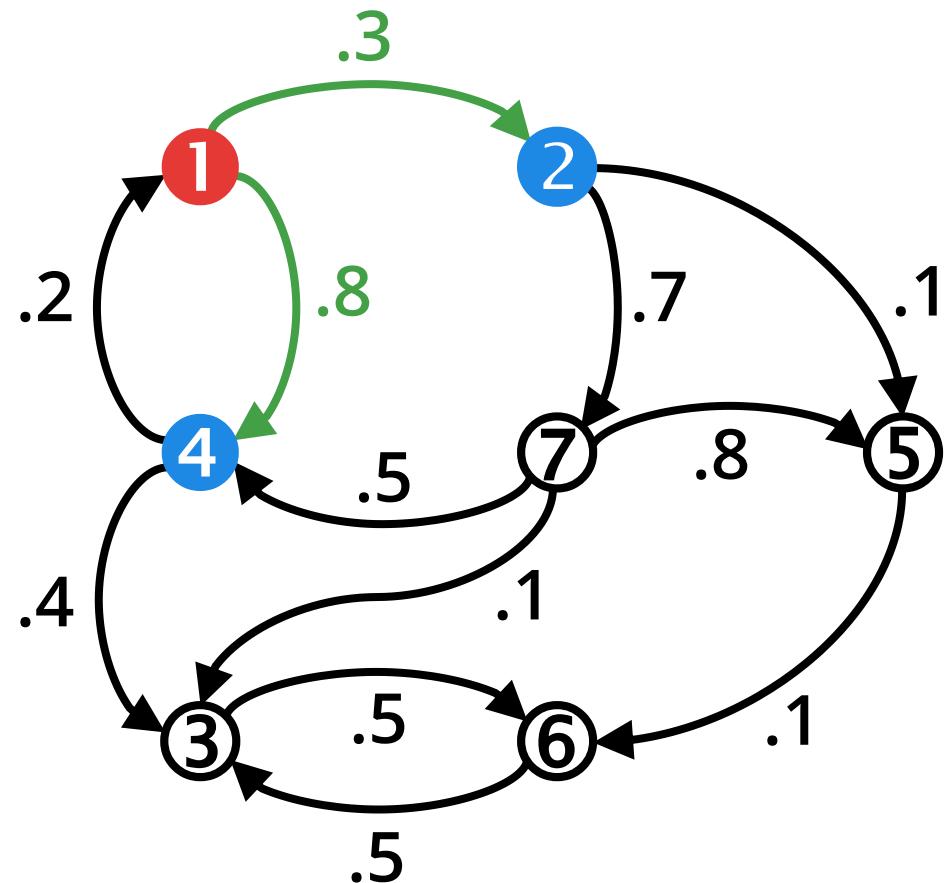
d	①	②	③	④	⑤	⑥	⑦
①	0						
②							

d min.+ A

# SSSP – ALGEBRAIC BELLMAN-FORD

semiring	domain	$\oplus$	$\otimes$	0
min-plus	$\mathbb{R} \cup \{+\infty\}$	min	+	$+\infty$



A	①	②	③	④	⑤	⑥	⑦
1	0	.3		.8			
2		0			.1		.7
3			0			.5	
4	.2		.4	0			
5					0	.1	
6			.5			0	
7			.1	.5	.9		0

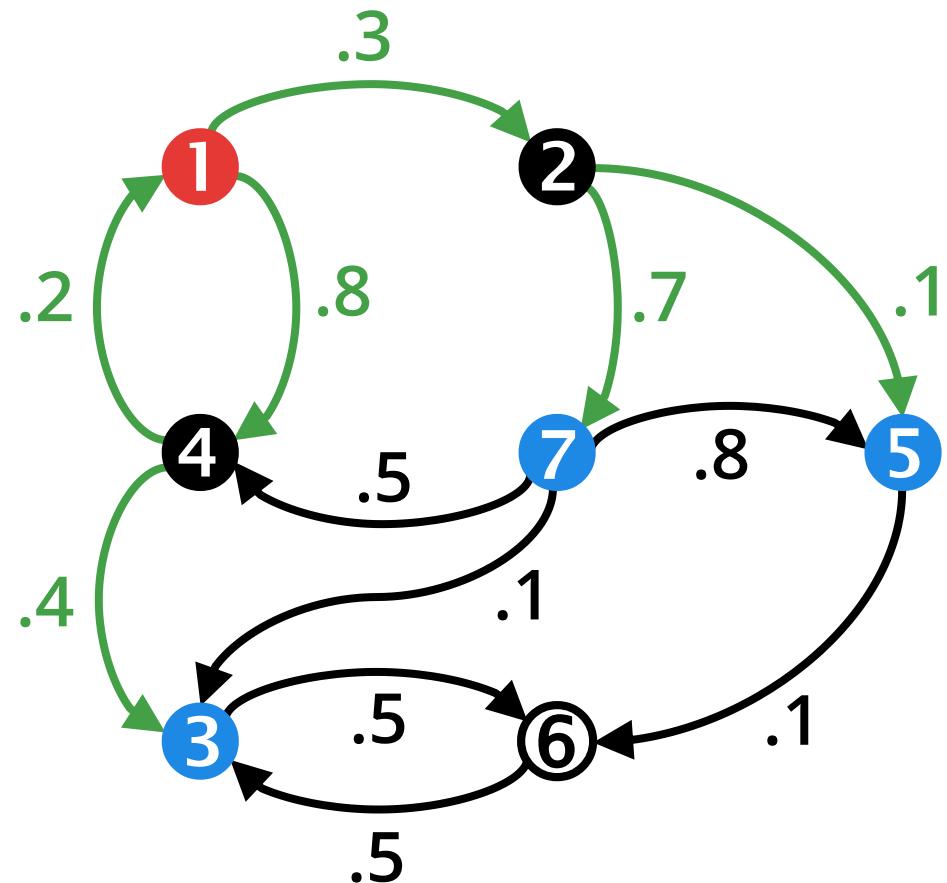
d	①	②	③	④	⑤	⑥	⑦
0							

d min.+ A
0 .3 .8

# SSSP – ALGEBRAIC BELLMAN-FORD

semiring	domain	$\oplus$	$\otimes$	0
min-plus	$\mathbb{R} \cup \{+\infty\}$	min	+	$+\infty$



A	①	②	③	④	⑤	⑥	⑦
①	0	.3		.8			
②		0			.1		.7
③			0				.5
④	.2		.4	0			
⑤					0	.1	
⑥			.5			0	
⑦			.1	.5	.9		0

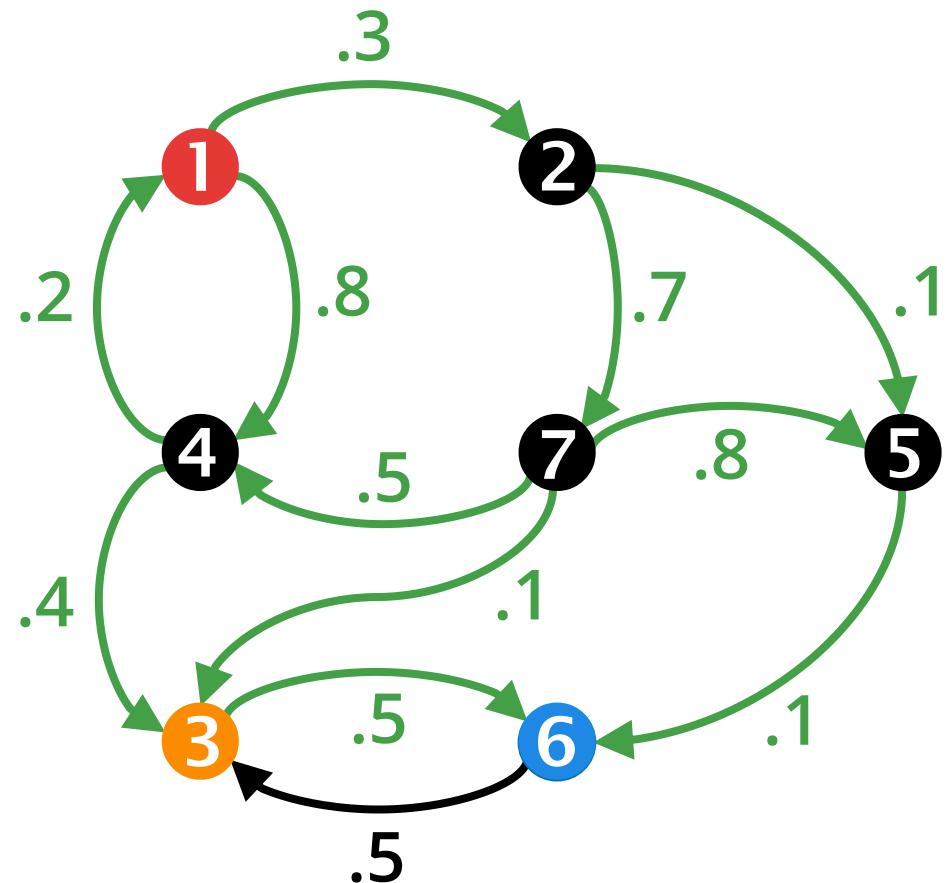
d	①	②	③	④	⑤	⑥	⑦
0	0	.3		.8			
0	.3	1.2	.8	.4			1

d min.+ A
-----------

# SSSP – ALGEBRAIC BELLMAN-FORD

semiring	domain	$\oplus$	$\otimes$	0
min-plus	$\mathbb{R} \cup \{+\infty\}$	min	+	$+\infty$



A	①	②	③	④	⑤	⑥	⑦
1	0	.3		.8			
2		0			.1		.7
3			0				.5
4	.2		.4	0			
5					0	.1	
6			.5			0	
7			.1	.5	.9		0

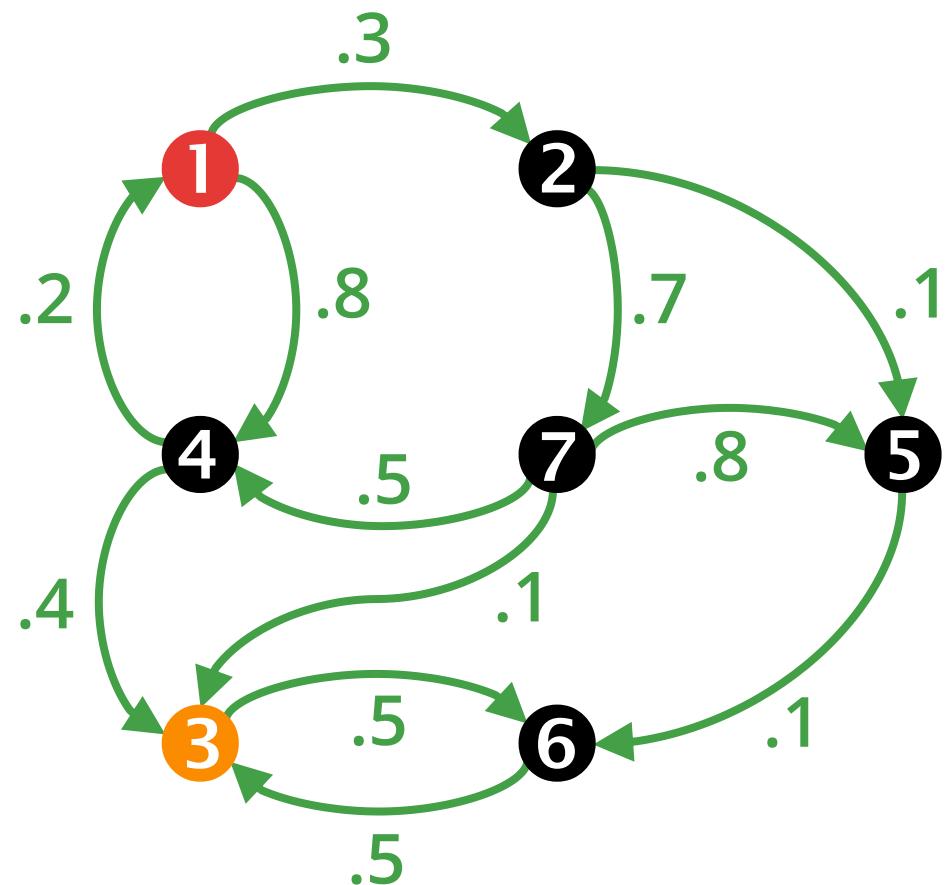
d	①	②	③	④	⑤	⑥	⑦
0	.0	.3	1.2	.8	.4		1

d min.+ A
0 .3 1.1 .8 .4 .5 1

# SSSP – ALGEBRAIC BELLMAN-FORD

semiring	domain	$\oplus$	$\otimes$	0
min-plus	$\mathbb{R} \cup \{+\infty\}$	min	+	$+\infty$



A	①	②	③	④	⑤	⑥	⑦
1	0	.3		.8			
2		0			.1		.7
3			0			.5	
4	.2		.4	0			
5					0	.1	
6		.5				0	
7			.1	.5	.9		0

d	①	②	③	④	⑤	⑥	⑦
0	.0	.3	1.1	.8	.4	.5	1

d min.+ A
0 .3 1 .8 .4 .5 1



# SSSP – ALGEBRAIC BELLMAN-FORD ALGO.

Input: adjacency matrix  $\mathbf{A}$ , source vertex  $s$ , #vertices  $n$

$$\mathbf{A}_{ij} = \begin{cases} 0 & \text{if } i = j \\ w(e_{ij}) & \text{if } e_{ij} \in E \\ \infty & \text{if } e_{ij} \notin E \end{cases}$$

Output: distance vector  $\mathbf{d}$  (real)

1.  $\mathbf{d} = [\infty \ \infty \dots \infty]$
2.  $\mathbf{d}(s) = 0$
3. for  $k = 1$  to  $n - 1$  \*terminate earlier if we reach a fixed point
4.  $\mathbf{d} = \mathbf{d} \min.+ \mathbf{A}$

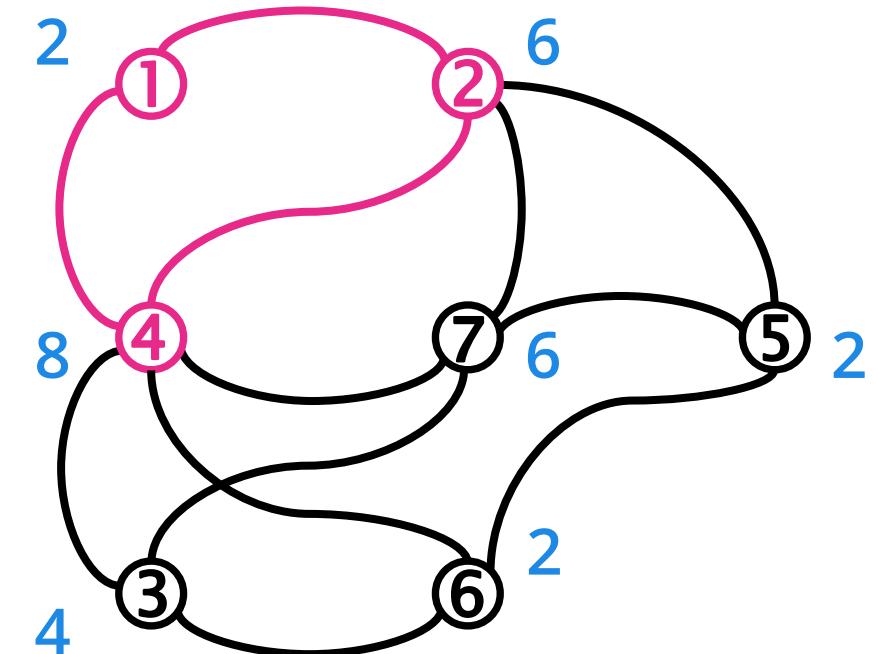
Optimizations for BFS (push/pull) also work here.

# Graph algorithms in GraphBLAS

Triangle count / Definition

# TRIANGLE COUNT

- IEEE GraphChallenge: an annual competition at the HPEC conference
- The task of the 2017 GraphChallenge was **triangle count**: given a graph  $G$ , count the number of triangles.
- Triangle = “*set of three mutually adjacent vertices in a graph*”
- Many solutions employed a linear algebraic computation model



Number of unique triangles:

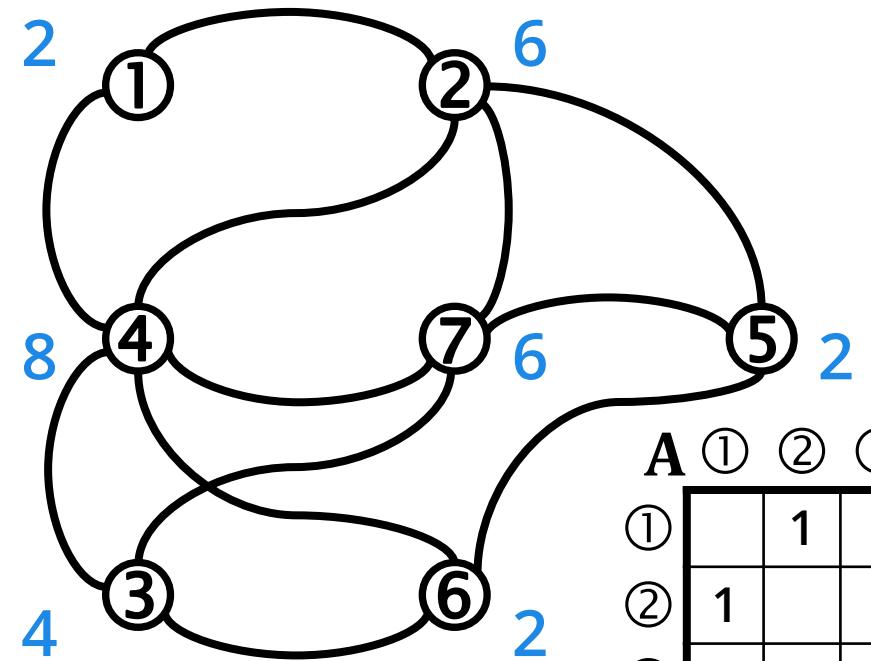
$$\frac{30}{6}$$

# Graph algorithms in GraphBLAS

Triangle count / Naïve algorithm

# TC EXAMPLE: NAÏVE APPROACH

$$\text{tri2} = \text{diag}^{-1}(A \oplus \cdot \otimes A \oplus \cdot \otimes A)$$



A	①	②	③	④	⑤	⑥	⑦
①	1	1	1				
②	1			1	1	1	
③				1	1	1	
④	1	1	1			1	1
⑤		1			1	1	
⑥			1	1	1		
⑦		1	1	1	1		

A	①	②	③	④	⑤	⑥	⑦	A	①	②	③	④	⑤	⑥	⑦
①	1	1	1					①	1	1	1				
②	1			1	1	1		②	1			1	1	1	
③				1	1	1		③	1			1	1	1	
④	1	1	1			1	1	④	1	1	1		1	1	1
⑤		1			1	1		⑤	1			1	1	1	
⑥			1	1	1			⑥	1	1	1	1			
⑦		1	1	1	1		⑦	1	1	1	1				

tri2	2	6	4	7	4	3	4	2	6	6	6	11	8	5	9	2
	6	6	4	8	4	7	9	4	6	4	8	4	7	9	8	4
	4	6	4	5	2	8	9	4	8	4	5	2	8	9	2	2

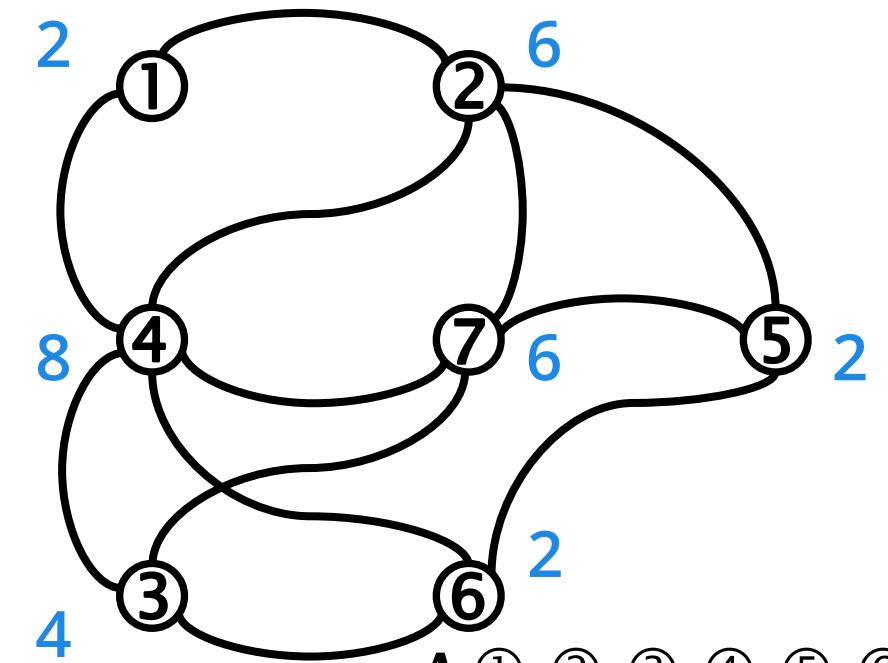
  

tri2	2	6	4	7	4	3	4	2	6	6	6	11	8	5	9	2
	6	6	4	8	4	7	9	4	6	4	8	4	7	9	8	4
	4	6	4	5	2	8	9	4	8	4	5	2	8	9	2	2

# Graph algorithms in GraphBLAS

Triangle count / Masked algorithm

# TC EXAMPLE: ELEMENT-WISE MULTIPLICATION



	①	②	③	④	⑤	⑥	⑦
①	1	1	1				
②	1		1	1			1
③			1		1	1	
④	1	1	1			1	1
⑤	1				1	1	
⑥		1	1	1			
⑦	1	1	1	1			

A	①	②	③	④	⑤	⑥	⑦
①	1		1				
②	1			1	1		1
③			1		1	1	
④	1	1	1			1	1
⑤	1				1	1	
⑥		1	1	1			
⑦	1	1	1	1			

2	1	1	1	1	1	1	2
1	4	2	2	1	2	2	
1	2	3	2	2	1	1	
1	2	2	5	3	1	2	
1	1	2	3	3		1	
1	2	1	1		3	3	
2	2	1	2	1	3	4	

$$\text{TRI} = A \oplus . \otimes A \otimes A$$

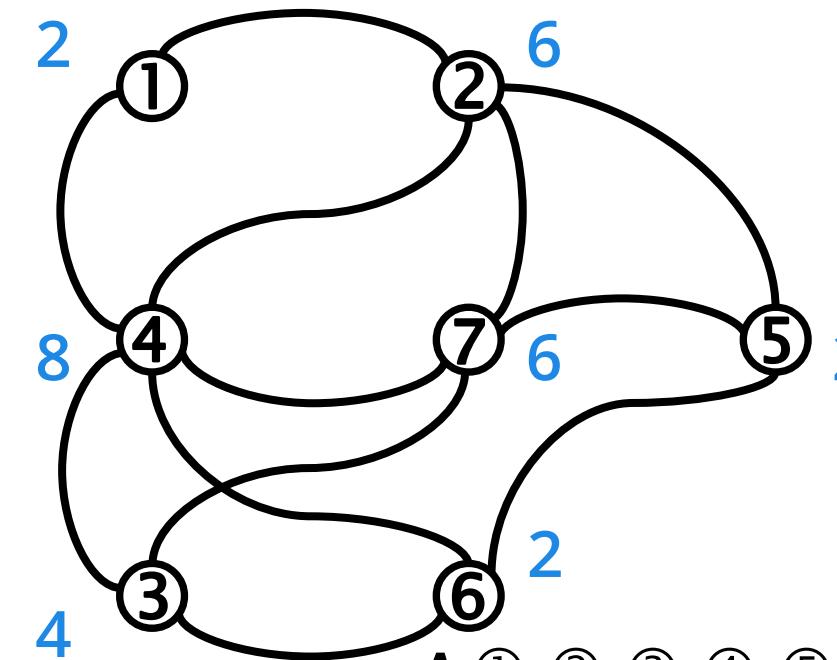
$$\text{tri2} = [\oplus_j \text{TRI}(:, j)]$$

	1	1					
1			2	1		2	
		2		1	1		
1	2	2			1	2	
	1					1	
		1					
2	1	1	1				
2	1	2	1	2	1		

$[\oplus_j \dots]$

2
6
4
8
2
2
6

# TC EXAMPLE: ELEMENT-WISE MULTIPLICATION



A	①	②	③	④	⑤	⑥	⑦
①	1						
②	1			1	1		1
③			1		1	1	
④	1	1	1			1	1
⑤	1				1	1	
⑥		1	1	1			
⑦	1	1	1	1			

A	①	②	③	④	⑤	⑥	⑦
①	1						
②	1			1	1		1
③			1		1	1	
④	1	1	1			1	1
⑤	1				1	1	
⑥		1	1	1			
⑦	1	1	1	1			

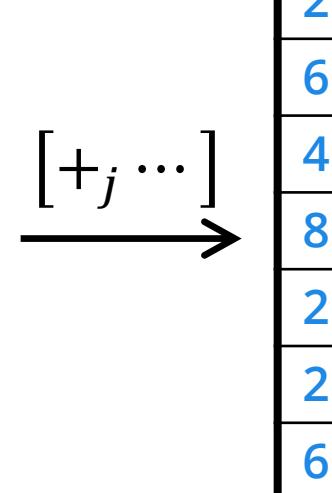
	1	1	2	1	2
1	1		2	1	2
2			2	1	1
3	1			1	2
4		2	2		
5	1				1
6		1	1		
7	2	1	2	1	

*Masking* limits where the operation is computed. Here, we use  $\mathbf{A}$  as a mask for  $\mathbf{A} \oplus.\otimes \mathbf{A}$ .

$$\text{TRI}(\mathbf{A}) = \mathbf{A} \oplus.\otimes \mathbf{A}$$

$$\text{tri2} = [\oplus_j \text{TRI}(:,j)]$$

**tri2**



# Graph algorithms in GraphBLAS

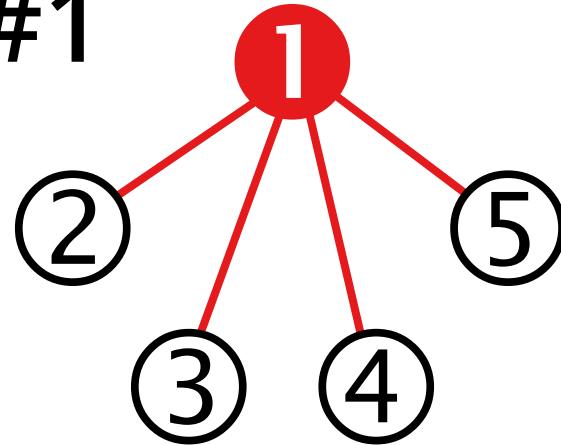
The importance of masking

# THE IMPORTANCE OF MASKING

**Q:** Is masking absolutely necessary?

**A:** Yes, it can reduce the complexity of some algorithms.  
We demonstrate this with two examples.

#1



A	1	2	3	4	5
1		1	1	1	1
2	1				
3	1				
4	1				
5	1				

A ① ② ③ ④ ⑤

1	1	1	1	1
2	1			
3	1			
4	1			
5	1			

A

4				
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

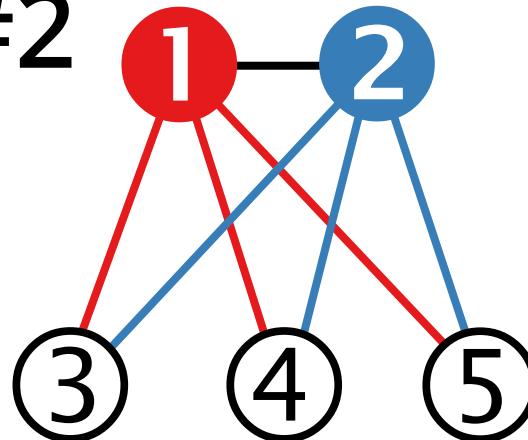
A  $\oplus \otimes$  A


A  $\oplus \otimes$  A  $\oplus \otimes$  A

A simple corner case is the star graph: there are  $(n - 1)^2$  wedges but none of them close into triangles.

We do quadratic work while it's clear that there are no triangles in the graph (it's a tree).

#2



A	1	2	3	4	5
1	1	1	1	1	1
2	1		1	1	1
3	1	1			
4	1	1			
5	1	1			

A 1 2 3 4 5

A	1	2	3	4	5
1	1	1	1	1	1
2	1		1	1	1
3	1	1			
4	1	1			
5	1	1			

1	1	1	1	1
2	1		1	1
3	1	1		
4	1	1		
5	1	1		

A

4	3	1	1	1
3	4	1	1	1
1	1	2	2	2
1	1	2	2	2
1	1	2	2	2

 $A \oplus \otimes A$ 

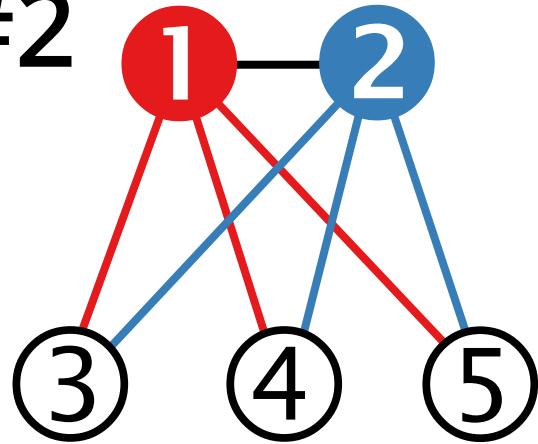
	3	1	1	1
3		1	1	1
1	1			
1	1			
1	1			

 $A \oplus \otimes A \oplus \otimes A$ 

A full bipartite graph  $K_{2,3}$  with the vertices in the top partition connected.

A bipartite graph only has cycles of even length, so it's easy to see that all triangles will contain the two vertices in the top partition. Still,  $A \oplus \otimes A$  enumerates all wedges starting and ending in the bottom partition, thus performing a lot of unnecessary work.

#2



A

	1	2	3	4	5
1	1	1	1	1	1
2	1		1	1	1
3	1	1			
4	1	1			
5	1	1			

A ① ② ③ ④ ⑤

①	1	1	1	1	1
②	1		1	1	1
③	1	1			
④	1	1			
⑤	1	1			

A

$$\text{TRI}\langle A \rangle = A \oplus . \otimes A$$

	3	1	1	1
3		1	1	1
1	1			
1	1			
1	1			

*Masking* avoids the materialization of large interim data sets by ensuring that we only enumerate wedges whose endpoints are already connected.

$$[\oplus_j \text{TRI}(:,j)]$$

6
6
2
2
2

tri

# Graph algorithms in GraphBLAS

Triangle count / Cohen's algorithm



# COHEN'S ALGORITHM: PSEUDOCODE

**Input:** adjacency matrix  $A$

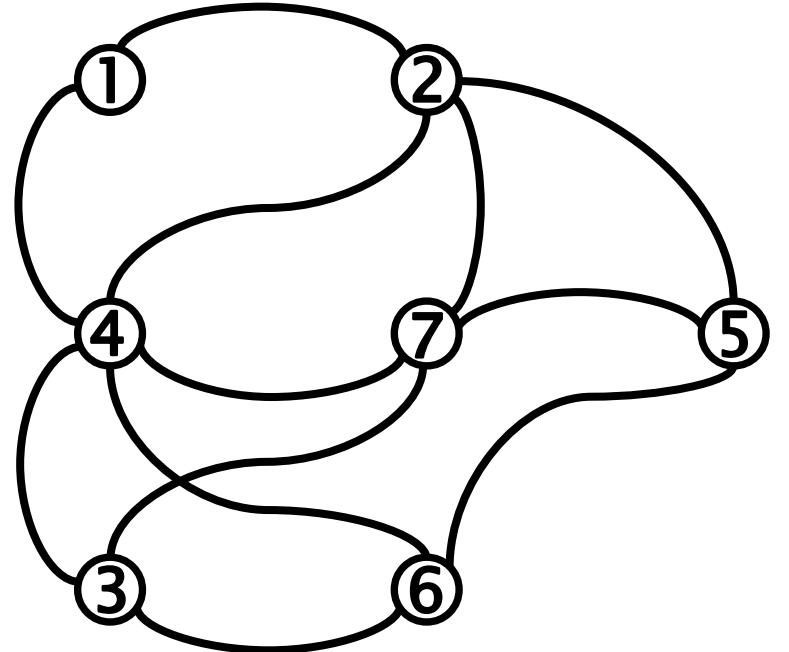
**Output:** triangle count  $t$

**Workspace:** matrices  $L, U, B, C$

1.  $L = \text{tril}(A)$  extract the lower triangle from  $A$
2.  $U = \text{triu}(A)$  extract the upper triangle from  $A$
3.  $B = L \oplus . \otimes U$
4.  $C = B \otimes A$
5.  $t = \sum C / 2$  sum the values in  $C$  and divide by 2



# COHEN'S ALGORITHM



L ① ② ③ ④ ⑤ ⑥ ⑦

①						
②	1					
③						
④	1	1	1			
⑤	1					
⑥		1	1	1	1	
⑦	1	1	1	1		

B

	1				
	1	3	1	1	2
	1	1			
	1		3	3	
	2	1	3	4	

U ① ② ③ ④ ⑤ ⑥ ⑦

①		1		1		
②			1	1	1	
③			1		1	1
④				1	1	
⑤					1	1
⑥						
⑦						

A ① ② ③ ④ ⑤ ⑥ ⑦

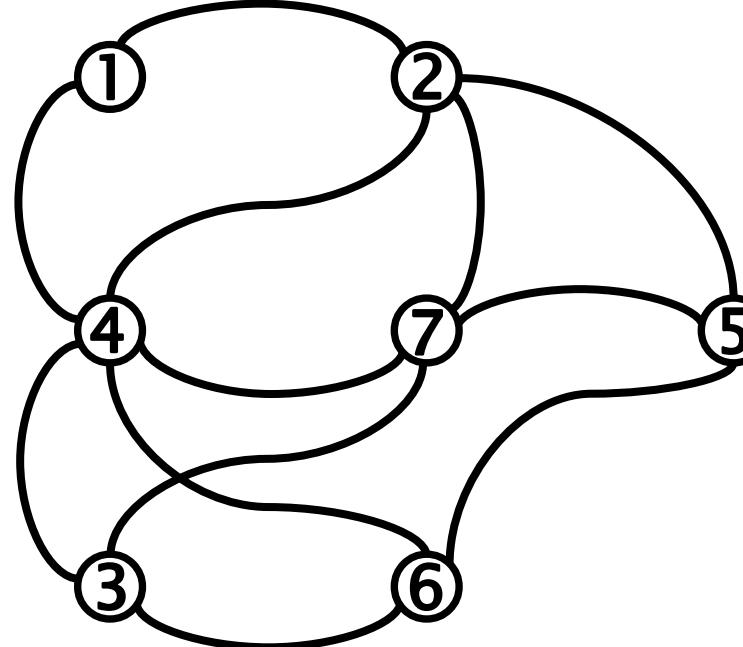
1		1		1		
1			1	1	1	
1			1		1	1
1	1	1	1		1	1
1			1	1	1	
1			1	1	1	1
1	1	1	1	1		

C

		1		
	1		1	2
	1			1
	1			
	2	1		

$$t = \sum \mathbf{C} / 2$$

# COHEN'S ALGORITHM: MASKING



**L** ① ② ③ ④ ⑤ ⑥ ⑦

①						
②	1					
③						
④	1	1	1			
⑤		1				
⑥		1	1	1	1	
⑦	1	1	1	1		

**U** ① ② ③ ④ ⑤ ⑥ ⑦

①		1		1			
②				1	1	1	
③				1	1	1	
④					1	1	
⑤						1	1
⑥							
⑦							

**C**


$$C\langle A \rangle = L \oplus . \otimes U$$

$$t = \sum C / 2$$

# Graph algorithms in GraphBLAS

Triangle count / Sandia algorithm



# SANDIA ALGORITHM

**Input:** adjacency matrix  $\mathbf{A}$

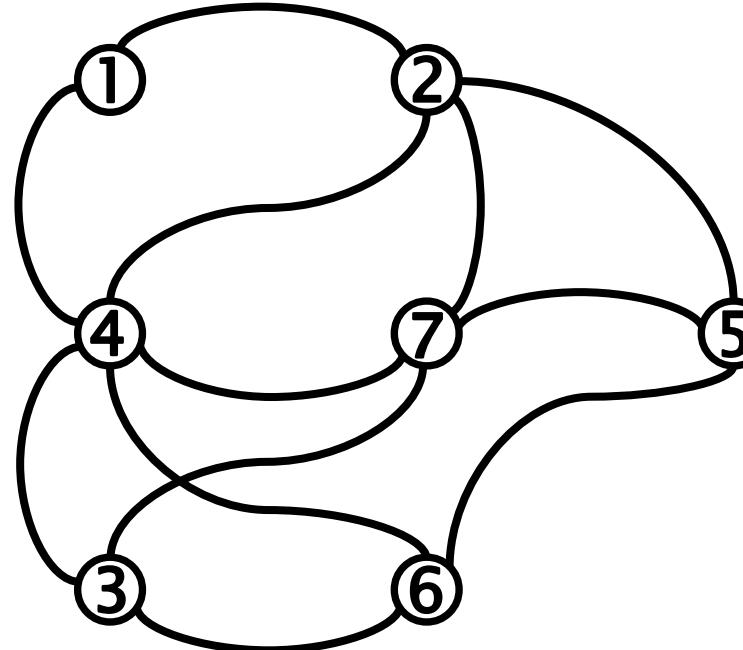
**Output:** triangle count  $t$

**Workspace:** matrices  $\mathbf{L}$ ,  $\mathbf{U}$ ,  $\mathbf{B}$ ,  $\mathbf{C}$

1.  $\mathbf{L} = \text{tril}(\mathbf{A})$  extract the lower triangle from  $\mathbf{A}$
2.  $\mathbf{C}(\mathbf{L}) = \mathbf{L} \oplus_{\cdot} \otimes \mathbf{L}$  multiply matrices  $\mathbf{L}$  and  $\mathbf{L}$  using mask  $\mathbf{L}$
3.  $t = \sum \mathbf{C}$  sum the values in  $\mathbf{C}$



# SANDIA ALGORITHM



L	①	②	③	④	⑤	⑥	⑦
①							
②	1						
③							
④	1	1	1				
⑤		1					
⑥			1	1	1	1	
⑦		1	1	1	1		

L	①	②	③	④	⑤	⑥	⑦
①							
②	1						
③							
④	1	1	1				
⑤		1					
⑥			1	1	1	1	
⑦		1	1	1	1		

C							
①							
②							
③							
④							
⑤							
⑥							
⑦							

C							
①							
②							
③							
④							
⑤							
⑥							
⑦							

$$C\langle L \rangle = L \oplus . \otimes L$$

$$t = \sum C$$

# Graph algorithms in GraphBLAS

Triangle count / CMU algorithm

# CMU ALGORITHM

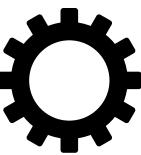
- Iterates on the vertices of the graph, extracts corresponding submatrices and computes  $t = t + a_{10}^\top \oplus \cdot \otimes A_{20} \oplus \cdot \otimes a_{12}$
- Tradeoffs:
  - does not require mxm, only vxm and mxv
  - slower than mxm-based algorithms
- The formula is derived using the matrix trace  $\text{tr}(\mathbf{A}) = \sum_{i=0}^{n-1} \mathbf{A}_{ii}$  and its invariant property under cyclic permutation, e.g.  $\text{tr}(\mathbf{ABC}) = \text{tr}(\mathbf{BCA}) = \text{tr}(\mathbf{CAB})$ . See the paper for details.

$A$	$i$	
$A_{00}$	$a_{01}$	$A_{20}^\top$
$a_{01}^\top$	0	$a_{21}^\top$
$A_{20}$	$a_{21}$	$A_{22}$

T.M. Low et al. (Carnegie Mellon University):

*First look: linear algebra-based triangle counting without matrix multiplication*, HPEC 2017





# CMU ALGORITHM: PSEUDOCODE

**Input:** adjacency matrix  $\mathbf{A}$

**Output:** triangle count  $t$

**Workspace:** matrices  $\mathbf{A}_{20}$ ,  $\mathbf{C}$ , vectors  $\mathbf{a}_{10}$ ,  $\mathbf{a}_{12}^\top$

1. **for**  $i = 2$  **to**  $n - 1$
2.    $\mathbf{A}_{20} = \mathbf{A}[i + 1:n, 0:i - 1]$
3.    $\mathbf{a}_{10} = \mathbf{A}[0:i - 1, i]$
4.    $\mathbf{a}_{12} = \mathbf{A}[i, i + 1:n]$
5.    $t = t + \mathbf{a}_{10}^\top \oplus \otimes \mathbf{A}_{20} \oplus \otimes \mathbf{a}_{12}$

$\mathbf{A}$	$i$	
$\mathbf{A}_{00}$	$\mathbf{a}_{01}$	$\mathbf{A}_{20}^\top$
$\mathbf{a}_{01}^\top$	0	$\mathbf{a}_{21}^\top$
$\mathbf{A}_{20}$	$\mathbf{a}_{21}$	$\mathbf{A}_{22}$

T.M. Low et al. (Carnegie Mellon University):

*First look: linear algebra-based triangle counting without matrix multiplication*, HPEC 2017



# PROVABLY CORRECT ALGORITHMS

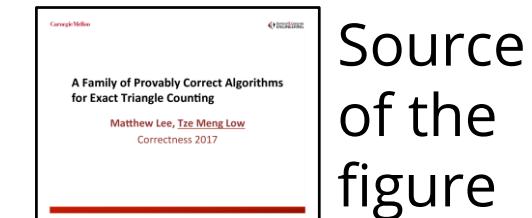
The “CMU algorithm” belongs to a family of algorithms which can be derived using the “FLAME approach”.

There are 8 similar algorithms in total, the one presented here is Algorithm 2.

Algorithm: $\tilde{\Delta} := \frac{1}{6}\Gamma(\tilde{A}^3)$	
$A \rightarrow \left( \begin{array}{c c} A_{TL} & A_{TR} \\ \hline A_{TR}^T & A_{BR} \end{array} \right)$	where $A_{TL}$ is a $0 \times 0$ matrix
while $m(A_{TL}) < m(A)$ do	
$\left( \begin{array}{c c} A_{TL} & A_{TR} \\ \hline A_{TR}^T & A_{BR} \end{array} \right) \rightarrow \left( \begin{array}{c c c} A_{00} & a_{01} & A_{02} \\ \hline a_{01}^T & \alpha_{11} & a_{12}^T \\ \hline A_{02}^T & a_{12} & A_{22} \end{array} \right)$	where $\alpha_{11}$ is a $1 \times 1$ matrix
Algorithm 1	Algorithm 2
$\Delta := \Delta + \frac{1}{2}a_{01}^T A_{00} a_{01}$	$\Delta := \Delta + a_{01}^T A_{02} a_{21}$
Algorithm 3	Algorithm 4
$\Delta := \Delta + \frac{1}{2}a_{01}^T A_{00} a_{01}$	$\Delta := \Delta + \frac{1}{2}a_{12}^T A_{22} a_{12}$
$\Delta := \Delta + \frac{1}{2}a_{12}^T A_{22} a_{12}$	
$\Delta := \Delta - a_{01}^T A_{02} a_{21}$	
$\left( \begin{array}{c c} A_{TL} & A_{TR} \\ \hline A_{TR}^T & A_{BR} \end{array} \right) \leftarrow \left( \begin{array}{c c c} A_{00} & a_{01} & A_{02} \\ \hline a_{01}^T & \alpha_{11} & a_{12}^T \\ \hline A_{02}^T & a_{12} & A_{22} \end{array} \right)$	
endwhile	

Algorithm: $t := \frac{1}{6}\Gamma(A^3)$	
$A \rightarrow \left( \begin{array}{c c} A_{TL} & A_{TR} \\ \hline A_{TR}^T & A_{BR} \end{array} \right)$	where $A_{BR}$ is a $0 \times 0$ matrix
while $m(A_{TL}) < m(A)$ do	
$\left( \begin{array}{c c} A_{TL} & A_{TR} \\ \hline A_{TR}^T & A_{BR} \end{array} \right) \rightarrow \left( \begin{array}{c c c} A_{00} & a_{01} & A_{02} \\ \hline a_{01}^T & \alpha_{11} & a_{12}^T \\ \hline A_{02}^T & a_{12} & A_{22} \end{array} \right)$	where $\alpha_{11}$ is a $1 \times 1$ matrix
Algorithm 5	Algorithm 6
$\Delta := \Delta + \frac{1}{2}a_{01}^T A_{00} a_{01}$	$\Delta := \Delta + a_{01}^T A_{02} a_{21}$
Algorithm 7	Algorithm 8
$\Delta := \Delta + \frac{1}{2}a_{01}^T A_{00} a_{01}$	$\Delta := \Delta + \frac{1}{2}a_{12}^T A_{22} a_{12}$
$\Delta := \Delta + \frac{1}{2}a_{12}^T A_{22} a_{12}$	
$\Delta := \Delta - a_{01}^T A_{02} a_{21}$	
$\left( \begin{array}{c c} A_{TL} & A_{TR} \\ \hline A_{TR}^T & A_{BR} \end{array} \right) \leftarrow \left( \begin{array}{c c c} A_{00} & a_{01} & A_{02} \\ \hline a_{01}^T & \alpha_{11} & a_{21}^T \\ \hline A_{02}^T & a_{21} & A_{22} \end{array} \right)$	
endwhile	

M. Lee, T.M. Low (Carnegie Mellon University):  
*A family of provably correct algorithms for exact triangle counting*,  
CORRECTNESS @ SC 2017



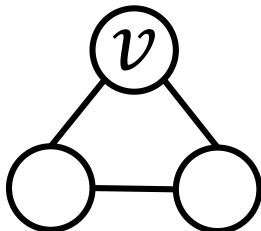
Source  
of the  
figure

# Graph algorithms in GraphBLAS

Vertex-wise triangle count

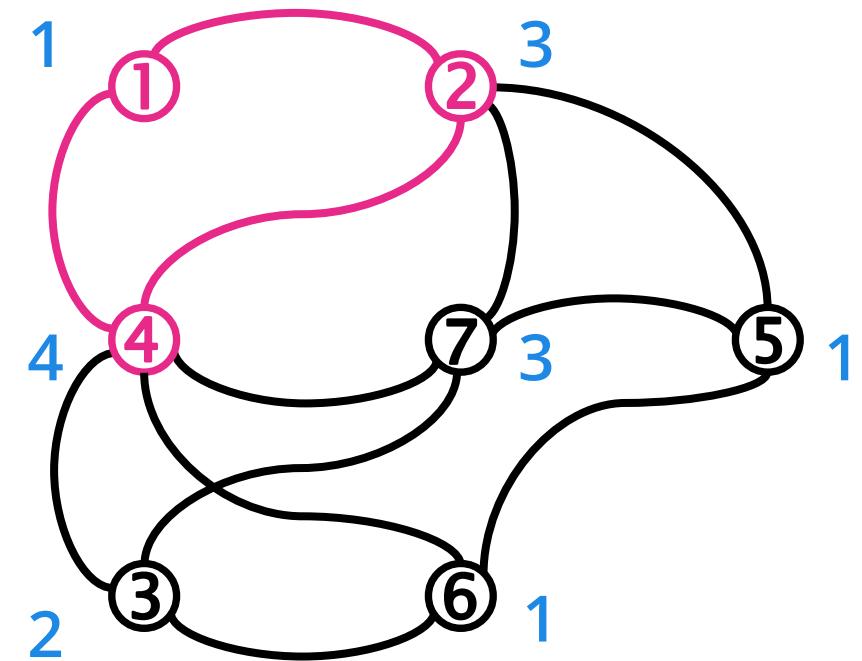
# VERTEX-WISE TRIANGLE COUNT

Triangle: a set of three mutually adjacent vertices.

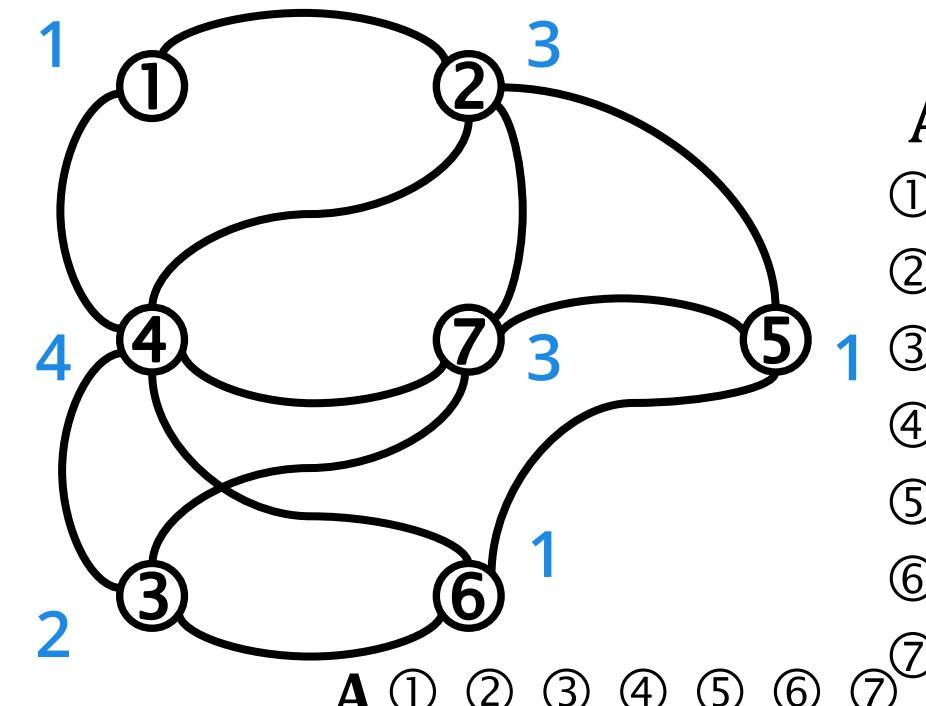


Usages:

- Global clustering coefficient
- Local clustering coefficient
- Finding communities



# TC: ELEMENT-WISE MULTIPLICATION



A	①	②	③	④	⑤	⑥	⑦
①	1		1				
②	1			1	1		1
③			1		1	1	
④	1	1	1			1	1
⑤	1				1	1	
⑥		1	1	1			
⑦	1	1	1	1			

A	①	②	③	④	⑤	⑥	⑦
①	1	1	1				
②	1		1	1	1		
③			1		1	1	
④	1	1	1			1	1
⑤	1				1	1	
⑥		1	1	1	1		
⑦	1	1	1	1			

A	①	②	③	④	⑤	⑥	⑦
①	2	1	1	1	1	1	2
②	1	4	2	2	1	2	2
③		1	2	3	2	2	1
④	1	2	2	5	3	1	2
⑤	1	1	2	3	3		1
⑥		1	2	1	1	3	3
⑦	2	2	1	2	1	3	4

$\otimes A$

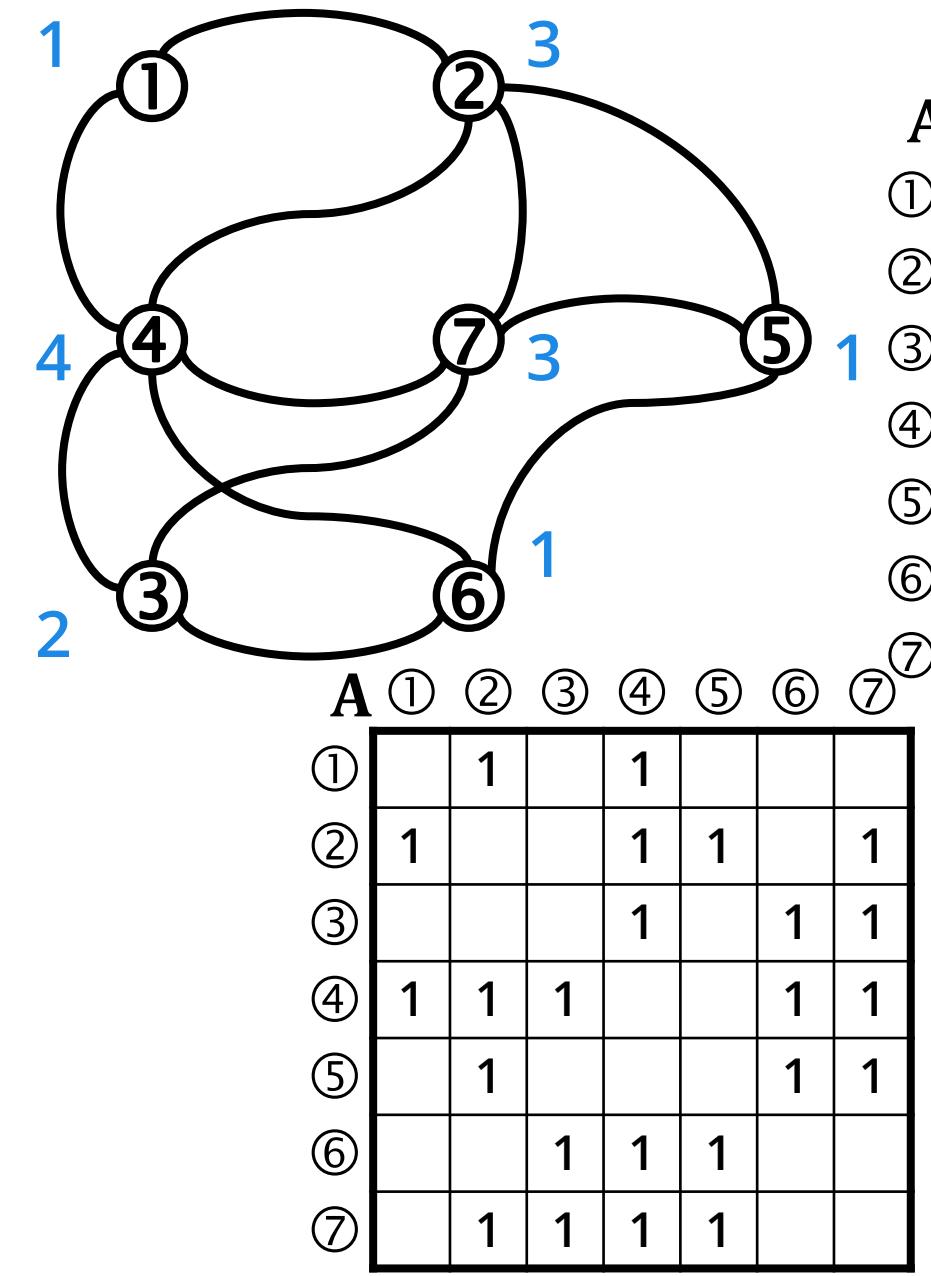
TRI	①	②	③	④	⑤	⑥	⑦
①	1	1	1				
②	1		2	1		2	
③		1	2	2	1	1	
④	1	2	2	5	3	1	2
⑤	1	1	2	3	3		1
⑥		1	2	1	1	3	3
⑦	2	2	1	2	1	3	4

$[\oplus_j \dots]$

/2

tri	①	②	③	④	⑤	⑥	⑦
①	1	3	2	4	1	1	3
②	3	2	1	1	2	1	1
③	2	1	1	1	2	1	1
④	1	1	1	1	1	1	1
⑤	1	2	1	2	1	2	1
⑥	2	1	2	1	2	1	2
⑦	1	1	1	1	1	1	1

# TC: ELEMENT-WISE MULTIPLICATION



A	1	2	3	4	5	6	7
①	1		1				
②	1			1	1		1
③				1		1	1
④	1	1	1			1	1
⑤		1				1	1
⑥			1	1	1		
⑦	1	1	1	1			

A	1	2	3	4	5	6	7
①	1		1				
②	1		1	1	1		1
③			1		1	1	
④	1	1	1			1	1
⑤		1				1	1
⑥			1	1	1		
⑦	1	1	1	1			

*Masking* limits where the operation is computed.  
Here, we use  $\mathbf{A}$  as a mask  
for  $\mathbf{A} \oplus \otimes \mathbf{A}$ .

$$\text{TRI}(\mathbf{A}) = \mathbf{A} \oplus \otimes \mathbf{A}$$

$$\text{tri} = [\oplus_j \text{TRI}(:,j)]/2$$

$\text{tri}$

$\xrightarrow{[\oplus_j \dots]}$

$/2$

1
3
2
4
1
1
3



# TC: ALGORITHM

**Input:** adjacency matrix  $\mathbf{A}$

**Output:** vector  $\mathbf{tri}$

**Workspace:** matrix  $\mathbf{TRI}$

1.  $\mathbf{TRI}(\mathbf{A}) = \mathbf{A} \odot \otimes \mathbf{A}$  compute the triangle count matrix
2.  $\mathbf{tri} = [\oplus_j \mathbf{TRI}(:, j)]/2$  compute the triangle count vector

**Optimization:** use  $\mathbf{L}$ , the lower triangular part of  $\mathbf{A}$  to avoid duplicates.

$$\mathbf{TRI}(\mathbf{A}) = \mathbf{A} \odot \otimes \mathbf{L}$$

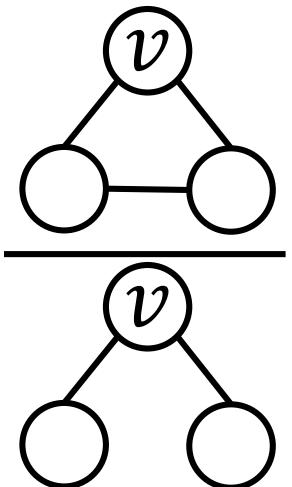
**Worst-case optimal joins:** There are deep theoretical connections between masked matrix multiplication and relational joins. It has been proven in 2013 that for the triangle query, binary joins always provide suboptimal runtime, which gave rise to new research on the family of worst-case optimal multi-way joins algorithms.

# Graph algorithms in GraphBLAS

Local clustering coefficient

# LCC: LOCAL CLUSTERING COEFFICIENT

$$LCC(v) = \frac{\text{\#edges between neighbours of } v}{\text{\#possible edges between neighbours of } v} = \frac{\text{tri}(v)}{\text{wed}(v)}$$



If  $|N(v)| \leq 1$ ,  $LCC(v) = 0$

Important metric in social network analysis.

The numerator is the number of *triangles* in  $v$ ,  $\text{tri}(v)$ .

The denominator is the number of *wedges* in  $v$ ,  $\text{wed}(v)$ .

$$LCC(v) = \frac{\text{tri}(v)}{\text{wed}(v)}$$

The difficult part is  $\text{tri}(v)$ .

# LCC: NUMBER OF WEDGES IN EACH VERTEX

$$\text{LCC}(\nu) = \frac{\text{tri}(\nu)}{\text{wed}(\nu)}$$

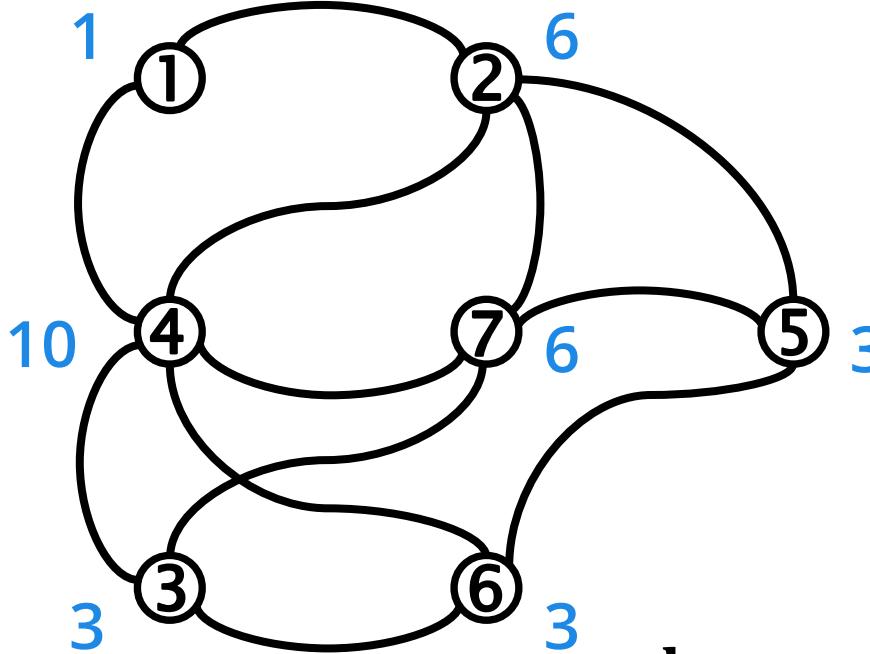
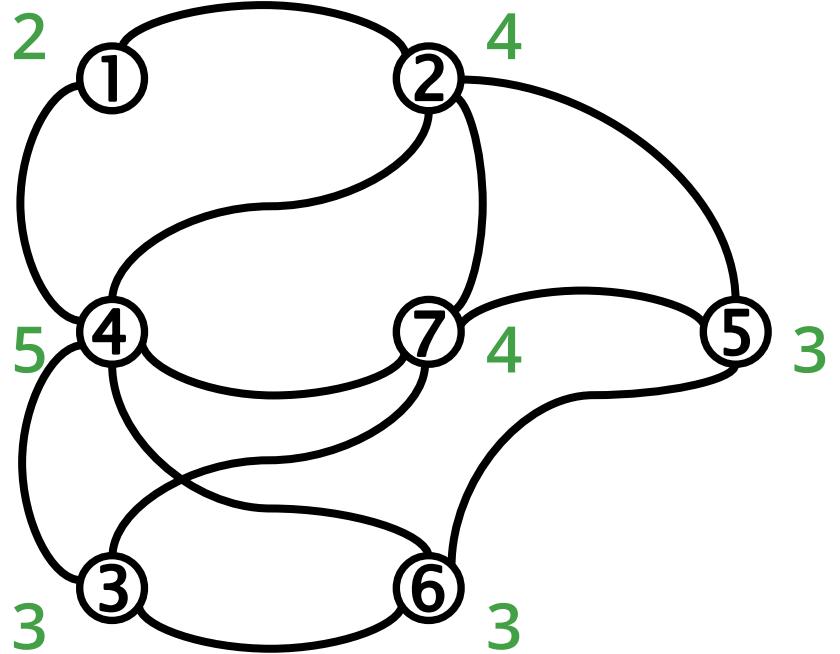
- For  $\text{wed}(\nu)$ , we determine the #wedges for each vertex as the 2-combination of its degree:

$$\text{comb2}(x) = \frac{x \cdot (x - 1)}{2}$$

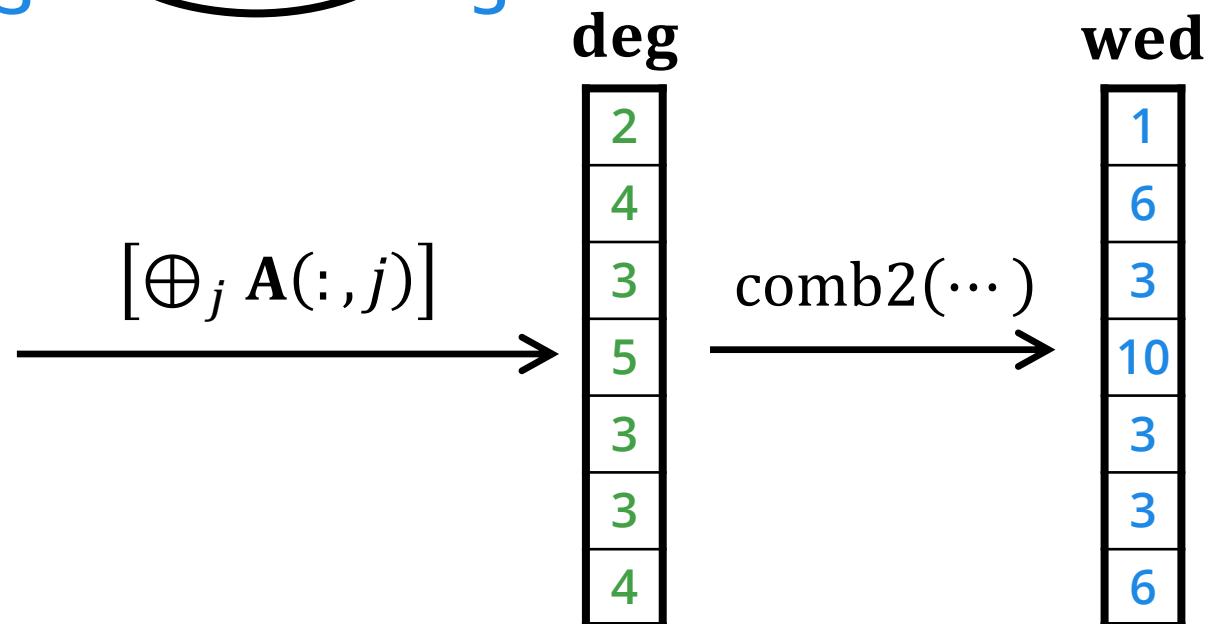
- Given the degrees  $\text{deg} = [\oplus_j \mathbf{A}(:, j)]$ , we compute  $\text{wed}$  by applying a unary function on the elements of the vector:

$$\text{wed} = \text{comb2}(\text{deg})$$

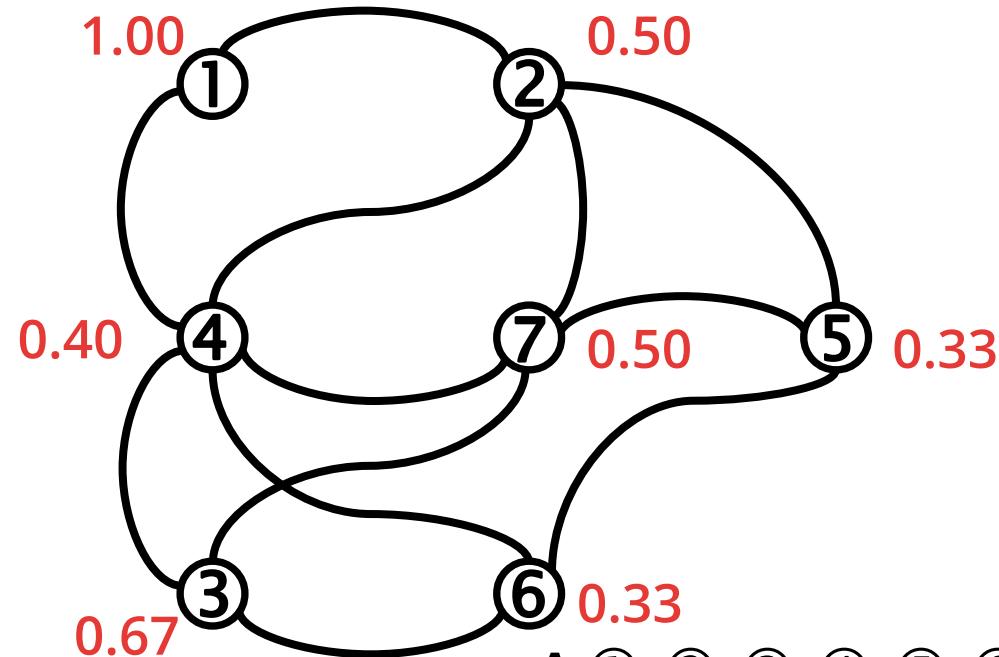
# LCC EXAMPLE: NUMBER OF WEDGES



A	①	②	③	④	⑤	⑥	⑦
①	1			1			
②	1			1	1		1
③				1		1	1
④	1	1	1			1	1
⑤		1			1	1	
⑥			1	1	1		
⑦		1	1	1	1		



# LCC EXAMPLE: COMPLETE ALGORITHM



$A \begin{smallmatrix} 1 & 2 & 3 & 4 & 5 & 6 & 7 \end{smallmatrix}$

①	1	1	1			
②	1		1	1	1	
③			1		1	1
④	1	1	1		1	1
⑤	1			1	1	
⑥			1	1	1	
⑦	1	1	1	1		

$A \begin{smallmatrix} 1 & 2 & 3 & 4 & 5 & 6 & 7 \end{smallmatrix}$

①	1	1	1			
②	1		1	1	1	
③			1		1	1
④	1	1	1		1	1
⑤	1			1	1	
⑥			1	1	1	
⑦	1	1	1	1		

$$\text{TRI}(A) = A \oplus . \otimes A$$

$\text{deg}$

2	
4	
3	
5	
3	
3	
4	

$$[\oplus_j \dots]$$

$\text{comb2}(\dots)$

$\text{tri}$

1	
3	
2	
4	
1	

$\text{wed}$

1	
6	
3	
10	
3	

$\text{lcc}$

1.00	
0.50	
0.67	
0.40	
0.33	
0.33	
0.50	

$$[\oplus_j \dots]$$

/2

1	1	1		
1		2	1	2
		2	1	1
1	2	2		1
1			1	1



# LCC: ALGORITHM

**Input:** adjacency matrix **A**

**Output:** vector **lcc**

**Workspace:** matrix **TRI**, vectors **tri**, **deg**, **wed**, and **lcc**

1. **TRI** $\langle \mathbf{A} \rangle = \mathbf{A} \oplus \cdot \otimes \mathbf{A}$  compute triangle count matrix
2. **tri** =  $[\oplus_j \mathbf{TRI}(:,j)]/2$  reduce to triangle count vector
3. **deg** =  $[\oplus_j \mathbf{A}(:,j)]$  reduce to vertex degree vector
4. **wed** = **comb2**(**deg**) apply **comb2** to get wedge count vector
5. **lcc** = **tri**  $\oslash$  **wed** LCC vector



# LCC: FURTHER OPTIMIZATIONS

Further optimization: use  $\mathbf{L}$ , the lower triangular part of  $\mathbf{A}$ .

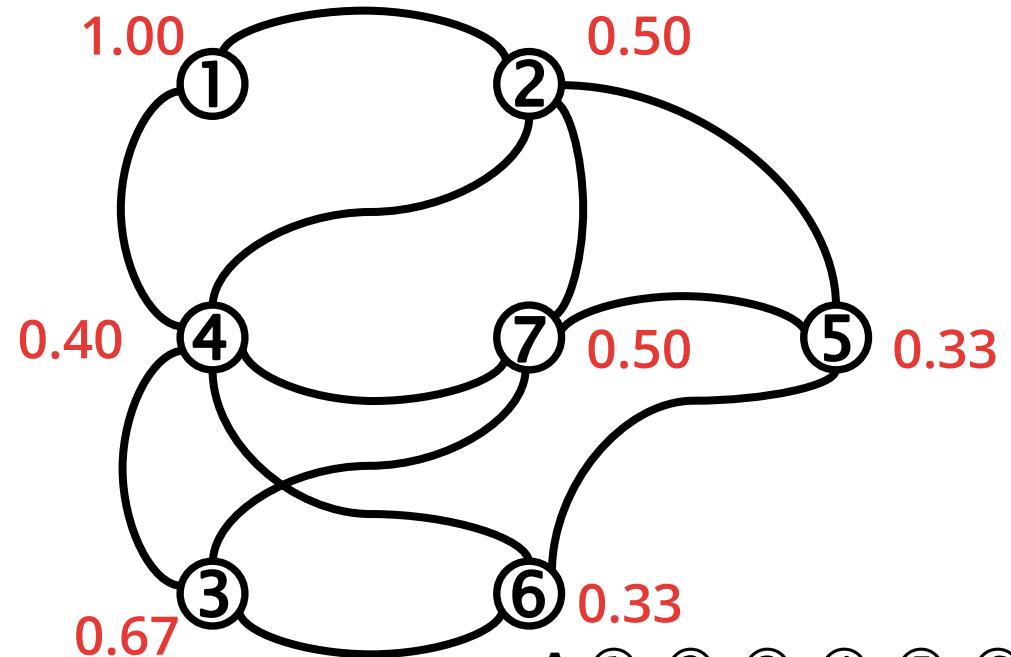
$$\text{TRI}\langle \mathbf{A} \rangle = \mathbf{A} \oplus_{\cdot} \otimes \mathbf{L}$$

The number of wedges is now the 2-combination of  $\deg$ .

$$\text{comb2}(x) = \frac{x \cdot (x - 1)}{2}$$

Permuting the adjacency matrix allows further optimizations.

# LCC EXAMPLE: LOWER TRIANGULAR PART OF MX.



A	1	2	3	4	5	6	7
1	1	1	1				
2	1		1	1	1		
3			1		1	1	
4	1	1	1		1	1	
5		1			1	1	
6		1	1	1	1		
7	1	1	1	1			

L	1	2	3	4	5	6	7
1							
2	1						
3							
4	1	1	1				
5		1					
6			1	1	1		
7	1	1	1	1			

1	1				
1			1	1	
			2		
1	1	2			
	1				
	1				
2	1				

$$\text{TRI}(A) = A \oplus . \otimes L$$

deg

2
4
3
5
3
3
4

$$[\oplus_j \dots]$$

comb2( ... )

tri

1
3
2
4
1

wed

1
6
3
10
3

lcc

1.00
0.50
0.67
0.40
0.33
0.33
0.50

$\ominus$

=

# Graph algorithms in GraphBLAS

## PageRank

# PAGERANK – DEFINITION (LDBC GRAPHALYTICS)

For  $k = 1$  to  $t$  iterations:

$$\text{PR}_0(v) = \frac{1}{n}$$

$$\text{PR}_k(v) = \underbrace{\frac{1 - \alpha}{n}}_{\text{teleport}} + \underbrace{\alpha \cdot \sum_{u \in N_{\text{in}}(v)} \frac{\text{PR}_{k-1}(u)}{|N_{\text{out}}(u)|}}_{\text{influence}} + \underbrace{\alpha \cdot \sum_{w \in dng} \text{PR}_{k-1}(w)}_{\text{dangling}}$$

$\alpha$ : damping factor

$dng$ : dangling vertices,  $dng = \{w \in V : |N_{\text{out}}(w)| = 0\}$

There are dozens of PR definitions, some treat dangling vertices differently.

# PAGERANK – IN LINEAR ALGEBRA

Initially:

$$\mathbf{pr}_0 = [1 \ 1 \dots 1] \oslash n, \quad \mathbf{outd} = [\oplus_j \mathbf{A}(:, j)]$$

In each iteration:

$$PR_k(v) = \frac{1 - \alpha}{n} + \alpha \cdot \sum_{u \in N_{\text{in}}(v)} \frac{PR_{k-1}(u)}{|N_{\text{out}}(u)|} + \frac{\alpha}{n} \cdot \sum_{w \in dng} PR_{k-1}(w)$$

$$\mathbf{pr}_k = \underbrace{\frac{1 - \alpha}{n}}_{\text{constant}} \oplus \underbrace{\alpha \otimes \left( \frac{\mathbf{pr}_{k-1}}{\mathbf{outd}} \right) \oplus \mathbf{.} \otimes \mathbf{A}}_{\text{SpMV}} \oplus \underbrace{\frac{\alpha}{n} \otimes \left[ \oplus_i (\mathbf{pr}_k \otimes \overline{\mathbf{outd}})(i) \right]}_{\text{element-wise sparse vector-dense vector multiplication}}$$



# PAGERANK – ALGORITHM

**Input:** adjacency matrix  $A$ , damping factor  $\alpha$ , #iterations  $t$ , #vertices  $n$

**Output:** PageRank vector  $pr$  (real); **Workspace:** vectors (real)

1.  $pr = [1 \ 1 \dots \ 1] \oslash n$
2. **outdegrees** =  $[\oplus_j A(:,j)]$
3. for  $k = 1$  to  $t$
4.     **importance** =  $\alpha \otimes (pr \oslash \text{outdegrees}) \oplus . \otimes A$
5.     **danglingVertexRanks** $\langle \neg \text{outdegrees} \rangle$  =  $pr(:)$
6.      $totalDanglingRank = \frac{\alpha}{n} \otimes [\oplus_i \text{danglingVertexRanks}(i)]$
7.      $pr = \frac{1-\alpha}{n} \oplus totalDanglingRank \oplus \text{importance}$

# Graph algorithms in GraphBLAS

$k$ -truss

# K-TRUSS

- **Definition:** the  $k$ -truss is a subset of the graph with the same number of vertices, where each edge appears in at least  $k - 2$  triangles in the original graph.



# K-TRUSS ALGORITHM

- Input: adjacency matrix  $\mathbf{A}$ , scalar  $k$
- Output:  $k$ -truss adjacency matrix  $\mathbf{C}$
- Helper:  $f(x, support) = x \geq support$

1.  $\mathbf{C} = \mathbf{A}$
2. for  $i = 1$  to  $n - 1$
3.    $\mathbf{C} \langle \mathbf{C} \rangle = \mathbf{C} \oplus \cdot \wedge \mathbf{C}$                           use the “plus-and” semiring
4.    $\mathbf{C} = f(\mathbf{C}, k - 2)$                           drop entries in  $\mathbf{C}$  less than  $k - 2$
5.   terminate if the number of non-zero values in  $\mathbf{C}$  did not change



# Graph algorithms in GraphBLAS

Community detection using label propagation

# CDLP: COMMUNITY DETECTION USING LABEL PROPAGATION

**Goal:** assign a label to each vertex representing the community it belongs to. The algorithm (originally published in network science) is slightly altered to ensure deterministic execution. Initially:

$$L_0(v) = v$$

In the  $k^{\text{th}}$  iteration:

$$L_k(v) = \min(\arg \max_l |\{u \in N(v) : L_{k-1}(u) = l\}|),$$

where  $N(v)$  is the set of neighbours of  $v$ .

Run for  $t$  iterations or until reaching a fixed point.

 U.N. Raghavan, R. Albert, S. Kumara: *Near linear time algorithm to detect community structures in large-scale networks*, Phys. Rev. E, 2007

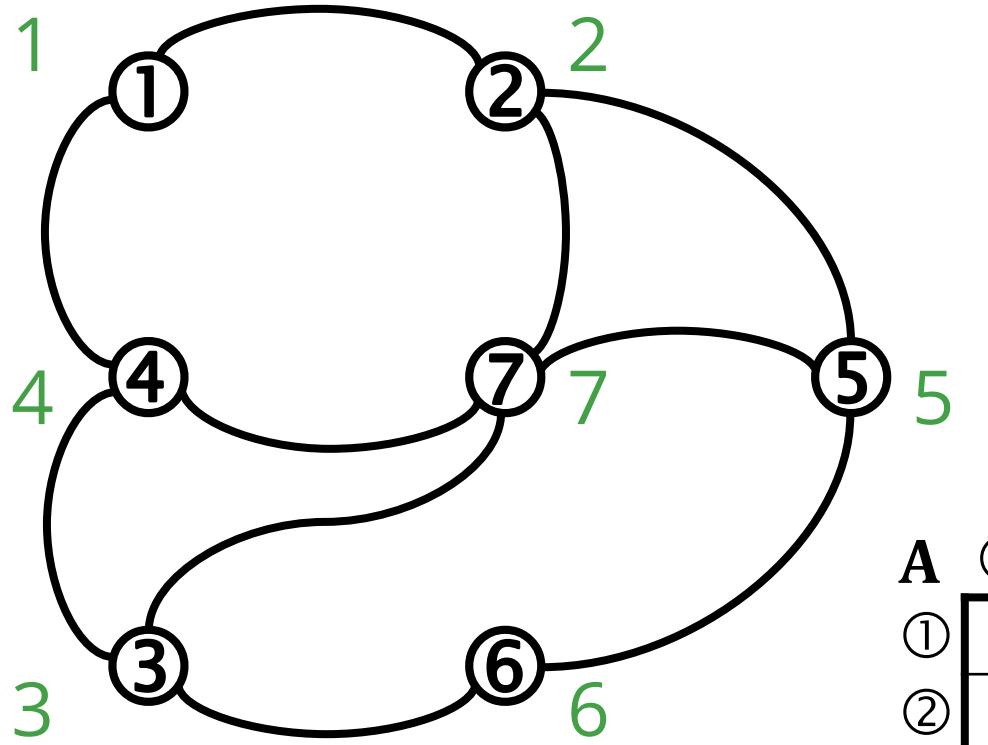
# IDEA: CAPTURE CDLP IN PURE GRAPHBLAS

- Define a semiring that operates over **occurrence vectors**
- $\oplus$  operator: combines two occurrence vectors
  - $\{6 \rightarrow 1, 9 \rightarrow 1\} \oplus \{6 \rightarrow 1, 7 \rightarrow 2\} = \{6 \rightarrow 2, 7 \rightarrow 2, 9 \rightarrow 1\}$
- Convert each element in a row to an occurrence vector
  - $\{6 \rightarrow 1\}, \{6 \rightarrow 1\}, \{7 \rightarrow 1\}, \{7 \rightarrow 1\}, \{9 \rightarrow 1\}$
- Reduce each row into a single occurrence vector:
  - $\{6 \rightarrow 2, 7 \rightarrow 2, 9 \rightarrow 1\}$
- Select the min. mode element from the occurrence vector: 6
- Works on paper, but occurrence vectors need dynamic memory allocation, which leads to poor performance with the current GraphBLAS API

# CDLP IN LINEAR ALGEBRA: FASTER ALGORITHM

- Extract each row from  $\mathbf{F}$ 
  - Easy if the matrix is stored in CSR format
- Select the minimum mode value in each row
  - Sort elements using parallel merge sort
  - Pick the min value that has the longest run (done in a single pass)
- Sort each row  $\mathbf{r}$
- Use the sorted list to compute  $\text{mode}(\mathbf{r})$
- The matrix multiplications are always performed with a diagonal matrix as the second operand so we never need the addition operator. Therefore, we set it to  $\oplus \equiv \text{any}$ .

# CDLP EXAMPLE



- Initially,  $\text{lab} = [1 \ 2 \ \dots \ n]$
- Propagate labels to create a “frequency matrix”:  
 $F = A \text{ any} . \text{sel2nd diag}(\text{lab})$

**diag(lab)**

	1	2	3	4	5	6	7
1	1						
2		2					
3			3				
4				4			
5					5		
6						6	
7							7

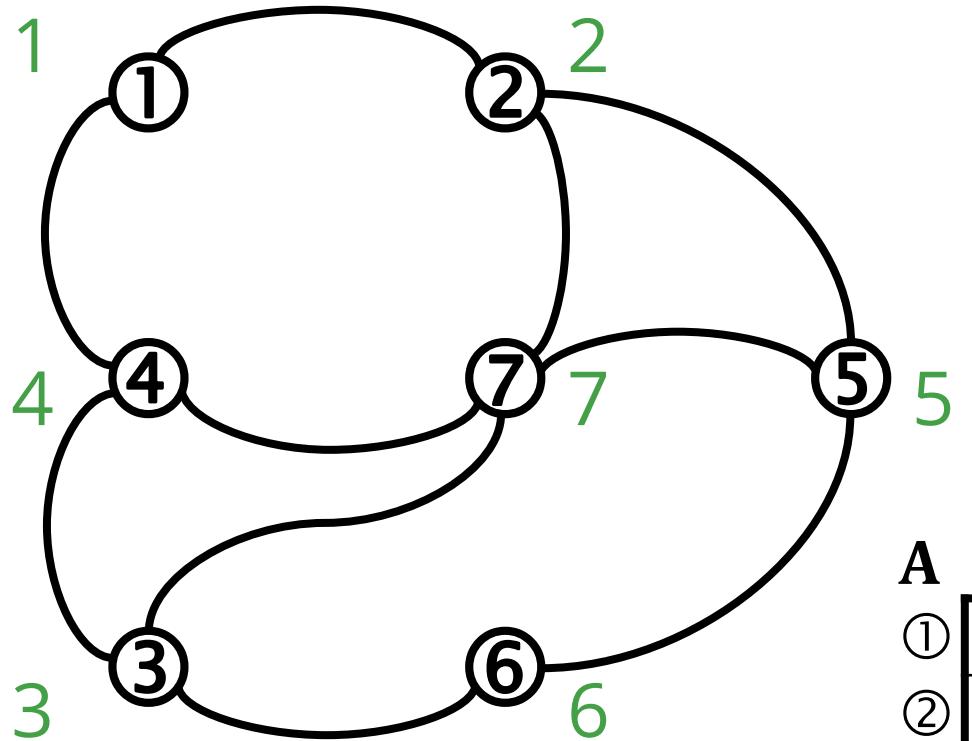
  

**A**

	1	2	3	4	5	6	7
1		•		•			
2	•			•		•	
3				•	•	•	
4	•		•			•	
5		•			•	•	
6			•	•			
7		•	•	•	•		

	2	4		5	7		
1	1			5	7		
2		2		4	6	7	
3			3			7	
4				4			
5					5		
6						6	7
7							7

# CDLP EXAMPLE



step: 1

diag(lab)

①	②	③	④	⑤	⑥	⑦
①	1					
②		2				
③			3			
④				4		
⑤					5	
⑥						6
⑦						7

A

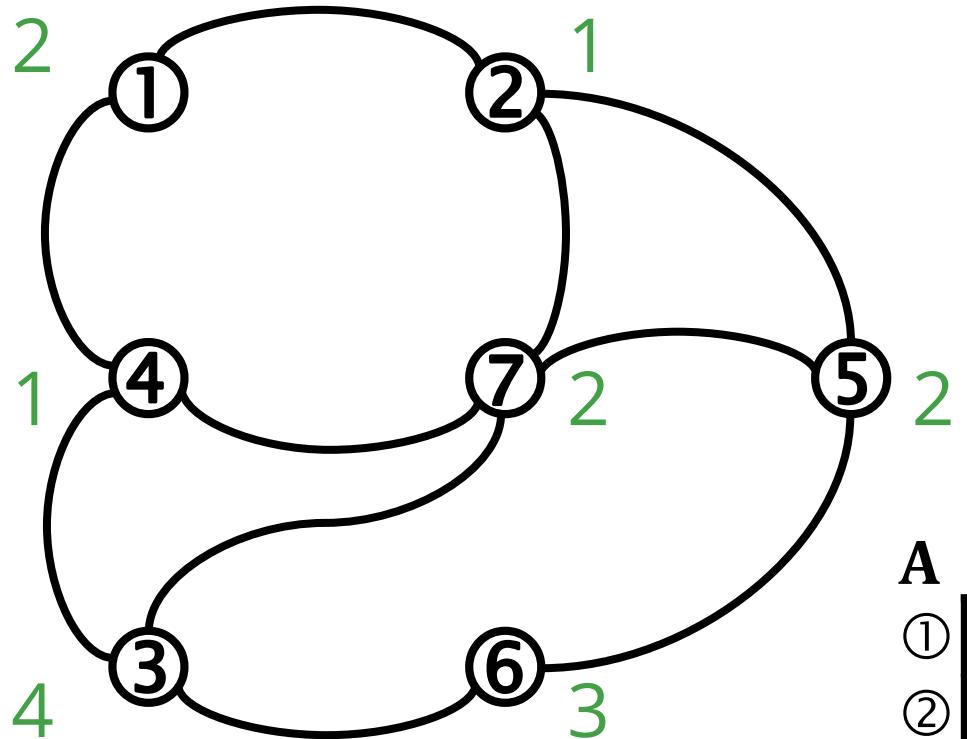
①	②	③	④	⑤	⑥	⑦
①		●		●		
②	●			●		●
③				●	●	●
④	●		●			●
⑤		●		●	●	
⑥			●	●		
⑦		●	●	●	●	

lab'

min. mode →

2	1	4	1	2	3	2
---	---	---	---	---	---	---

# CDLP EXAMPLE



step: 2

A	1	2	3	4	5	6	7
1		•		•			
2	•			•		•	
3				•		•	•
4	•		•			•	
5		•			•		•
6			•		•		
7		•	•	•	•		

diag(lab)

1	2	3	4	5	6	7
1	2					
2		1				
3			4			
4				1		
5					2	
6						3
7						2

lab'

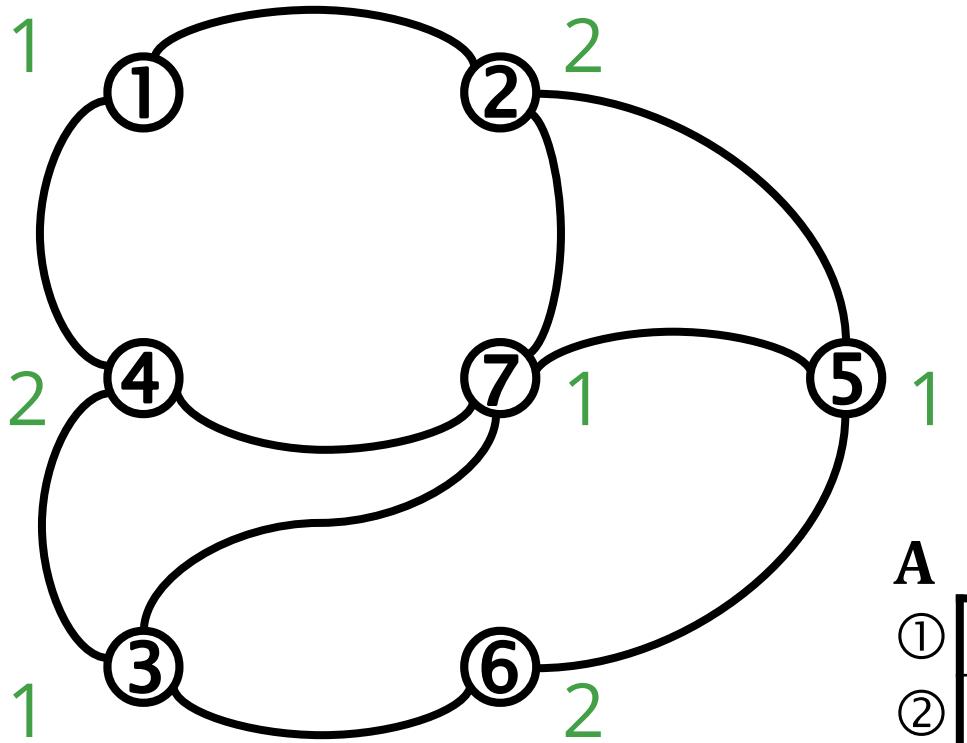
1	1	1	2	2
2			1	3
			2	2
2	4			
		1	3	2
			2	
2				
1	4	1	2	

min.  
mode



1
2
1
2
1
2
1

# CDLP EXAMPLE



step: 3

A	1	2	3	4	5	6	7
1		•		•			
2	•			•		•	
3				•	•	•	
4	•		•			•	
5		•			•	•	
6			•	•	•		
7		•	•	•	•		

diag(lab)

1	2	3	4	5	6	7
1	1					
2		2				
3			1			
4				2		
5					1	
6						2
7						1

lab'

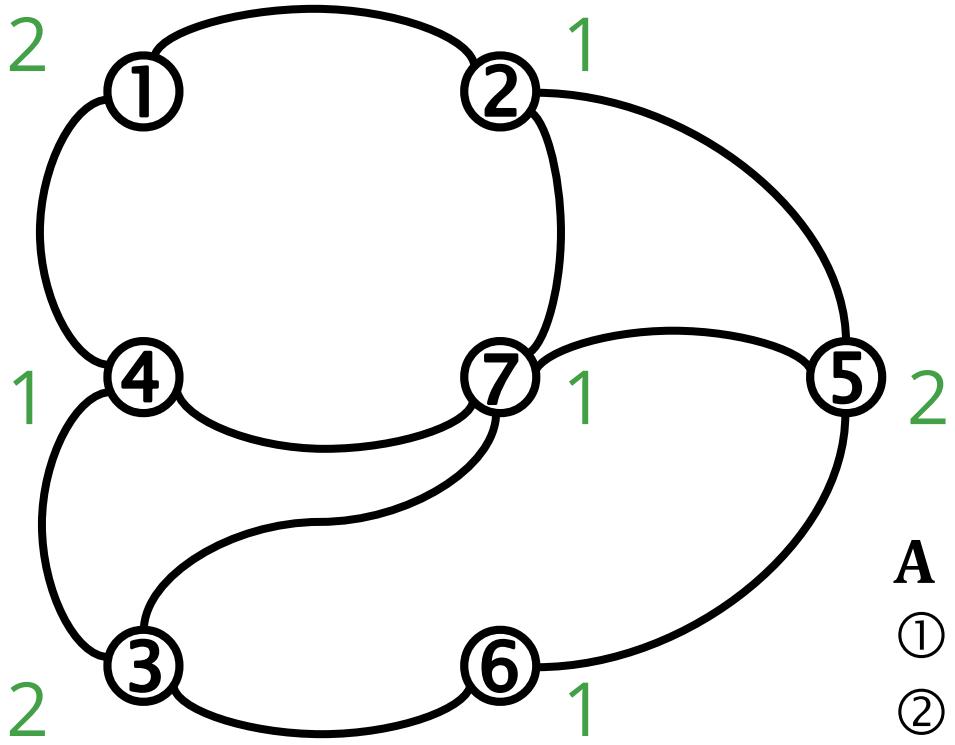
2	2	2	1	1
1		2	2	1
	1	1	2	1
1		1	1	
2	2	2	1	

min.  
mode



2	1	2	1	1
1		2	2	1
2	2	2	1	1
1		1	1	
2	2	2	1	1

# CDLP EXAMPLE



	1	2	3	4	5	6	7
1	2						
2		1					
3			2				
4				1			
5					2		
6						1	
7							1

A	1	2	3	4	5	6	7
1	•		•				
2	•			•		•	
3				•		•	•
4	•		•			•	•
5		•			•	•	
6			•		•		
7		•	•	•	•		

1	1	1	2	1
2			2	1
2		1	1	1
2	2		1	1
1	2	2	2	

lab'	1	2	1	2	1	2	1
min. mode	1	2	1	2	1	2	1

step: 4 – same result as in step 2  
 The original non-deterministic variant of the algorithm is better at avoiding such oscillations.



# CDLP: ALGORITHM

**Input:** adjacency matrix  $\mathbf{A}$ , #vertices  $n$ , #iterations  $t$

**Output:** vector  $\mathbf{lab}$

**Workspace:** matrix  $\mathbf{F}$ , vector  $\mathbf{r}$

1.  $\mathbf{lab} = [1 \ 2 \ \dots \ n]$
2. for  $k = 1$  to  $t$
3.    $\mathbf{F} = \mathbf{A}$  any.sel2nd diag( $\mathbf{lab}$ )
4.   for  $i = 1$  to  $n$
5.      $\mathbf{r} = \mathbf{F}(i, :)$
6.     sort( $\mathbf{r}$ )
7.      $\mathbf{lab}(i) = \text{select\_min\_mode}(\mathbf{r})$



Can be batched and parallelized



# CDLP: ALGORITHM

**Input:** adjacency matrix **A**, #vertices  $n$ , #iterations  $t$

**Output:** vector **lab**

**Workspace:** matrix **F**, vector **r**, array of row indices **I**, array of values **X**

1. **lab** = [1 2 ...  $n$ ]
2. for  $k = 1$  to  $t$
3.   **F** = **A** any.sel2nd diag(**lab**)
4.    $\langle \mathbf{I}, \mathbf{X} \rangle$  = extract\_tuples(**F**)
5.   merge\_sort\_pairs( $\langle \mathbf{I}, \mathbf{X} \rangle$ )
6.   **lab** = for each row in **I**, select min mode value from **X**

# CDLP ON DIRECTED GRAPHS

For directed graphs, we compute the labels  $L_k(v)$  as:

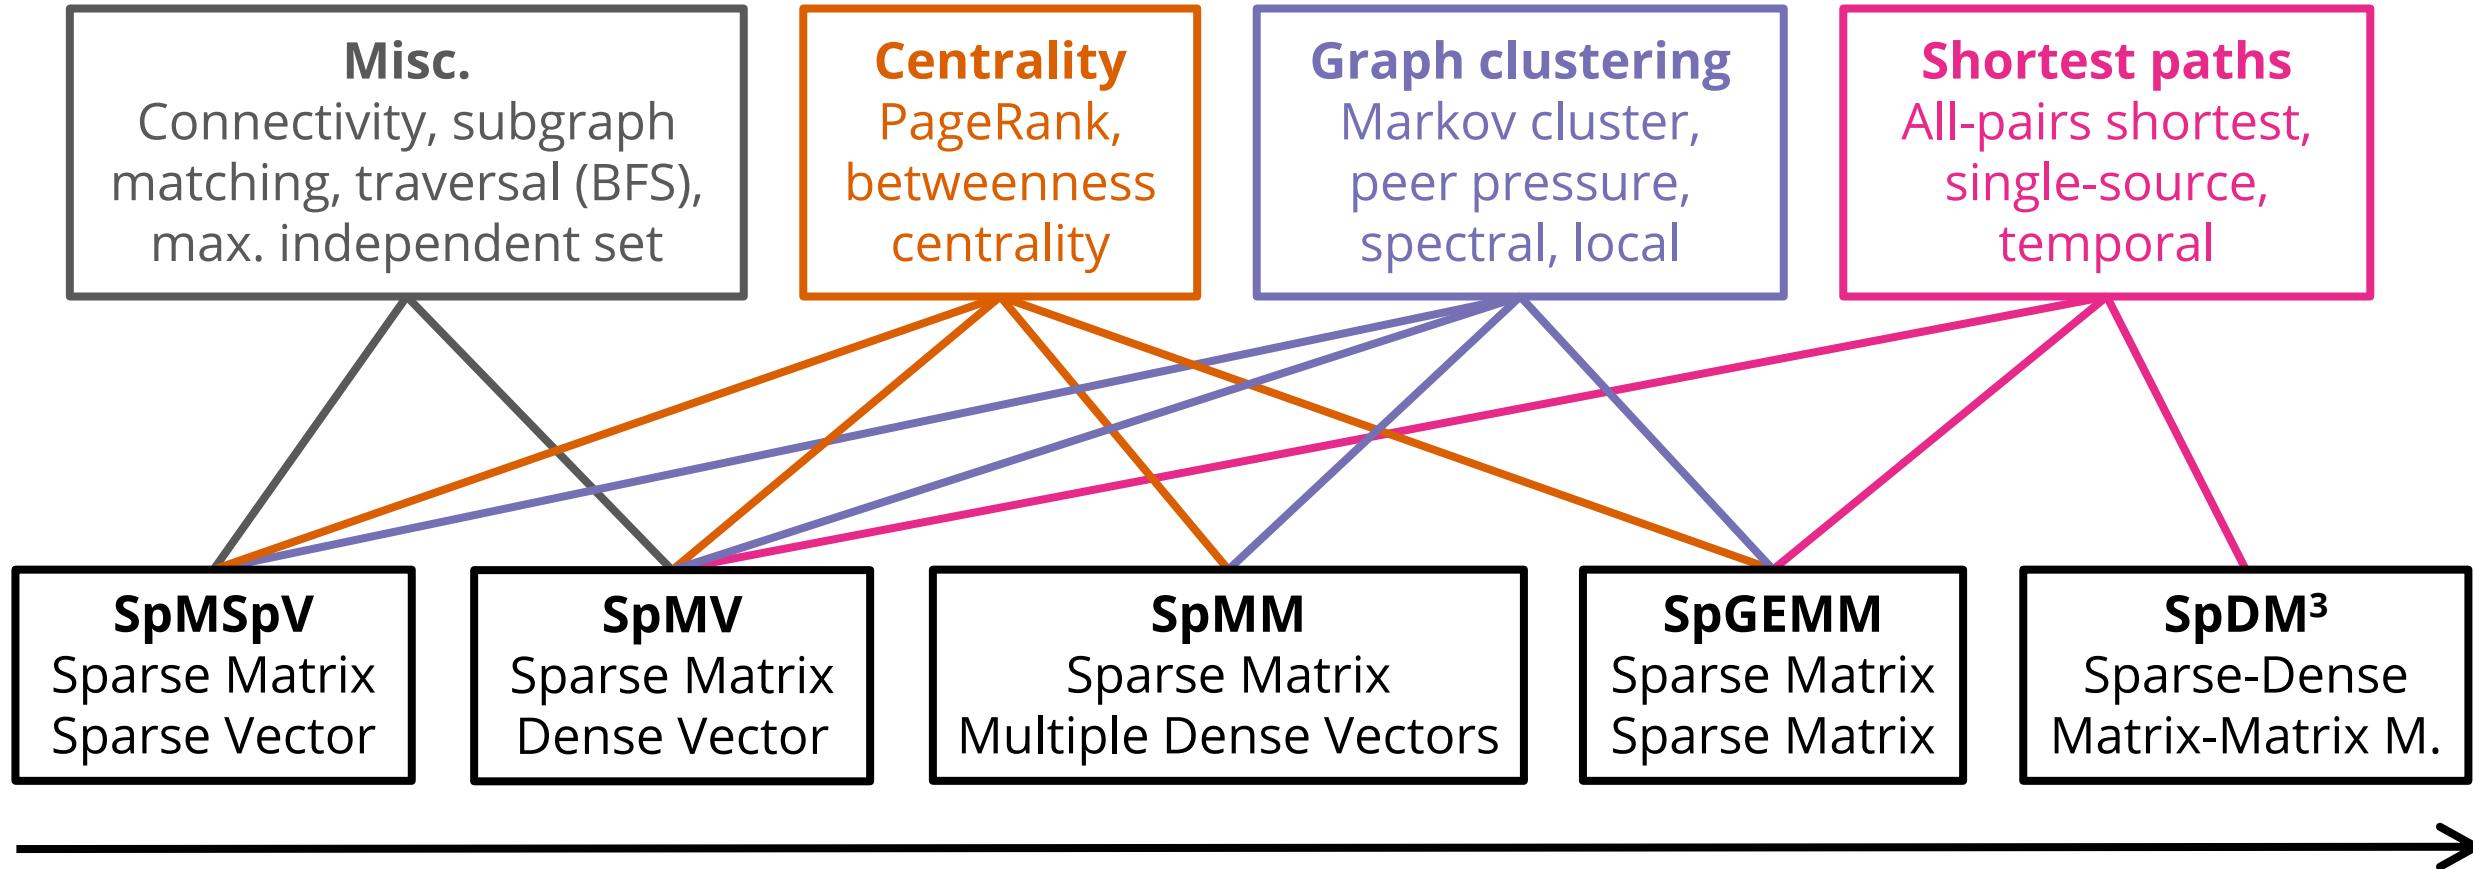
$$\min(\arg \max_l [|\{u \in N_{\text{in}}(v) : L_{k-1}(u) = l\}| + |\{u \in N_{\text{out}}(v) : L_{k-1}(u) = l\}|])$$

- In linear algebra, this can be expressed with two matrices:
    - $\mathbf{F}_{\text{in}} = \mathbf{A}$  any. sel2nd diag(**lab**)
    - $\mathbf{F}_{\text{out}} = \mathbf{A}^T$  any. sel2nd diag(**lab**)
  - Simultaneously iterate over rows  $\mathbf{r}_{\text{in}}$  of  $\mathbf{F}_{\text{in}}$  and  $\mathbf{r}_{\text{out}}$  of  $\mathbf{F}_{\text{out}}$
  - For each row pair, sort  $\mathbf{r}_{\text{in}} \cup \mathbf{r}_{\text{out}}$  and select the minimum mode value
  - Batching also works:
    - $\langle \mathbf{I}_{\text{in}}, \mathbf{X}_{\text{in}} \rangle = \text{extract\_tuples}(\mathbf{F}_{\text{in}})$
    - $\langle \mathbf{I}_{\text{out}}, \mathbf{X}_{\text{out}} \rangle = \text{extract\_tuples}(\mathbf{F}_{\text{out}})$
- } merge\_sort\_pairs( $\langle \mathbf{I}_{\text{in}} \cup \mathbf{I}_{\text{out}}, \mathbf{X}_{\text{in}} \cup \mathbf{X}_{\text{out}} \rangle$ )

# Graph algorithms in GraphBLAS

Graph algorithms & GraphBLAS primitives

# GRAPH ALGORITHMS & GRAPHBLAS PRIMITIVES



Based on the figure in A. Buluç:

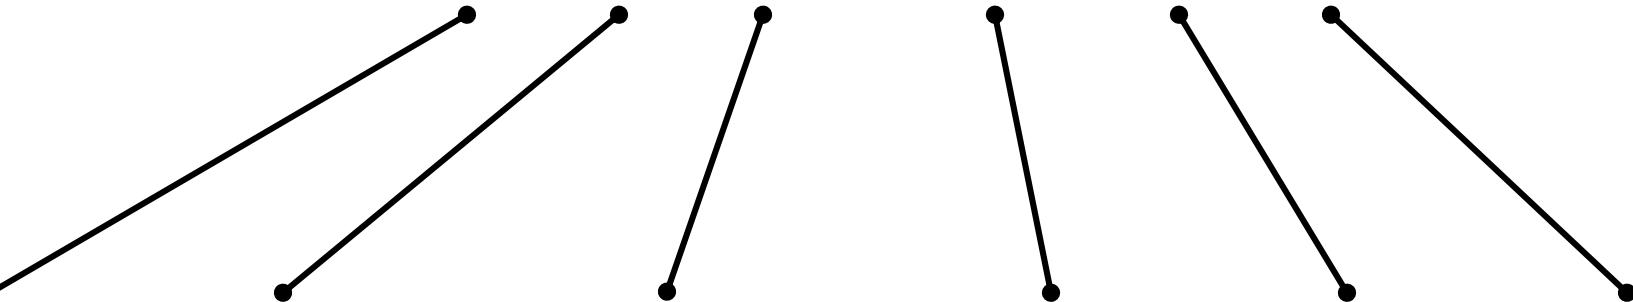
*Graph algorithms, computational motifs, and GraphBLAS*, ECP Meeting 2018

# GraphBLAS and SuiteSparse internals

# GRAPHBLAS C API

- “A crucial piece of the GraphBLAS effort is to translate the mathematical specification to an API that
  - is faithful to the mathematics as much as possible, and
  - enables efficient implementations on modern hardware.”

$$C \leftarrow M \odot= \oplus \cdot \otimes (A^T, B^T)$$



`mxm(Matrix *C, Matrix M, BinaryOp accum, Semiring op, Matrix A, Matrix B, Descriptor desc)`

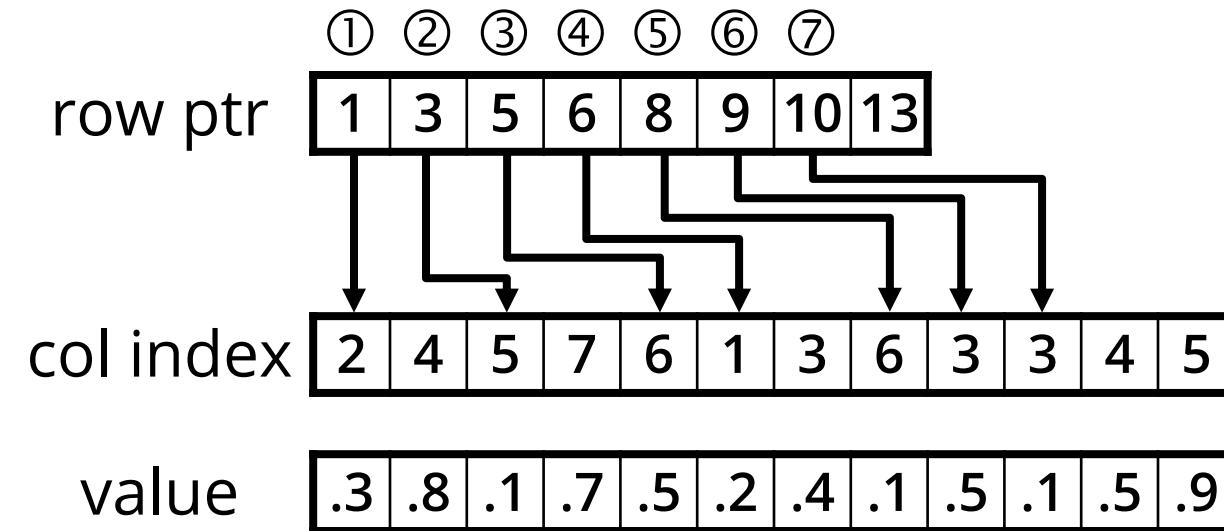


# GRAPHBLAS OBJECTS

- GraphBLAS objects are opaque: the matrix representation can be adjusted to suit the data distribution, hardware, etc.
- The typical representations compressed formats are:
  - CSR: Compressed Sparse Row (also known as CRS)
  - CSC: Compressed Sparse Column (also known as CCS)

A	①	②	③	④	⑤	⑥	⑦
①		.3		.8			
②					.1		.7
③						.5	
④	.2		.4				
⑤						.1	
⑥			.5				
⑦			.1	.5	.9		

CSR representation of A:

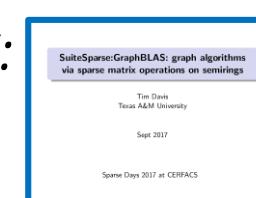


# SUITESPARSE:GRAPHBLAS INTERNALS

- Authored by Prof. Tim Davis at Texas A&M University, based on his SuiteSparse library (used in MATLAB).
- Design decisions, algorithms and data structures are discussed in the TOMS paper and in the User Guide.
- Extensions: methods and types prefixed with GxB.
- Sophisticated load balancer for multi-threaded execution.
- A GPU implementation is work-in-progress.



T.A. Davis: *Algorithm 1000: SuiteSparse:GraphBLAS: graph algorithms in the language of sparse linear algebra*, ACM TOMS, 2019

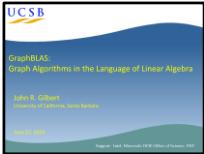


T.A. Davis: *SuiteSparse:GraphBLAS: graph algorithms via sparse matrix operations on semirings*, Sparse Days 2017

# Further reading and libraries

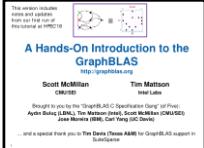
# RESOURCES

## Presentations and tutorials for learning GraphBLAS:



J.R. Gilbert:

*GraphBLAS: Graph Algorithms in the Language of Linear Algebra*, Seminar talk since 2014



S. McMillan and T.G. Mattson:

*A Hands-On Introduction to the GraphBLAS*, Tutorial at HPEC since 2018



A. Buluç:

*GraphBLAS: Concepts, algorithms, and applications*, Scheduling Workshop, 2019



M. Kumar, J.E. Moreira, P. Pattnaik:

*GraphBLAS: Handling performance concerns in large graph analytics*, Computing Frontiers 2018



List of GraphBLAS-related books, papers, presentations, posters, software, etc.  
[szarnyasz/graphblas-pointers](https://github.com/szarnyasz/graphblas-pointers)

# THE LAGRAPH LIBRARY

- Similar to the LAPACK library for BLAS
- Uses SuiteSparse:GraphBLAS
- Implementations of common algorithms
  - BFS, SSSP, LCC, PageRank, Borůvka
  - Triangle count,  $k$ -truss
  - CDLP (community detection using label propagation)
  - Weakly connected components, Strongly Connected Components
  - Betweenness centrality
  - Deep neural network

T.G. Mattson et al.: *LAGraph: A Community Effort to Collect Graph Algorithms Built on Top of the GraphBLAS*, GrAPL@IPDPS 2019



[GraphBLAS/LAGraph](#)

# REQUIREMENTS BY GRAPH COMPUTATIONS

Libraries for linear-algebra based graph processing support the following features (prioritized):

1. Sparse matrices      For reasonable performance
2. Arbitrary semirings    For expressive power
3. Masking                A big reduction in complexity for some algos
4. Parallel execution     Constant speedup, ideally by #threads

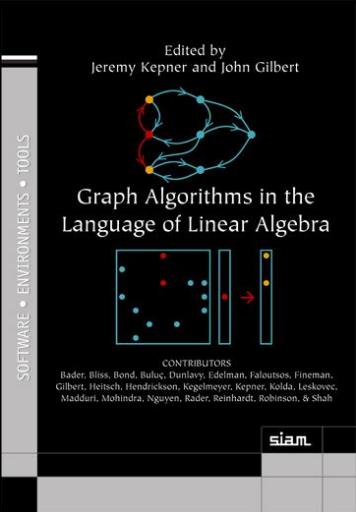
Most libraries **only satisfy requirement #1**: Intel MKL, Eigen, Boost uBLAS, MTL4, Armadillo, NIST Sparse BLAS, GMM++, CUSP, Numpy

Exceptions are the Efficient Java Matrix Library (EJML) and Julia's SparseArrays library, where arbitrary semirings can be used.

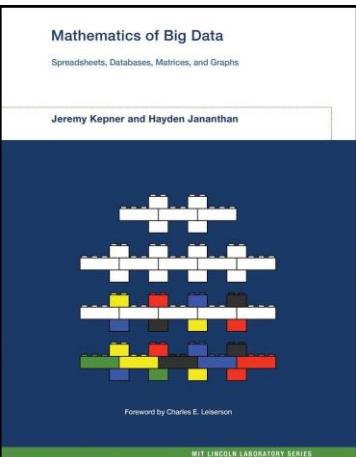
# GRAPHBLAS PAPERS AND BOOKS

- *Standards for Graph Algorithm Primitives*
  - Position paper by 19 authors @ IEEE HPEC 2013
- *Novel Algebras for Advanced Analytics in Julia*
  - Technical paper on semirings in Julia @ IEEE HPEC 2013
- *Mathematical Foundations of the GraphBLAS*
  - Theory paper by 16 authors @ IEEE HPEC 2016
- *Design of the GraphBLAS C API*
  - Design decisions and overview of the C API @ GABB@IPDPS 2017
- *Algorithm 1000: SuiteSparse:GraphBLAS: graph algorithms in the language of sparse linear algebra*
  - Algorithms in the SuiteSparse implementation @ ACM TOMS 2019

# BOOKS



- *Graph Algorithms in the Language of Linear Algebra*
  - Edited by J. Kepner and J.R. Gilbert, published by SIAM in 2011
  - Algorithms for connected components, shortest paths, max-flow, betweenness centrality, spanning tree, graph generation, etc.
  - Algorithms and data structure for fast matrix multiplication
  - Predates GraphBLAS: preliminary notation, no API usage



- *Mathematics of Big Data*
  - Authored by Jananthan & Kepner, published by MIT Press in 2018
  - Generalizes the semiring-based approach for associative arrays
  - Contains important papers, including the HPEC'16 paper above
  - Discusses D4M (Dynamic Distributed Dimensional Data Model)

# GRAPHBLAS COMMUNITY

Wiki: [graphblas.org](http://graphblas.org) | Communication: primarily mailing list

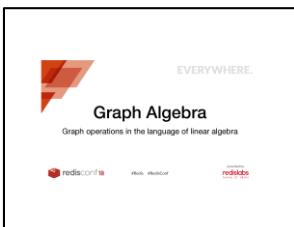
## Annual events:

- May: IEEE IPDPS conference
  - GrAPL workshop (Graphs, Architectures, Programming and Learning), a merger of
    - GABB (Graph Algorithms Building Blocks)
    - GraML (Graph Algorithms and Machine Learning)
  - See [graphanalysis.org](http://graphanalysis.org) for previous editions
- Sep: IEEE HPEC conference
  - GraphBLAS BoF meeting
- Nov: IEEE/ACM Supercomputing conference
  - GraphBLAS Working Group
  - IA<sup>3</sup> workshop (Workshop on Irregular Applications: Architectures and Algorithms)

Blog: [AldenMath](http://AldenMath) by Timothy Alden Davis

# REDISGRAPH

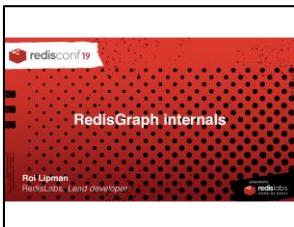
- Graph database built on top of Redis with partial (but extending) support for the Cypher language
- Uses SuiteSparse:GraphBLAS for graph operations
- Preliminary benchmark results show good performance on traversal-heavy workloads



R. Lipman, T.A. Davis:

*Graph Algebra – Graph operations in the language of linear algebra*

RedisConf 2018



R. Lipman:

*RedisGraph internals*

RedisConf 2019

# GRAPHBLAS IMPLEMENTATIONS

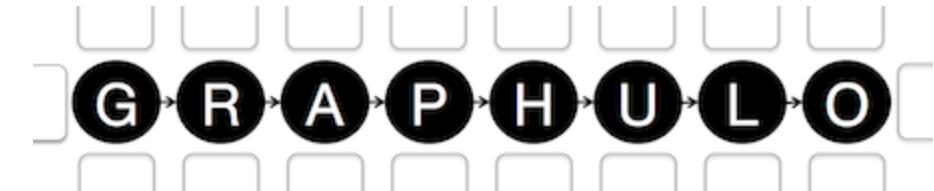
- SuiteSparse:GraphBLAS
  - v1.0.0: Nov 2017 – sequential
  - v3.0.1: July 2019 – parallel
  - v4.0.1draft: Dec 2020 – many optimizations, incl. bitmap format
- IBM GraphBLAS
  - Complete implementation in C++, released in May 2018
  - Concise but sequential
- GBTL (GraphBLAS Template Library): C++
  - v1.0: parallel but no longer maintained
  - v2.0, v3.0: sequential
- GraphBLAST: GPU implementation, based on GBTL

# GRAPHULO

- Build on top of the Accumulo distributed key-value store
- Written in Java
- Focus on scalability

V. Gadepally et al.:

*Graphulo: Linear Algebra Graph Kernels  
for NoSQL Databases*, GABB@IPDPS 2015



# COMBBLAS: COMBINATORIAL BLAS

- “an extensible distributed memory parallel graph library offering a small but powerful set of linear algebra primitives”
- Not a GraphBLAS implementation but serves as an incubator for new ideas that may later find their way into GraphBLAS
- Scales to 250k+ CPU cores
- Used on supercomputers such as Cray



A. Buluç, J.R. Gilbert: *The Combinatorial BLAS: design, implementation, and application*, International Journal of High Performance Computing Applications, 2011

# PYGRAPHBLAS: PYTHON WRAPPER

- Goal: Pythonic GraphBLAS wrapper, close to pseudo-code
- See example code for SSSP and triangle count
- Comes with Jupyter notebooks

```
def sssp(matrix, start):  
    v = Vector.from_type(matrix.gb_type, matrix.nrows)  
    v[start] = 0  
  
    with min_plus, Accum(min_int64):  
        for _ in range(matrix.nrows):  
            w = Vector.dup(v)  
            v @= matrix  
            if w == v:  
                break  
  
    return v
```

```
def sandia(A, L):  
    return L.mxm(L, mask=L).reduce_int()
```

```
sandia(M, M.tril())
```



[michelp/pygraphblas](https://github.com/michelp/pygraphblas)

# GRBLAS: PYTHON WRAPPER

- Goal: wrapper with an almost 1-to-1 mapping to the GrB API
  - Comes with a Conda package
  - Compiles user-defined functions to C
  - Supports visualization

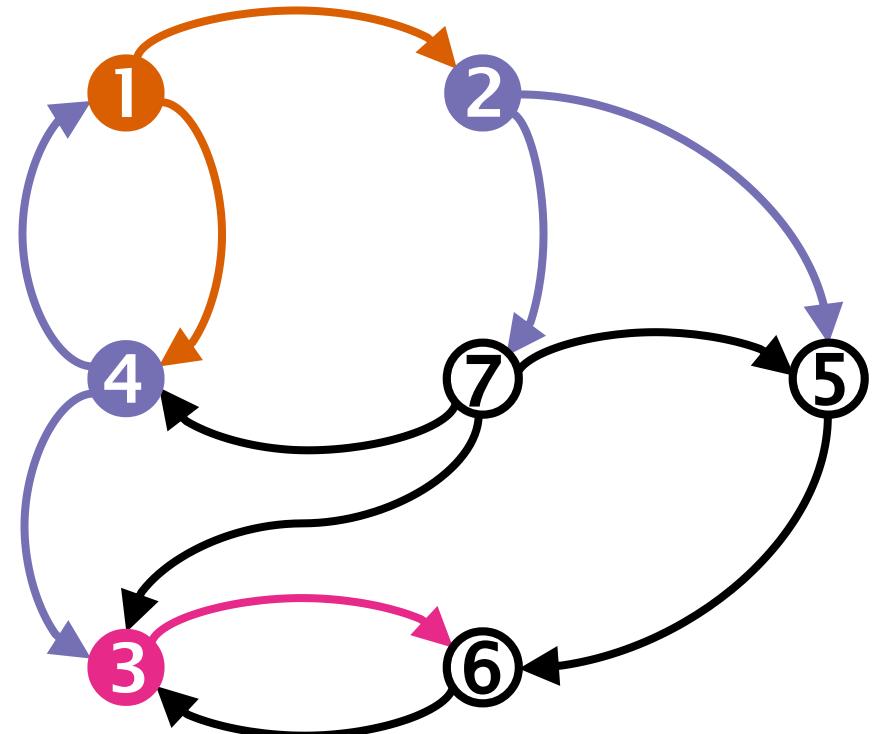
```
M(mask, accum) << A.mxm(B, semiring)           # mxm
w(mask, accum) << A.mxv(v, semiring)           # mxv
w(mask, accum) << v.vxm(B, semiring)           # vxm
M(mask, accum) << A.ewise_add(B, binaryop)     # eWiseAdd
M(mask, accum) << A.ewise_mult(B, binaryop)    # eWiseMult
M(mask, accum) << A.kronecker(B, binaryop)      # kronecker
M(mask, accum) << A.T                           # transpose
```



[metagraph-dev/grblas](https://github.com/metagraph-dev/grblas)

# Parallelism in GraphBLAS

# PARALLELISM IN GRAPHBLAS



traversals

F

t1

t2

t3

vertices

A

1	2	3	4	5	6	7
1				●		
2					●	●
3						●
4	●				●	
5						●
6			●			
7				●	●	●

Frontier any . pair A

# THE CASE FOR LINEAR ALGEBRA-BASED GRAPH ALGORITHMS

Many irregular applications contain coarse-grained parallelism that can be exploited by abstractions at the proper level.

Traditional graph computation	Graphs in the language of linear algebra
Data-driven, unpredictable communication	Fixed communication patterns
Irregular and unstructured, poor locality of reference	Operations on matrix blocks exploit memory hierarchy
Fine-grained data accesses, dominated by latency	Coarse-grained parallelism, bandwidth-limited



D. Bader et al., *The GraphBLAS effort and its implications for Exascale*,  
SIAM Workshop on Exascale Applied Mathematics Challenges and Opportunities, 2014

# Summary

# SUMMARY

- Linear algebra is a powerful abstraction
  - Good expressive power
  - Concise formulation of most graph algorithms
  - Good performance
  - Still lots of ongoing research
- Trade-offs:
  - Learning curve (maths, C programming, GraphBLAS API)
  - Some algorithms are difficult to formulate in linear algebra
  - Only a few GraphBLAS implementations (yet)
- Overall: GraphBLAS is a good abstraction layer for graph algorithms in the age of heterogeneous hardware

*“Nuances”* – Some important  
adjustments to the definitions

# GRAPHBLAS SEMIRINGS\*

The GraphBLAS specification defines semirings as follows:

$\langle D_{\text{out}}, D_{\text{in}_1}, D_{\text{in}_2}, \oplus, \otimes, 0 \rangle$  structure is a *GraphBLAS semiring* defined by

- $D_{\text{out}}$ ,  $D_{\text{in}_1}$ , and  $D_{\text{in}_2}$  three domains
- $\oplus: D_{\text{out}} \times D_{\text{out}} \rightarrow D_{\text{out}}$  an associative and commutative addition operation
- $\otimes: D_{\text{in}_1} \times D_{\text{in}_2} \rightarrow D_{\text{out}}$  a multiplicative operation
- $0 \in D_{\text{out}}$  an identity element for  $\oplus$

$A = \langle D_{\text{out}}, \oplus, 0 \rangle$  is a commutative monoid.

$F = \langle D_{\text{out}}, D_{\text{in}_1}, D_{\text{in}_2}, \otimes \rangle$  is a closed binary operator.

“It is expected that implementations will utilize IEEE-754 floating point arithmetic, which is not strictly associative.” (C API specification)

# NOTATION\*

## Symbols:

- **A, B, C, M** – matrices
- **u, v, w, m** – vectors
- $s, k$  – scalar
- $i, j$  – indices
- $\langle \mathbf{m} \rangle, \langle \mathbf{M} \rangle$  – masks

## Operators:

- $\oplus$  – addition
- $\otimes$  – multiplication,  $\oslash$  – division
- $^T$  – transpose
- $\odot$  – accumulator

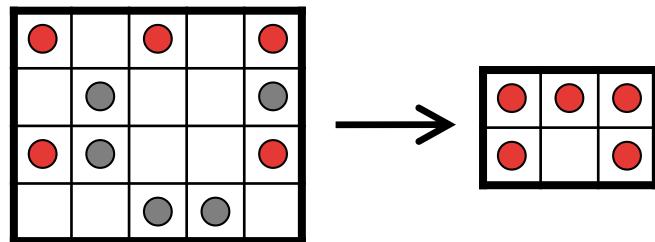
This table contains all GrB and GxB (SuiteSparse-specific) operations.

symbol	operation	notation
$\odot \cdot \otimes$	matrix-matrix multiplication	$C\langle \mathbf{M} \rangle \odot= A \oplus \cdot \otimes B$
$\oplus \cdot \otimes$	vector-matrix multiplication	$w\langle \mathbf{m} \rangle \odot= v \oplus \cdot \otimes A$
	matrix-vector multiplication	$w\langle \mathbf{m} \rangle \odot= A \oplus \cdot \otimes v$
$\otimes$	element-wise multiplication (set intersection of patterns)	$C\langle \mathbf{M} \rangle \odot= A \otimes B$ $w\langle \mathbf{m} \rangle \odot= u \otimes v$
$\oplus$	element-wise addition (set union of patterns)	$C\langle \mathbf{M} \rangle \odot= A \oplus B$ $w\langle \mathbf{m} \rangle \odot= u \oplus v$
$f$	apply unary operator	$C\langle \mathbf{M} \rangle \odot= f(A)$ $w\langle \mathbf{m} \rangle \odot= f(v)$
$[\oplus \dots]$	reduce to vector	$w\langle \mathbf{m} \rangle \odot= [\oplus_j A(:,j)]$
	reduce to scalar	$s \odot= [\oplus_{ij} A(i,j)]$
$A^T$	transpose matrix	$C\langle \mathbf{M} \rangle \odot= A^T$
-	extract submatrix	$C\langle \mathbf{M} \rangle \odot= A(i,j)$ $w\langle \mathbf{m} \rangle \odot= v(i)$
	assign submatrix with submask for $C(I,J)$	$C\langle \mathbf{M} \rangle(i,j) \odot= A$ $w\langle \mathbf{m} \rangle(i) \odot= v$
-	assign submatrix with mask for $C$	$C(i,j)\langle \mathbf{M} \rangle \odot= A$ $w(i)\langle \mathbf{m} \rangle \odot= v$
	apply select operator (GxB)	$C\langle \mathbf{M} \rangle \odot= f(A,k)$ $w\langle \mathbf{m} \rangle \odot= f(v,k)$
-	Kronecker product	$C\langle \mathbf{M} \rangle \odot= \text{kron}(A,B)$

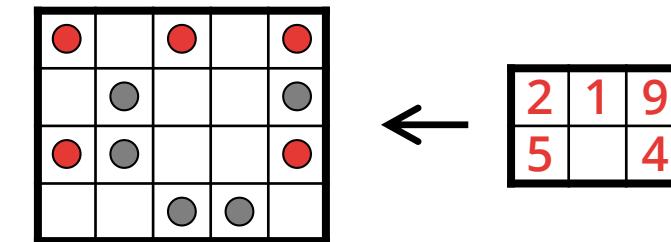
Not included in the simplified table

# LINEAR ALGEBRAIC PRIMITIVES FOR GRAPHS #3\*

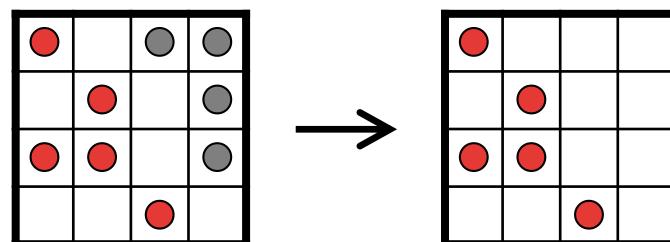
Sparse matrix extraction:  
induced subgraph



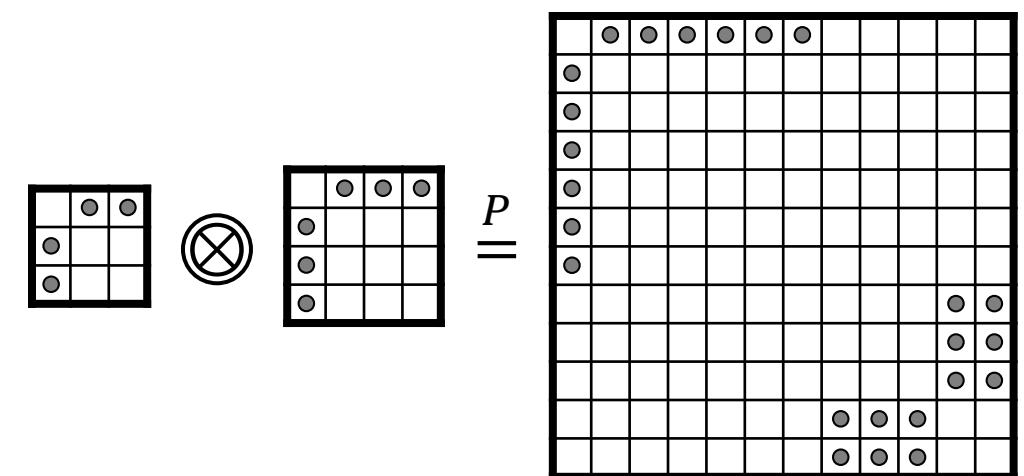
Sparse submatrix assignment:  
replace subgraph



Sparse matrix selection:  
filtering edges



Kronecker product:  
graph generation



# MATRIX-VECTOR MULTIPLICATION\*

The operation  $\mathbf{v} \oplus_{\cdot} \otimes \mathbf{A}$  gives the vertices reachable from the ones in  $\mathbf{v}$ . However, GraphBLAS publications and implementations often use  $\mathbf{A}^T \oplus_{\cdot} \otimes \mathbf{v}$  instead. The difference between these is that the former produces a row vector, while the latter produces a column vector:

$$\mathbf{v} \oplus_{\cdot} \otimes \mathbf{A} \equiv (\mathbf{A}^T \oplus_{\cdot} \otimes \mathbf{v}^T)^T$$

The GraphBLAS does not distinguish row/column vectors, therefore the notations are (formally) equivalent:

$$\mathbf{v} \oplus_{\cdot} \otimes \mathbf{A} \equiv \mathbf{A}^T \oplus_{\cdot} \otimes \mathbf{v}$$

# ELEMENT-WISE SUBTRACTION

*Element-wise subtraction* can be defined as an *element-wise addition* on the INT64\_MINUS monoid. It has the following semantics  $\mathbf{C} = \mathbf{A} \ominus \mathbf{B}$  is computed on the union of the patterns of the input matrices  $\mathbf{A}$  and  $\mathbf{B}$ .

For cells where only one input matrix has a non-zero value but the other does not (e.g.  $\mathbf{B}[0,0] = 1$  but  $\mathbf{A}[0,0]$  is empty), the result is the non-zero value:  $\mathbf{C}[0,0] = 1$ .

$$\begin{array}{c} \mathbf{A} \\ \boxed{\begin{array}{|c|} \hline 1 \\ \hline \end{array}} \end{array} \quad \ominus \quad \begin{array}{c} \mathbf{B} \\ \boxed{\begin{array}{|c|c|} \hline 1 & 1 \\ \hline 1 & 1 \\ \hline \end{array}} \end{array} \quad = \quad \begin{array}{c} \mathbf{C} \\ \boxed{\begin{array}{|c|c|} \hline 1 & 0 \\ \hline 1 & 1 \\ \hline \end{array}} \end{array}$$

This might come across as counter-intuitive first but it confirms the specification:

3070 The intermediate matrix  $\tilde{\mathbf{T}} = \langle \mathbf{D}_{out}(\text{op}), \text{nrows}(\tilde{\mathbf{A}}), \text{ncols}(\tilde{\mathbf{A}}), \{(i, j, T_{ij}) : \text{ind}(\tilde{\mathbf{A}}) \cap \text{ind}(\tilde{\mathbf{B}}) \neq \emptyset\} \rangle$   
3071 is created. The value of each of its elements is computed by

3072  $T_{ij} = (\tilde{\mathbf{A}}(i, j) \oplus \tilde{\mathbf{B}}(i, j)), \forall (i, j) \in \text{ind}(\tilde{\mathbf{A}}) \cap \text{ind}(\tilde{\mathbf{B}})$

3073  $T_{ij} = \tilde{\mathbf{A}}(i, j), \forall (i, j) \in (\text{ind}(\tilde{\mathbf{A}}) - (\text{ind}(\tilde{\mathbf{B}}) \cap \text{ind}(\tilde{\mathbf{A}})))$

3074  $T_{ij} = \tilde{\mathbf{B}}(i, j), \forall (i, j) \in (\text{ind}(\tilde{\mathbf{B}}) - (\text{ind}(\tilde{\mathbf{B}}) \cap \text{ind}(\tilde{\mathbf{A}})))$

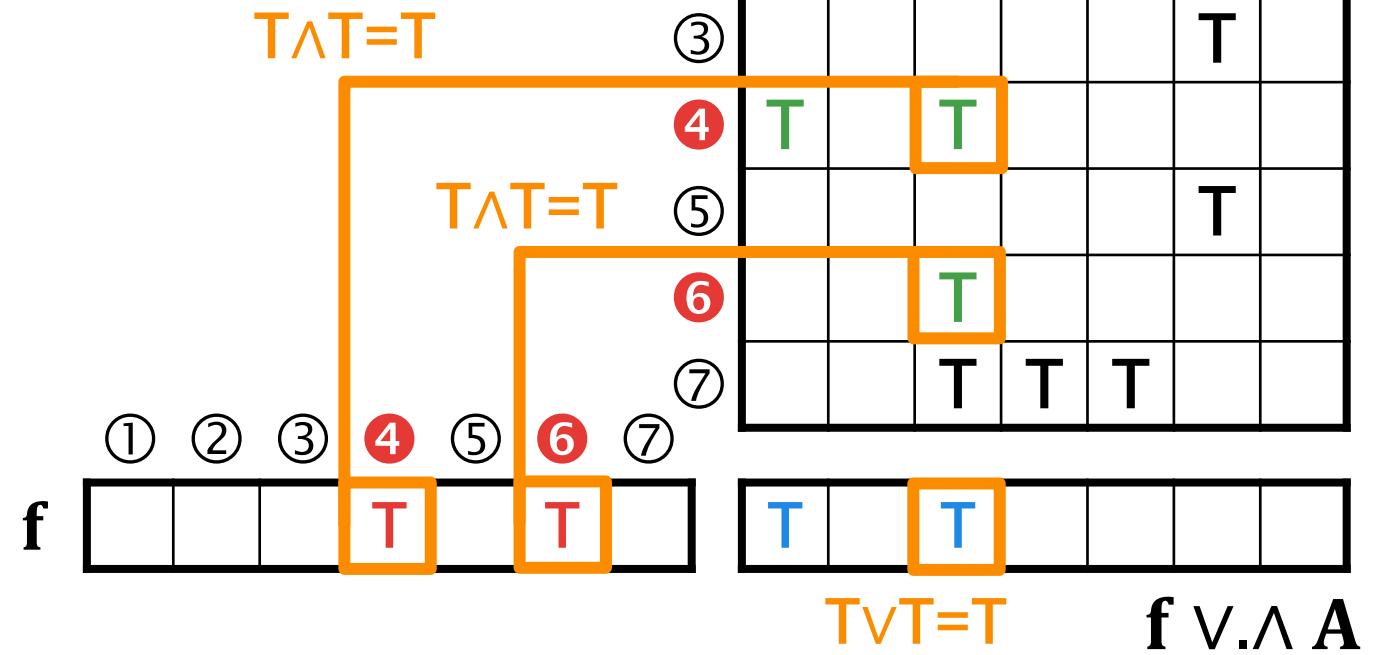
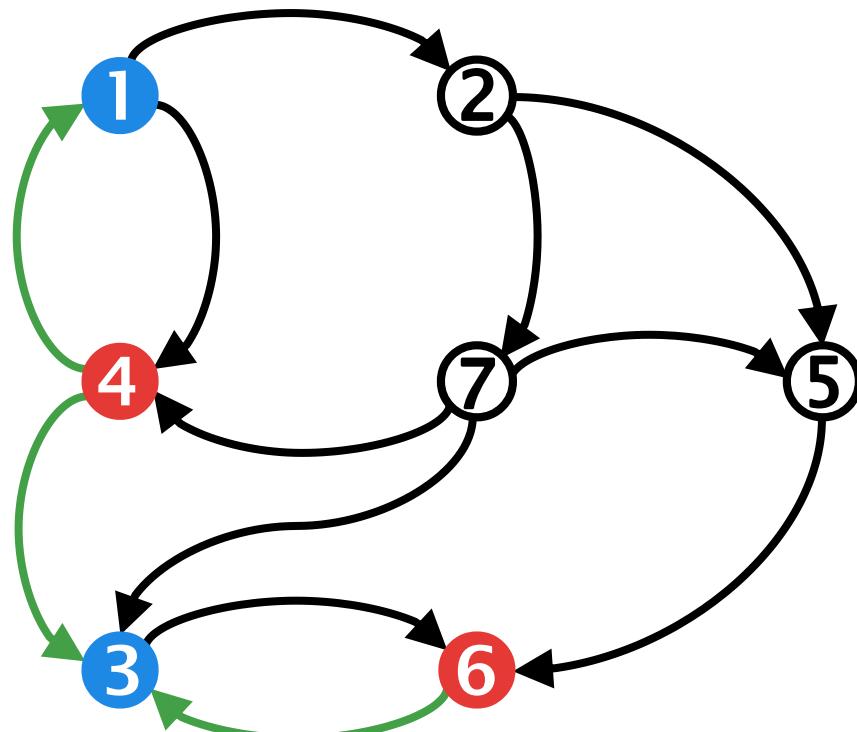
3075  
3076  
3077 where the difference operator in the previous expressions refers to set difference.

# More semirings

# MATRIX-VECTOR MULTIPLICATION SEMANTICS

semiring	set	$\oplus$	$\otimes$	0
Ior-land	{T, F}	$\vee$	$\wedge$	F

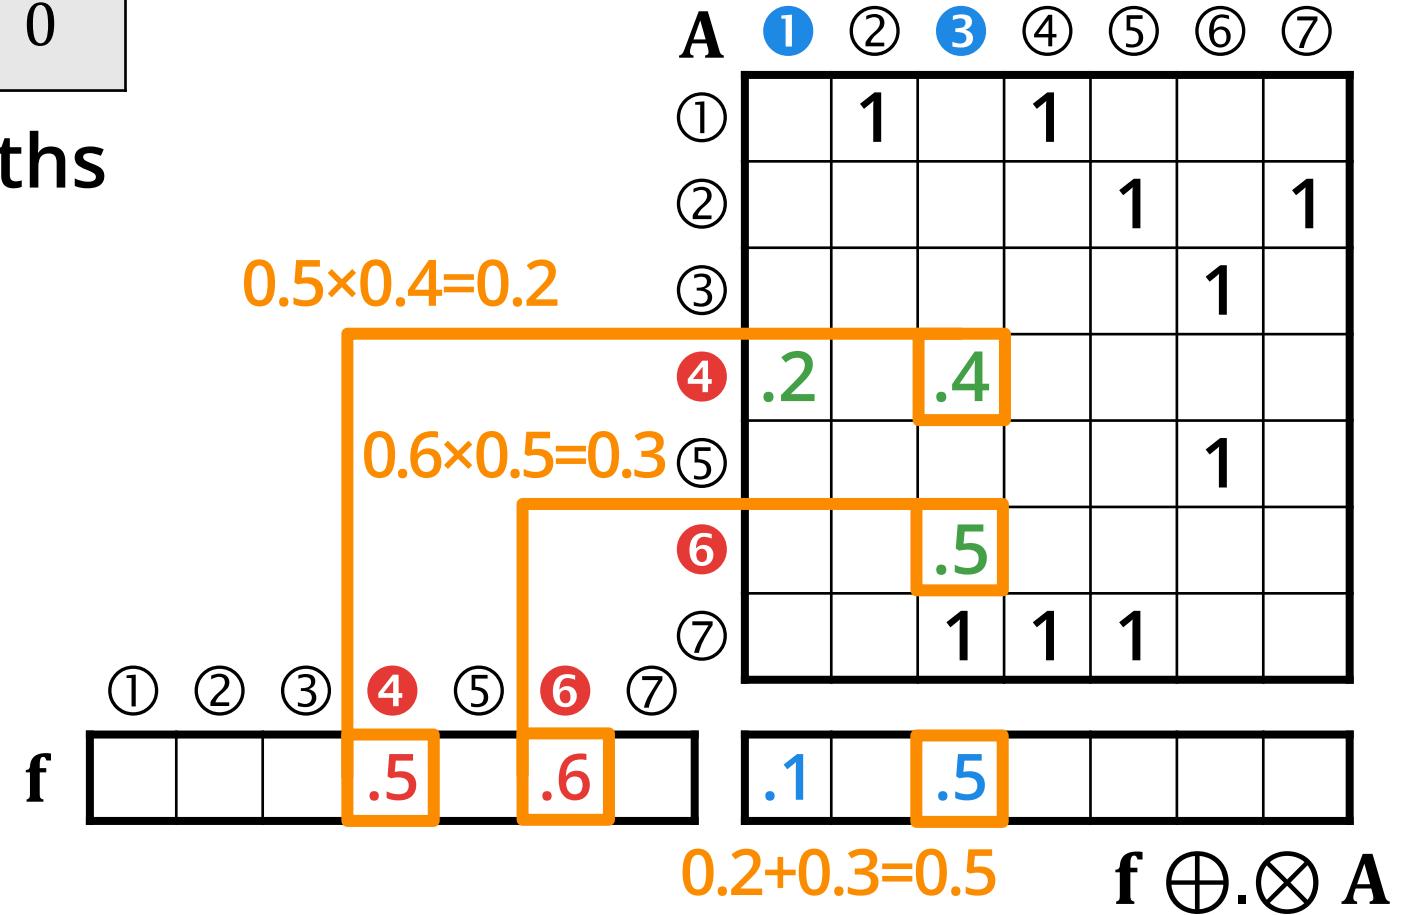
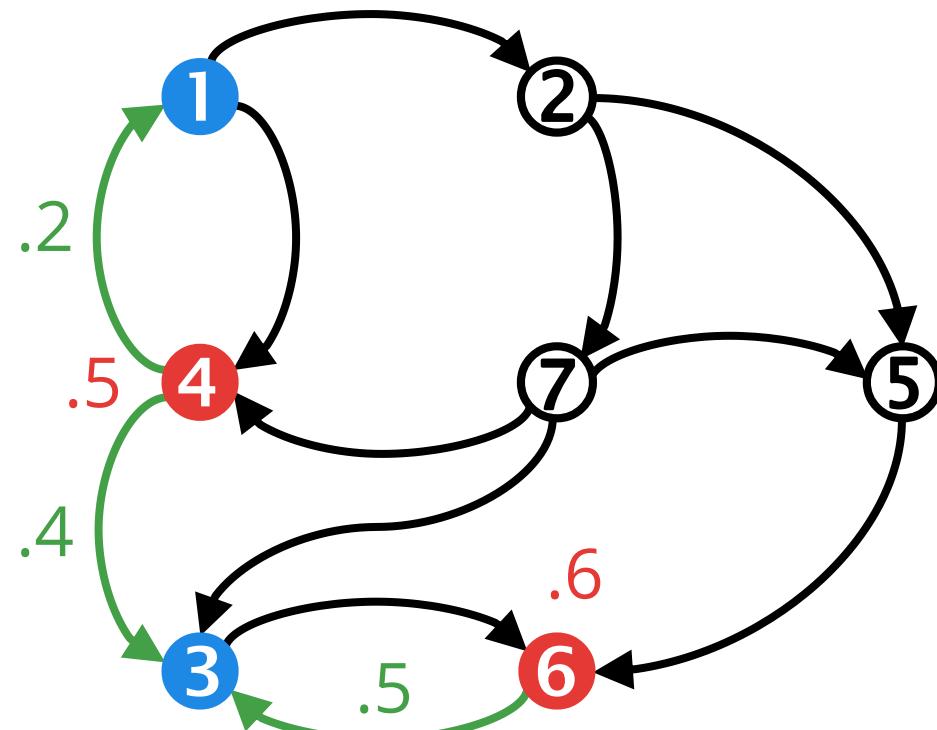
Semantics: reachability



# MATRIX-VECTOR MULTIPLICATION SEMANTICS

semiring	set	$\oplus$	$\otimes$	0
real arithmetic	$\mathbb{R}$	$+$	$\times$	0

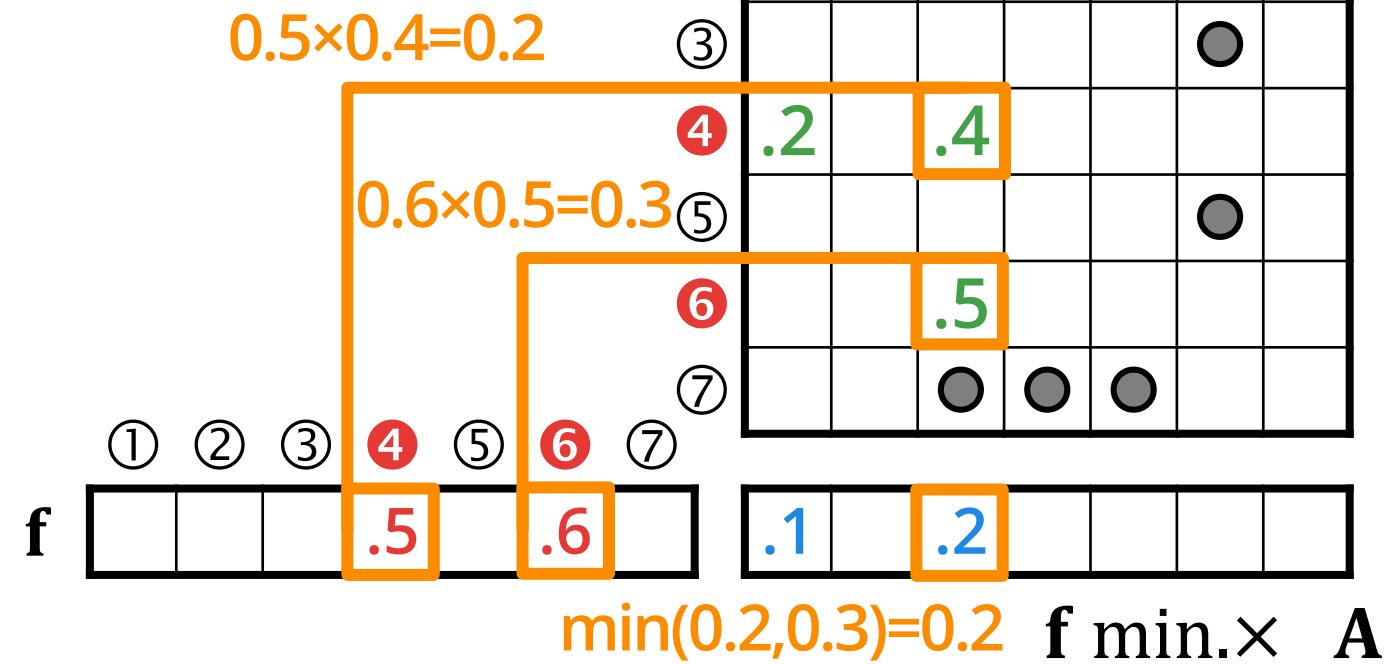
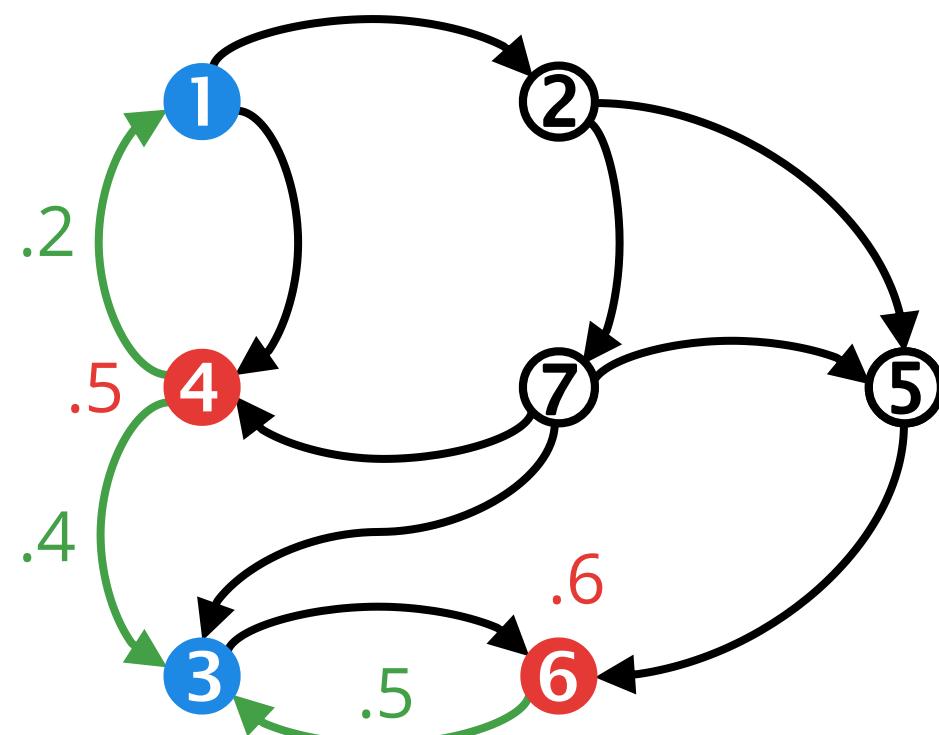
Semantics: strength of all paths



# MATRIX-VECTOR MULTIPLICATION SEMANTICS

semiring	set	$\oplus$	$\otimes$	0
min-times	$\mathbb{R} \cup \{+\infty\}$	min	$\times$	$+\infty$

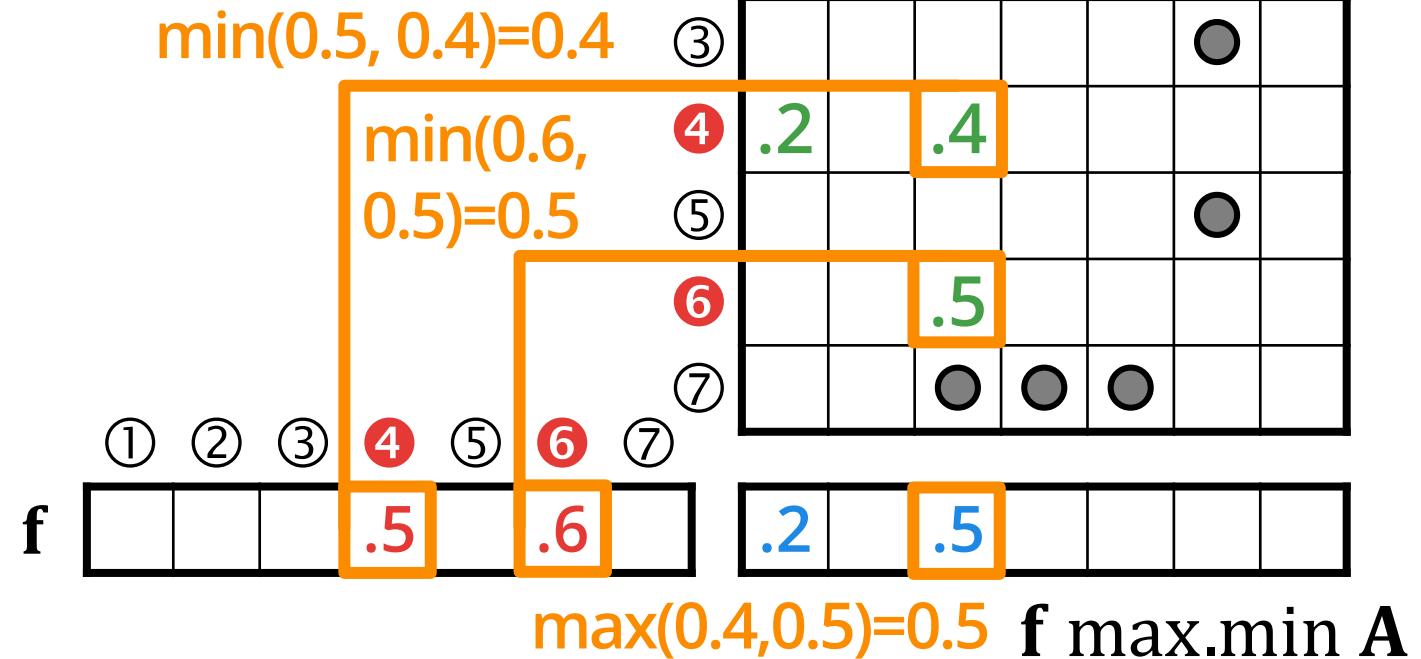
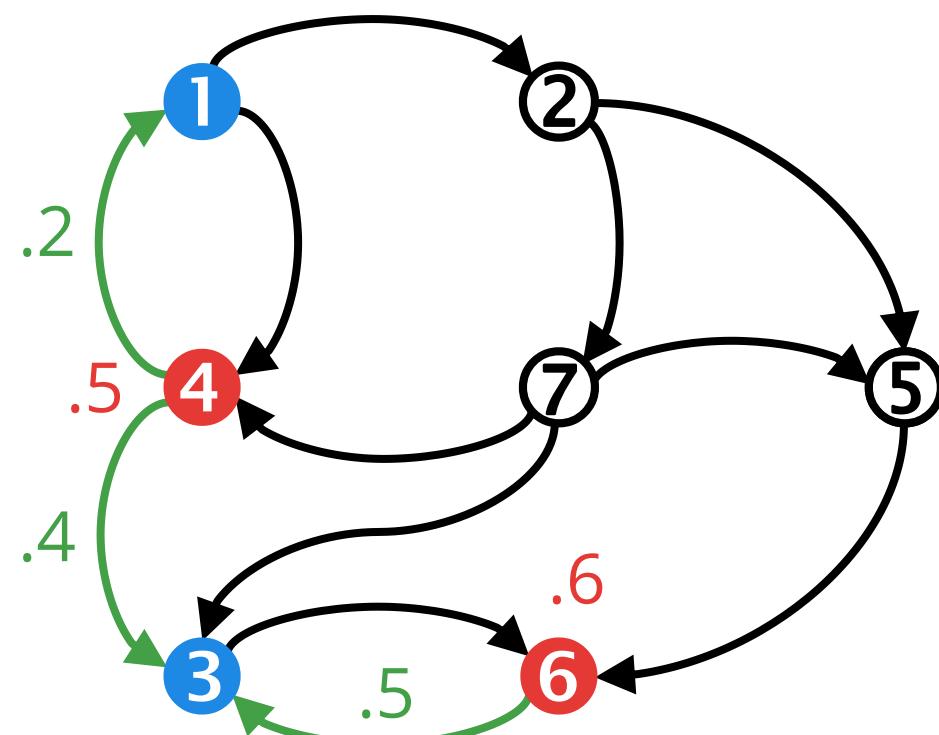
Semantics: shortest product of connections



# MATRIX-VECTOR MULTIPLICATION SEMANTICS

semiring	set	$\oplus$	$\otimes$	0
max-min	$\{0, +\infty\}$	max	min	0

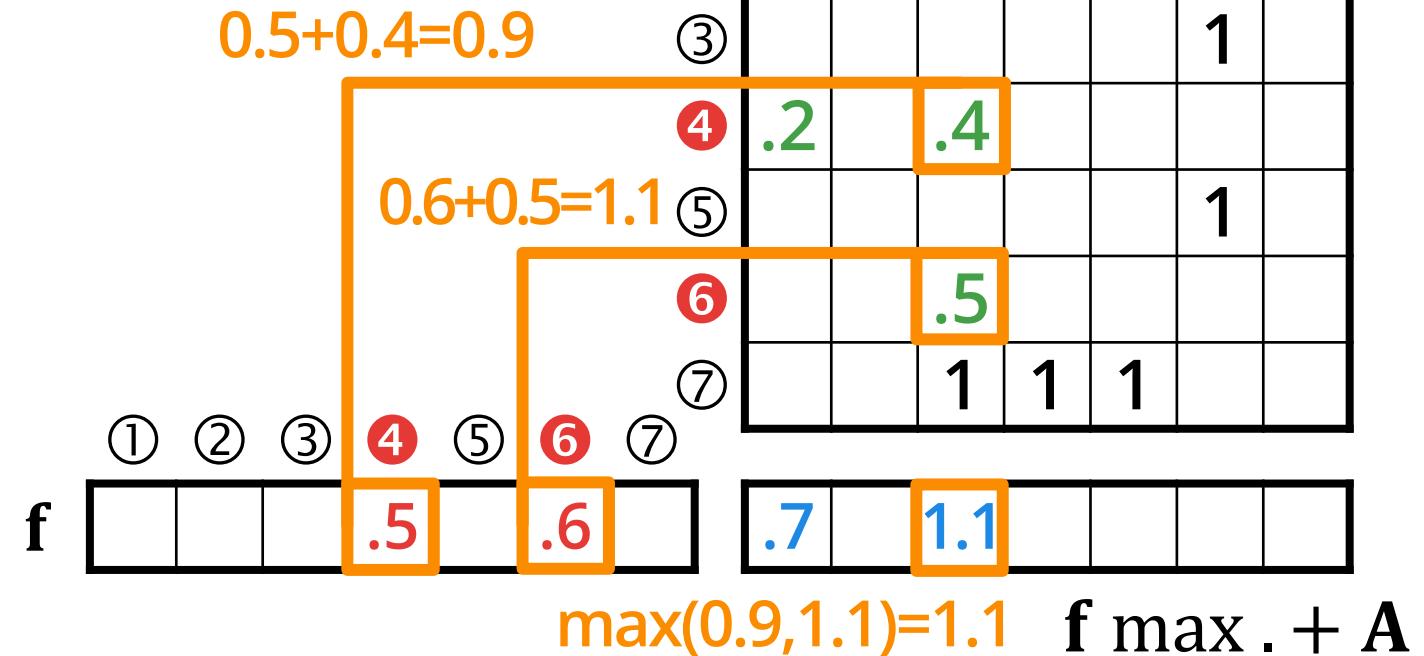
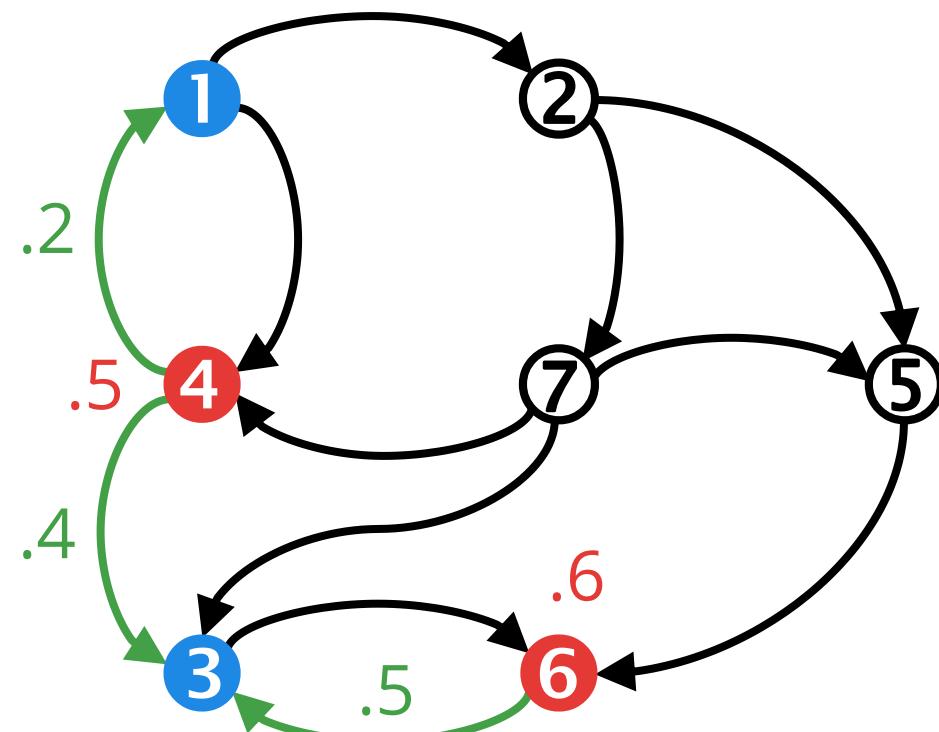
Semantics: longest of all shortest connections



# MATRIX-VECTOR MULTIPLICATION SEMANTICS

semiring	set	$\oplus$	$\otimes$	0
max-plus	$\mathbb{R} \cup \{-\infty\}$	max	+	$-\infty$

Semantics: matching (independent edge set)



# Case study: SIGMOD 2014 Contest

## Overview

# SIGMOD 2014 PROGRAMMING CONTEST

## Annual contest

- Teams compete on database-related programming tasks
- Highly-optimized C++ implementations

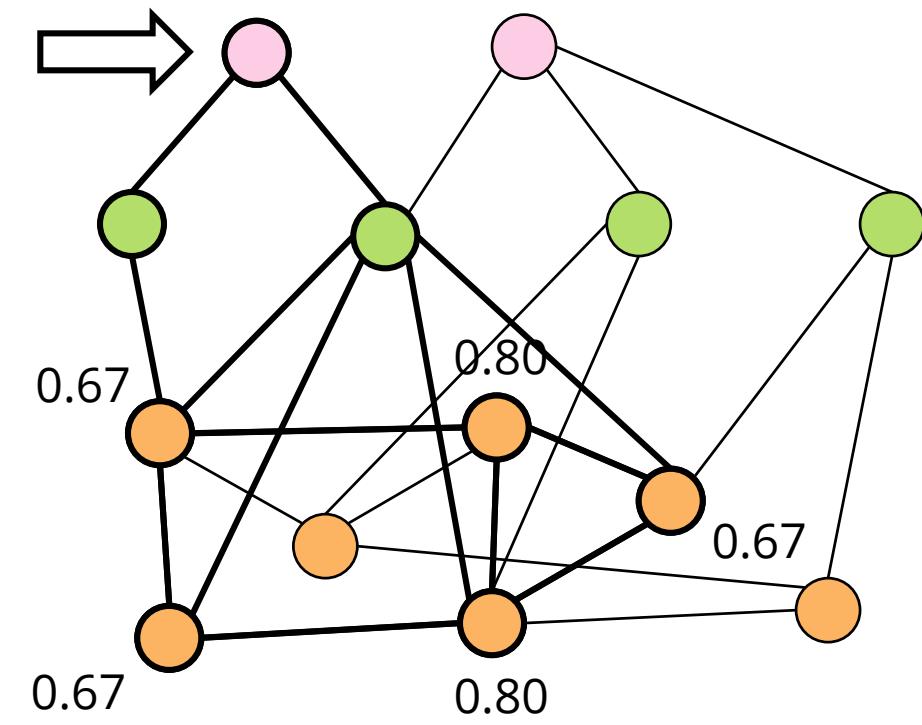
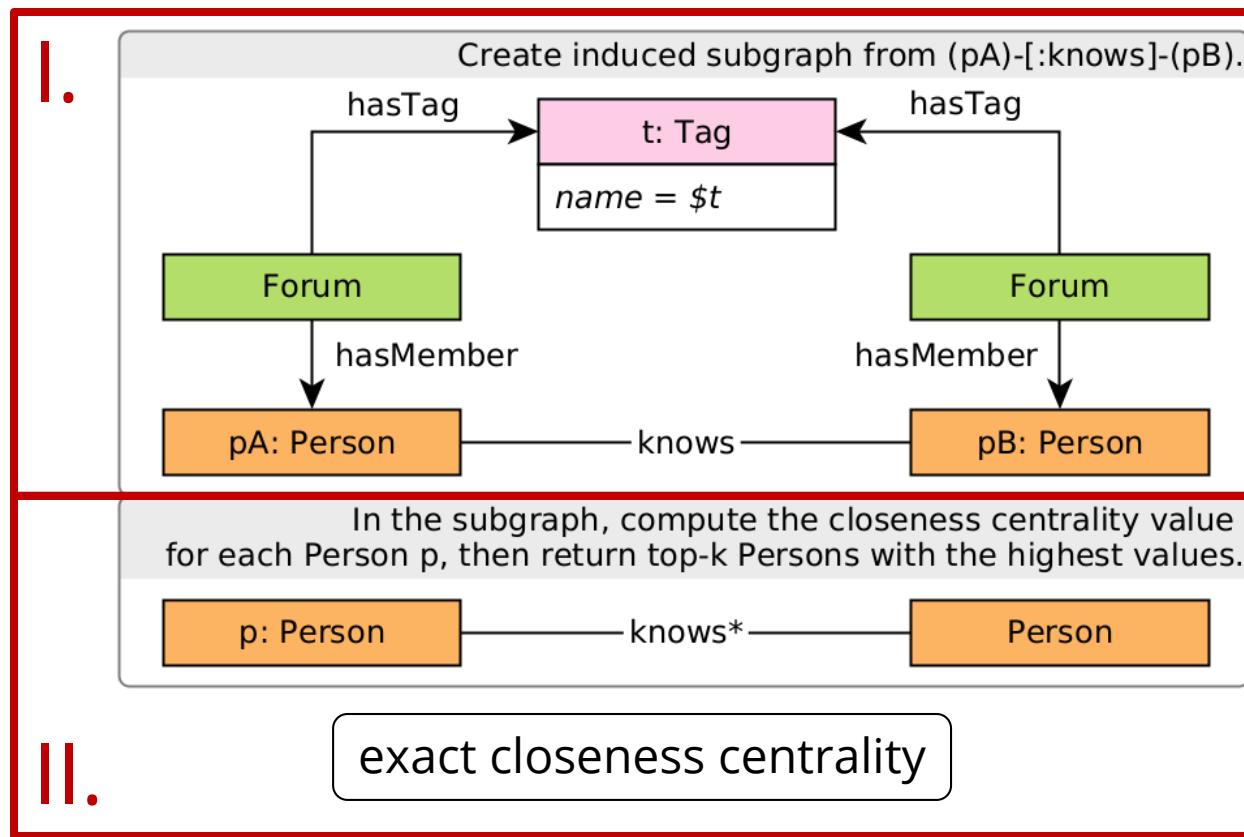
## 2014 event

- Tasks on the LDBC social network graph
  - Benchmark data set for property graphs
  - People, forums, comments, hashtags, etc.
- 4 queries
  - Mix of filtering operations and graph algorithms



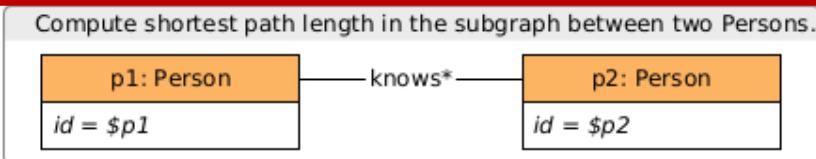
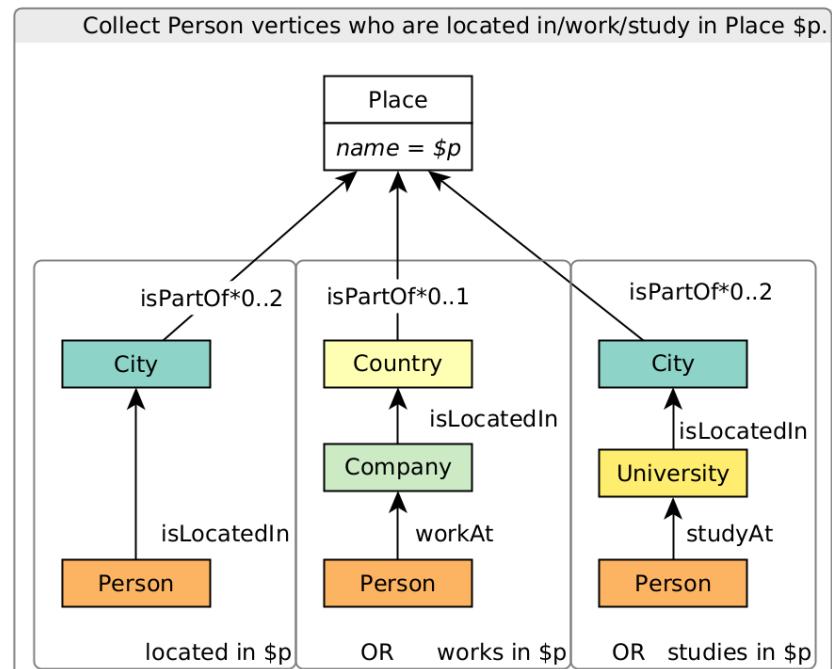
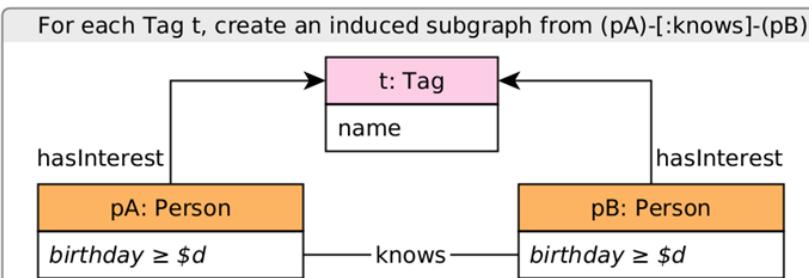
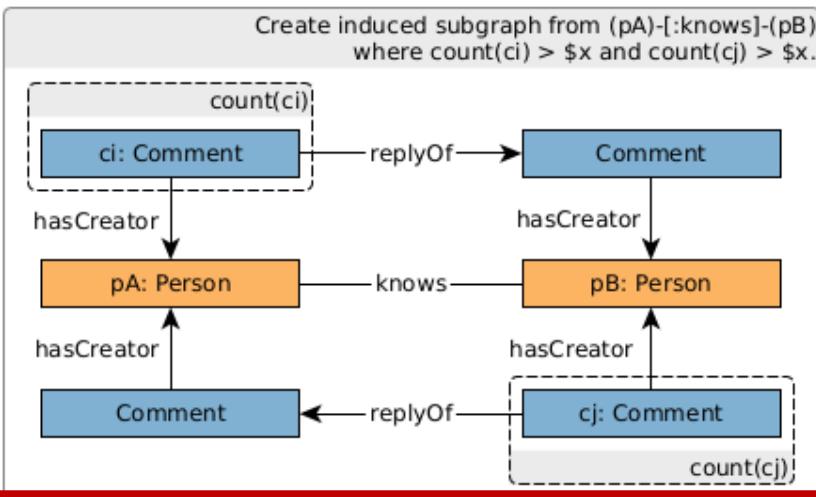
# QUERY TEMPLATE

- I. Compute an induced subgraph over Person-knows-Person
- II. Run the graph algorithm on the subgraph

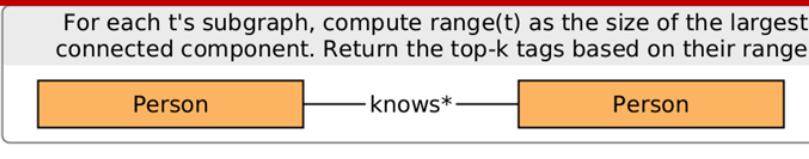


# OVERVIEW OF QUERIES 1, 2, 3

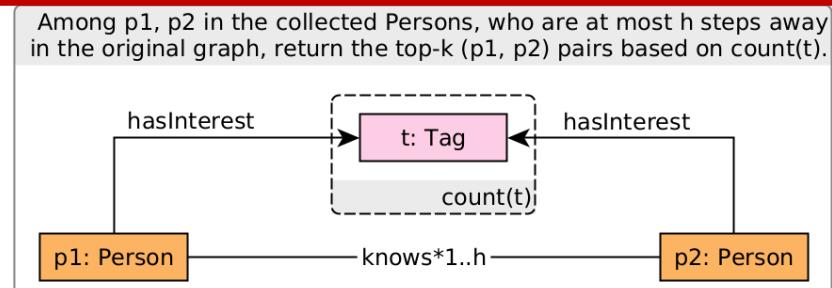
## I. Filter the induced subgraph(s)



unweighted shortest path



connected components



pairwise reachability

## II. Run the graph algorithm

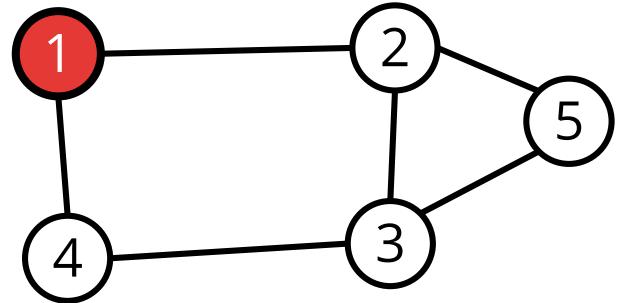
# GRAPHBLAS SOLUTION OF THE QUERIES

- Loading includes relabelling UINT64 vertex IDs to a contiguous sequence  $0 \dots N - 1$ .
- Filtering the induced subgraph from the property graph is mostly straightforward and composable with the algorithms.
- The algorithms can be concisely expressed in GraphBLAS:
  - Connected components ✓ → FastSV [[Zhang et al., PPSC'20](#)]
  - BFS ✓
  - Bidirectional BFS
  - All-source BFS + bitwise optimization
  - Multi-source bidirectional BFS

# Case study: SIGMOD 2014 Contest

BFS

# BFS: BREADTH-FIRST SEARCH



frontier

Boolean  
matrices  
and vectors

1	●
2	
3	
4	
5	

$\langle \neg \text{seen} \rangle$  mask

seen

1	●
2	
3	
4	
5	

A

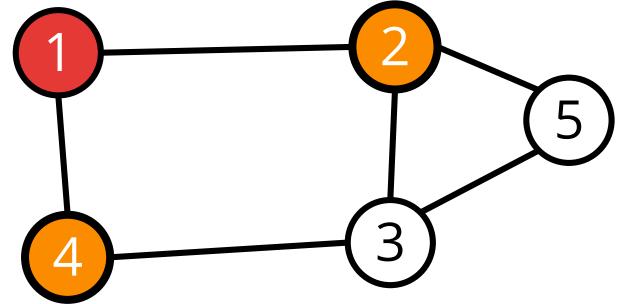
1	2	3	4	5
1		○		○
2	○		○	○
3		○		○
4	○		○	
5		○	○	

X	X	X	X	X

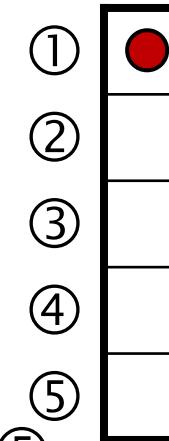
$\oplus$ : any  
 $\otimes$ : pair

$\text{next}(\neg \text{seen}) =$   
A any . pair frontier

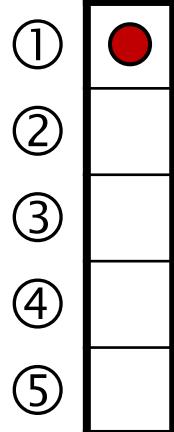
# BFS: BREADTH-FIRST SEARCH



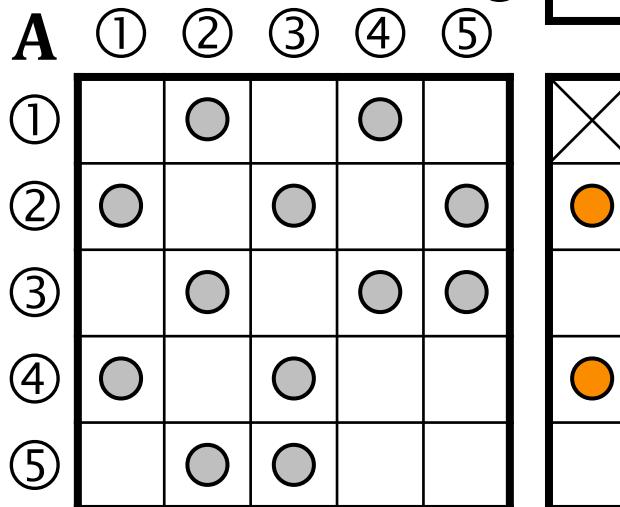
frontier



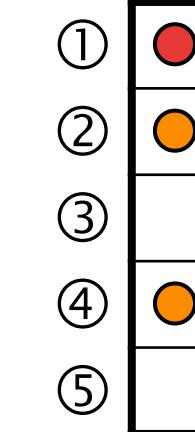
seen



A



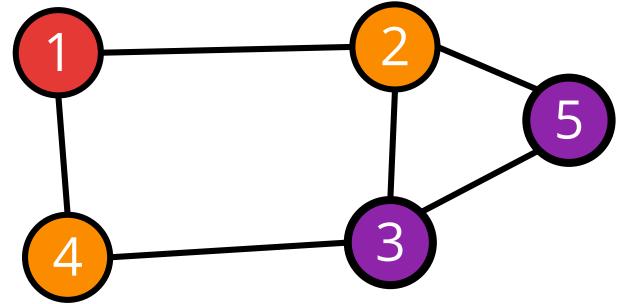
seen'



next( $\neg$ seen) =  
A any . pair frontier

seen' =  
seen any next

# BFS: BREADTH-FIRST SEARCH



**seen**

①	●
②	○
③	
④	○
⑤	

**frontier**

①
②
③
④
⑤

mask prevents  
redundant  
computations

**A**

①		○			○
②	○		○		○
③		○		○	○
④	○		○		
⑤		○	○		

**seen'**

①	●
②	○
③	●
④	○
⑤	●

**next( $\neg$ seen) =**  
**A any . pair frontier**

**seen' =**  
**seen any next**

# Case study: SIGMOD 2014 Contest

## All-source BFS

## Q4: CLOSENESS CENTRALITY VALUES

Q4 computes the top- $k$  Person vertices based on their exact *closeness centrality values*:

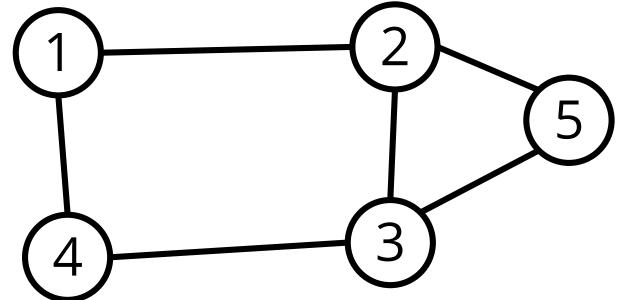
$$CCV(p) = \frac{(C(p) - 1)^2}{(n - 1) \cdot s(p)}$$

where

- $C(p)$  is the size of the connected component of vertex  $p$ ,
- $n$  is the number of vertices in the induced graph,
- $s(p)$  is the sum of geodesic distances to all other reachable persons from  $p$ .

$s(p)$  is challenging: needs unweighted all-pairs shortest paths.

# BOOLEAN ALL-SOURCE BFS ALGORITHM



**Frontier** t<sub>1</sub> t<sub>2</sub> t<sub>3</sub> t<sub>4</sub> t<sub>5</sub>

traversals

①	●			
②		●		
③			●	
④				●
⑤				●

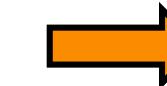
**Seen** t<sub>1</sub> t<sub>2</sub> t<sub>3</sub> t<sub>4</sub> t<sub>5</sub>

①	●			
②		●		
③			●	
④				●
⑤				●

**A** ① ② ③ ④ ⑤

①		●		●
②	●		●	●
③		●		●
④	●		●	
⑤		●	●	

		●	●	
	●		●	●
		●		●
	●		●	
		●	●	



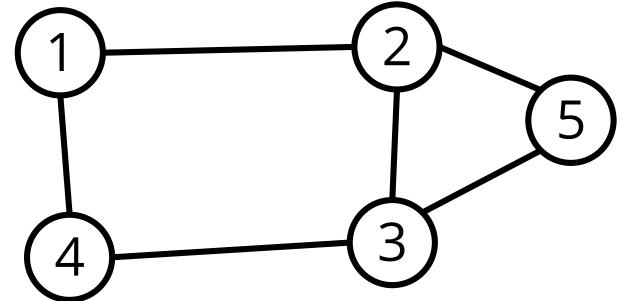
**Seen'** t<sub>1</sub> t<sub>2</sub> t<sub>3</sub> t<sub>4</sub> t<sub>5</sub>

①	●	●	●	
②	●	●	●	●
③		●	●	●
④	●		●	●
⑤		●	●	●

**Next**(¬**Seen**) =  
A any . pair Frontier

**Seen'** =  
Seen any Next

# BOOLEAN ALL-SOURCE BFS ALGORITHM



Frontier t<sub>1</sub> t<sub>2</sub> t<sub>3</sub> t<sub>4</sub> t<sub>5</sub>

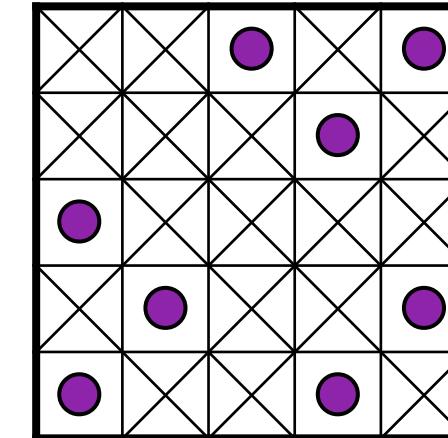
1			●	●	
2	●			●	●
3		●		●	●
4	●		●		
5		●	●		

Seen t<sub>1</sub> t<sub>2</sub> t<sub>3</sub> t<sub>4</sub> t<sub>5</sub>

1	●		●		
2		●	●		●
3		●	●	●	●
4	●		●	●	
5		●	●		●

A ① ② ③ ④ ⑤

1		●		●	
2	●		●		●
3		●	●	●	●
4	●		●	●	
5		●	●		



Seen' t<sub>1</sub> t<sub>2</sub> t<sub>3</sub> t<sub>4</sub> t<sub>5</sub>

1	●		●	●	●
2		●	●	●	●
3	●	●	●	●	●
4	●	●	●	●	●
5	●	●	●	●	●

Next( $\neg$ Seen) =  
A any . pair Frontier

Seen' =  
Seen any Next

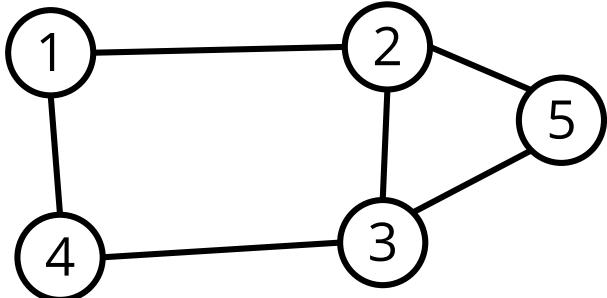
# Case study: SIGMOD 2014 Contest

## Bitwise all-source BFS

# BITWISE ALL-SOURCE BFS ALGORITHM

- For large graphs, the all-source BFS algorithm might need to run 500k+ traversals
- Two top-ranking teams used bitwise operations to process traversals in batches of 64 [[Then et al., VLDB'15](#)]
- This idea can be adopted in the GraphBLAS algorithm by
  - using UINT64 values
  - performing the multiplication on the `bor . second` semiring, where `bor` is “bitwise or” and  $\text{second}(x, y) = y$
- 5-10x speedup compared to the Boolean all-source BFS

# BITWISE ALL-SOURCE BFS ALGORITHM



**Frontier** t<sub>1-t<sub>4</sub></sub> t<sub>5</sub>

Using UINT4s here

①	1000	0000
②	0100	0000
③	0010	0000
④	0001	0000
⑤	0000	1000

**Seen** t<sub>1-t<sub>4</sub></sub> t<sub>5</sub>

①	1000	0000
②	0100	0000
③	0010	0000
④	0001	0000
⑤	0000	1000

**A** ① ② ③ ④ ⑤

①		●		●	
②	●		●		●
③		●		●	●
④	●		●		
⑤		●	●		

0101	0000
1010	1000
0101	1000
1010	0000
0110	0000

**Next =**

A bor . second Frontier



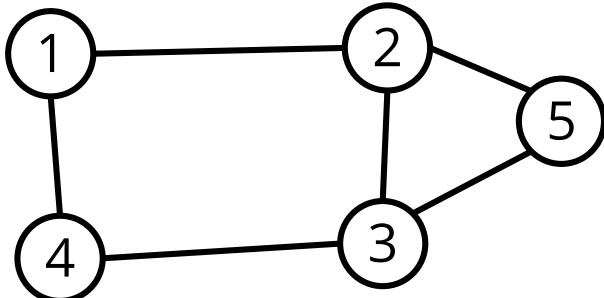
**Seen'** t<sub>1-t<sub>4</sub></sub> t<sub>5</sub>

①	0101	0000
②	1010	1000
③	0101	1000
④	1010	0000
⑤	0110	0000

**Seen' =**

Seen bor Next

# BITWISE ALL-SOURCE BFS ALGORITHM



Frontier	t1-t4	t5
①	0101	0000
②	1010	1000
③	0101	1000
④	1010	0000
⑤	0110	0000

Full VLDB paper  
on this algorithm

vs.

9 GrB operations

Seen	t1-t4	t5
①	0101	0000
②	1010	1000
③	0101	1000
④	1010	0000
⑤	0110	0000

A	①	②	③	④	⑤
①		●		●	
②	●		●		●
③		●		●	●
④	●		●		
⑤		●	●		

	t1-t4	t5
0010	1000	
0001	0000	
1000	0000	
0100	1000	
1001	0000	

Next =

A bor . second Frontier



Seen'	t1-t4	t5
①	1111	1000
②	1111	1000
③	1111	1000
④	1111	1000
⑤	1111	1000

Seen' =

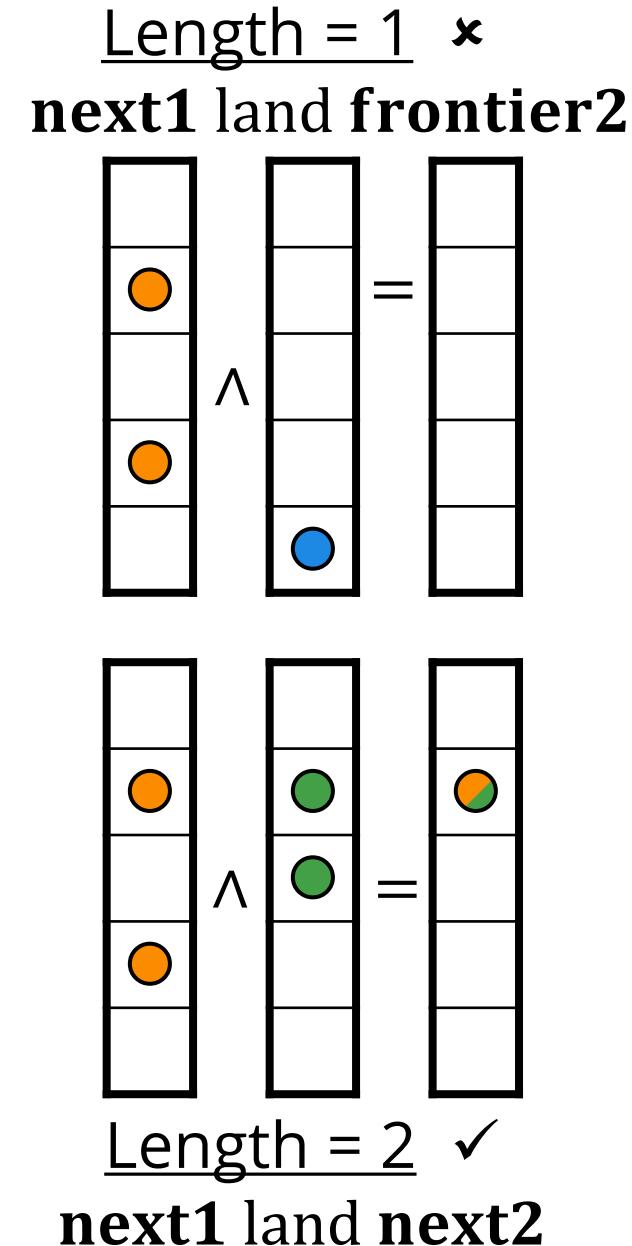
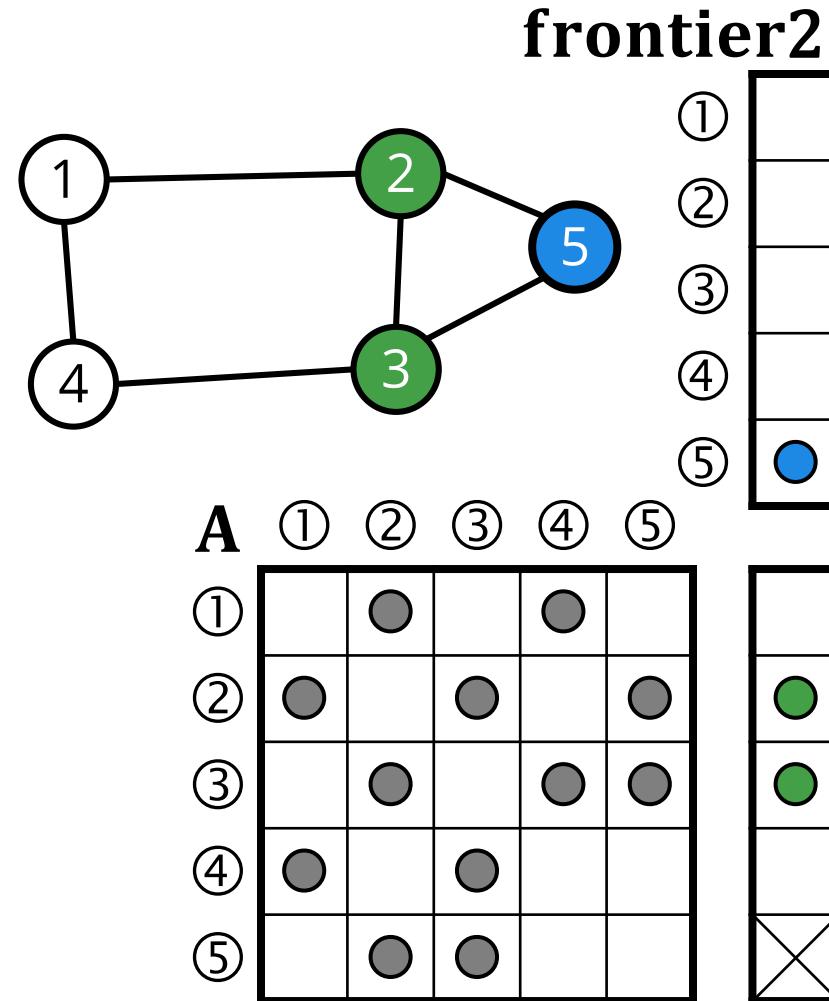
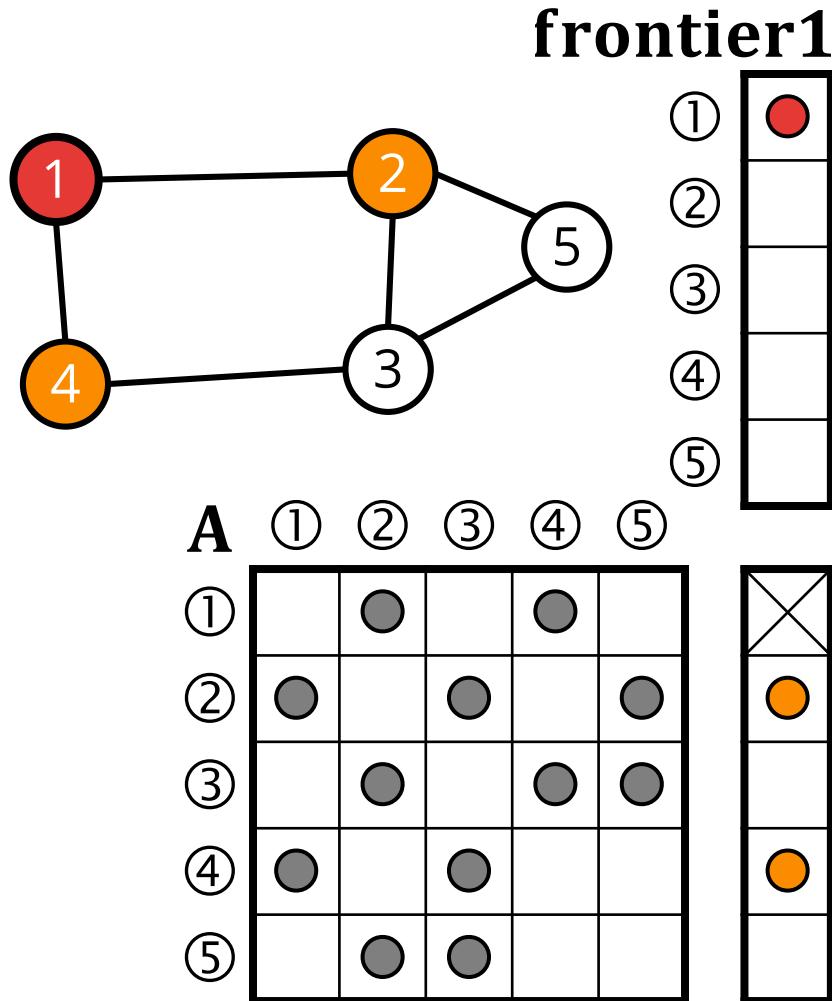
Seen bor Next

# Case study: SIGMOD 2014 Contest

## Bidirectional BFS

# BIDIRECTIONAL BFS

Advance frontiers alternately  
and intersect them

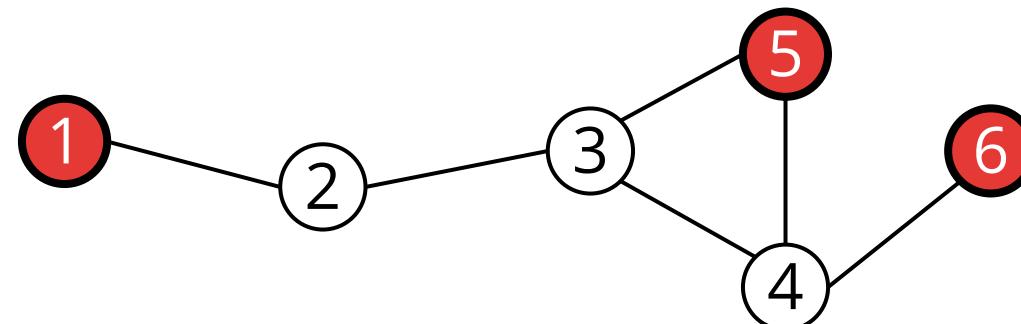


# Case study: SIGMOD 2014 Contest

## Bidirectional MSBFS

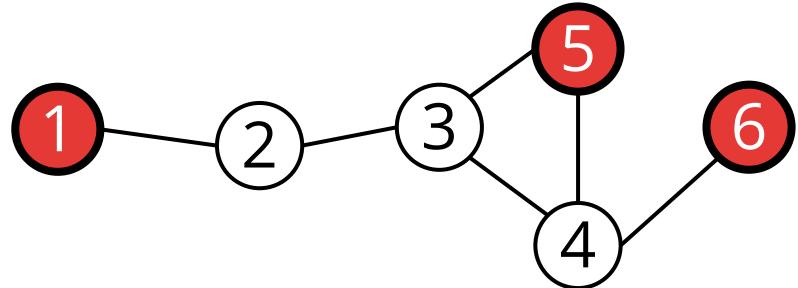
# BIDIRECTIONAL MSBFS ALGORITHM

- **Pairwise reachability problem:**  
From a given set of  $k$  vertices, which pairs of vertices are reachable from each other with at most  $h$  hops?
- **Naïve solution:**  
Run a  $k$ -source MSBFS for  $h$  steps and check reachability.  
The frontiers get large as they grow exponentially.
- **Better solution:**  
Advance all frontiers simultaneously for  $\lceil h/2 \rceil$  iterations.



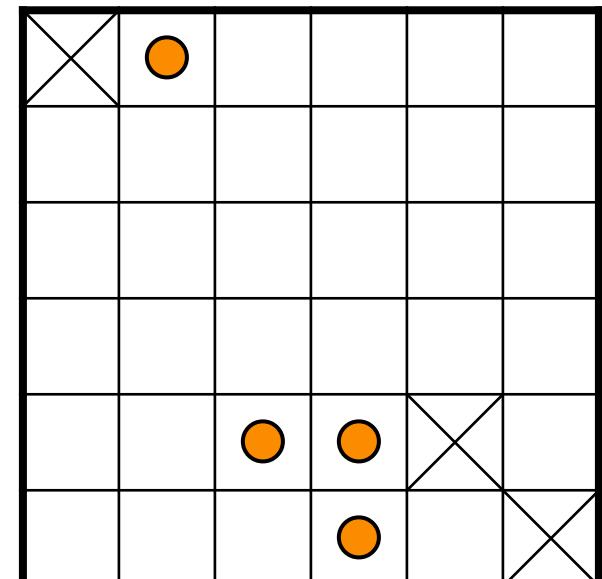
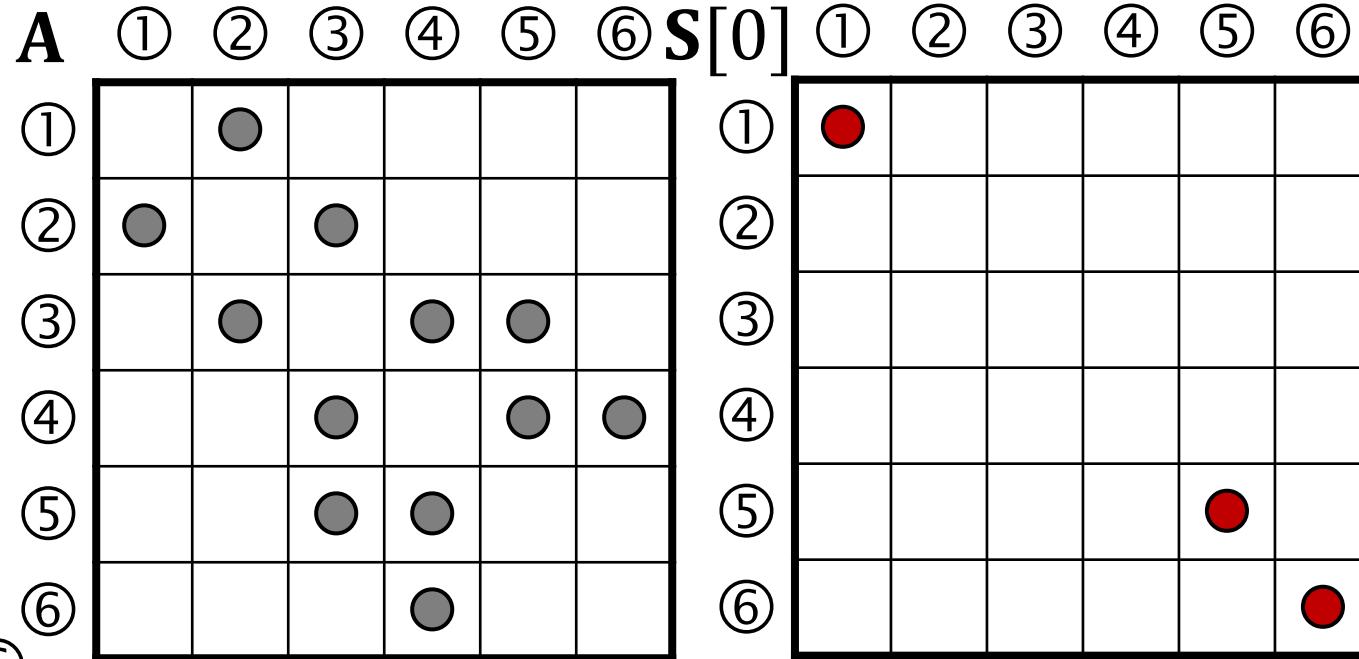
# BIDIRECTIONAL MSBFS

Seen[1]: reachability with  $\leq 1$  hops

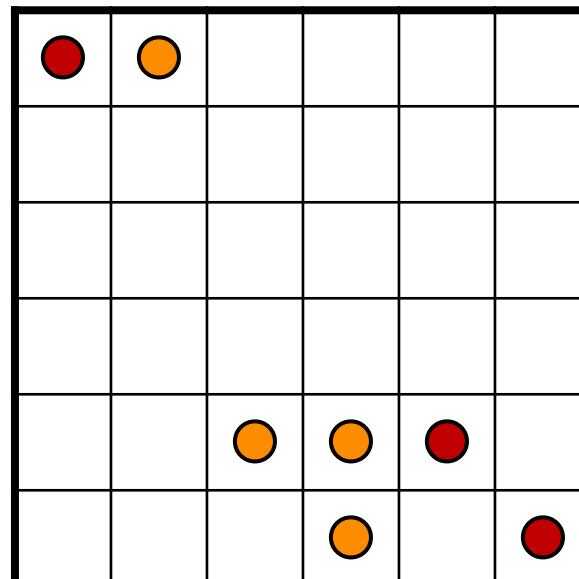


F ① ② ③ ④ ⑤ ⑥

1	●					
2						
3						
4						
5						
6						



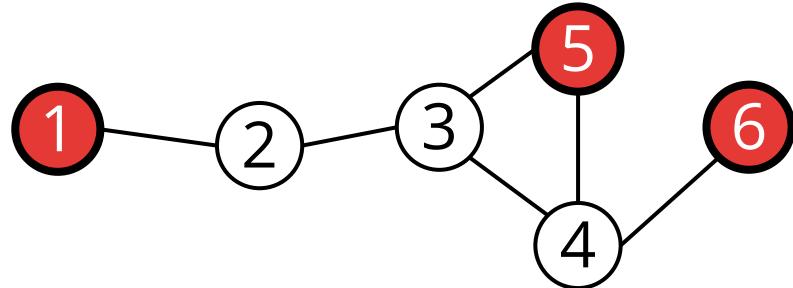
Next[1]



Seen[1]

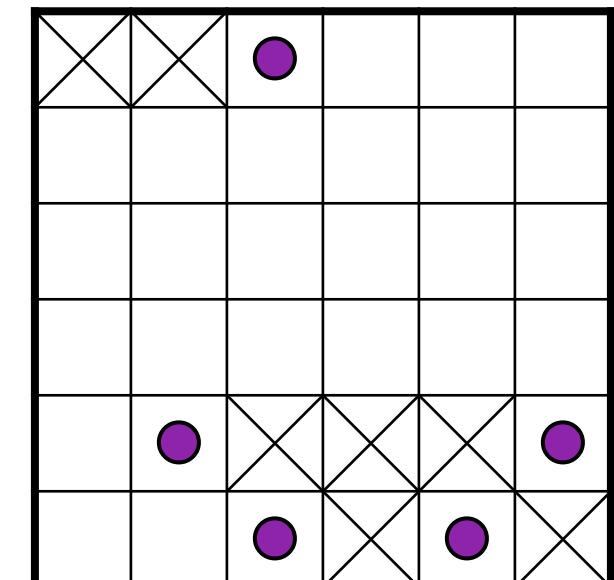
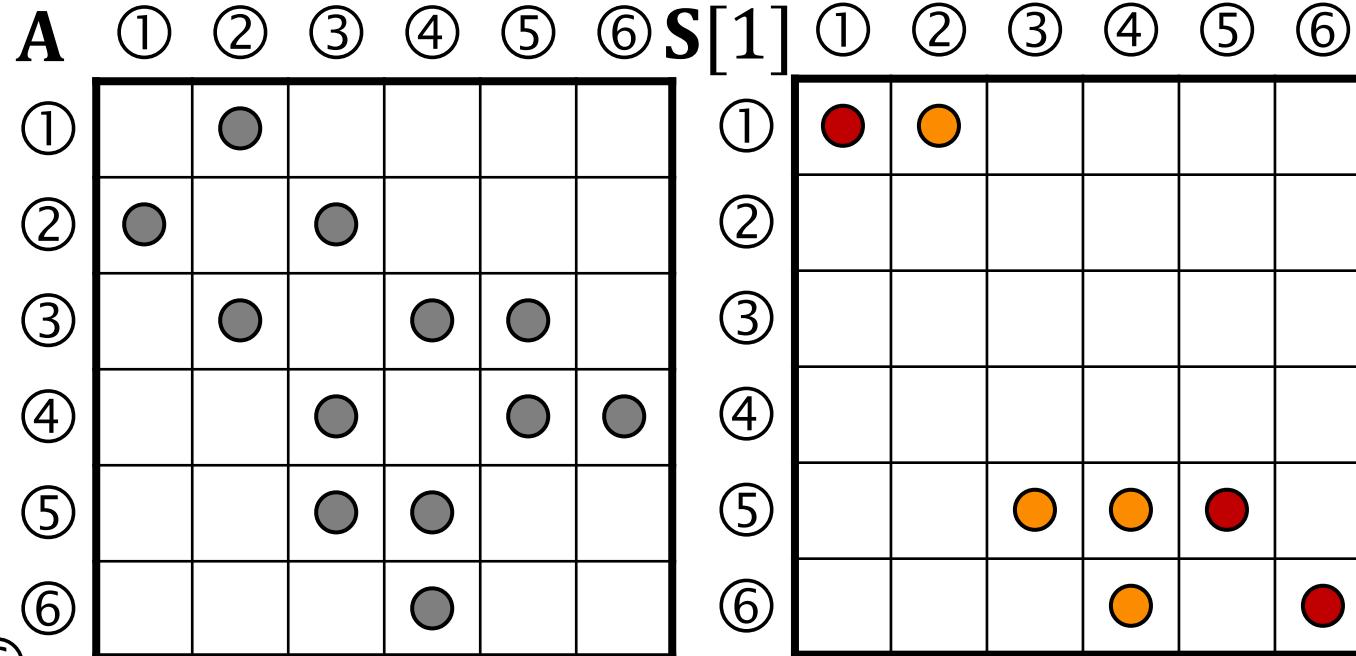
# BIDIRECTIONAL MSBFS

Seen[2]: reachability with  $\leq 2$  hops

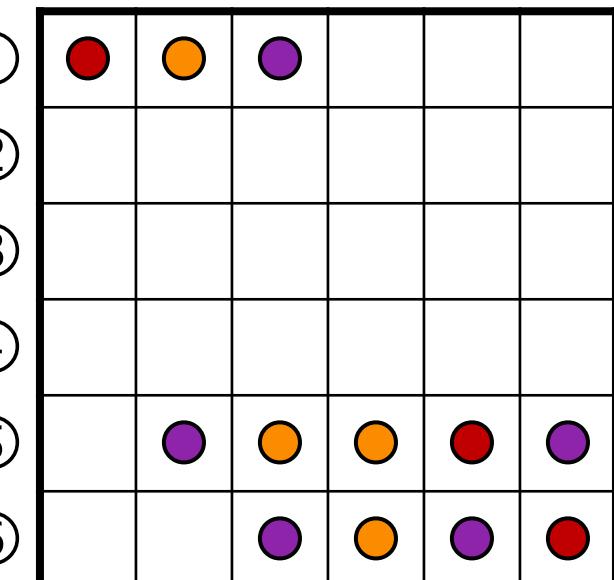


F ① ② ③ ④ ⑤ ⑥

1		●				
2						
3						
4						
5		●	●			
6			●			

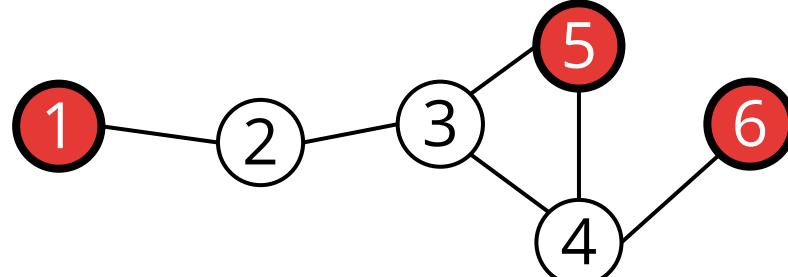


Next[2]<Seen[1]>



Seen[2]

# BIDIRECTIONAL MSBFS: PATHS OF LENGTH $\leq 4$



**Seen**[2]<sup>T</sup>

	①	②	③	④	⑤	⑥
①	●					
②		●				
③			●		●	●
④				●	●	●
⑤					●	●
⑥					●	●

From vertex 1, we could get to these vertices with  $\leq 2$  hops

**Seen**[2] ① ② ③ ④ ⑤ ⑥

	①	②	③	④	⑤	⑥
①	●	●	●			
②						
③						
④						
⑤		●	●	●	●	●
⑥			●	●	●	●

From vertex 5, we could get to these vertices with  $\leq 2$  hops

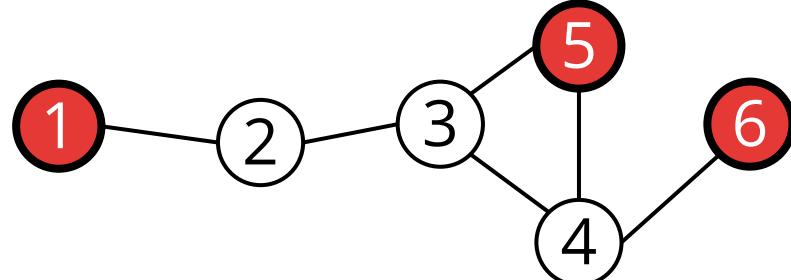
To get paths of at most 4 hops, we compute

**Seen**[2] any.pair **Seen**[2]<sup>T</sup>

Here, we found paths between all pairs:

- from ① to ⑤,
- from ① to ⑥,
- from ⑤ to ⑥.

# BIDIRECTIONAL MSBFS: PATHS OF LENGTH = 3



$\text{Next}[2]^\top$

①	②	③	④	⑤	⑥
①					
②					●
③	●				●
④					
⑤					●
⑥				●	

From vertex 1, we could get to this vertex with 2 hops

$\text{Next}[1]$  ① ② ③ ④ ⑤ ⑥

①		●				
②						
③						
④						
⑤						
⑥						

From vertex 5, we could get to these vertices with 1 hop

X	X	X	X	X	X
X	X	X	X	X	X
X	X	X	X	X	X
X	X	X	X	X	X
X	X	X	X	X	X

●				●	

To get exactly 3-length paths we compute

$\text{Next}[1] \text{ any.pair } \text{Next}[2]^\top$

We found two 3-length paths:

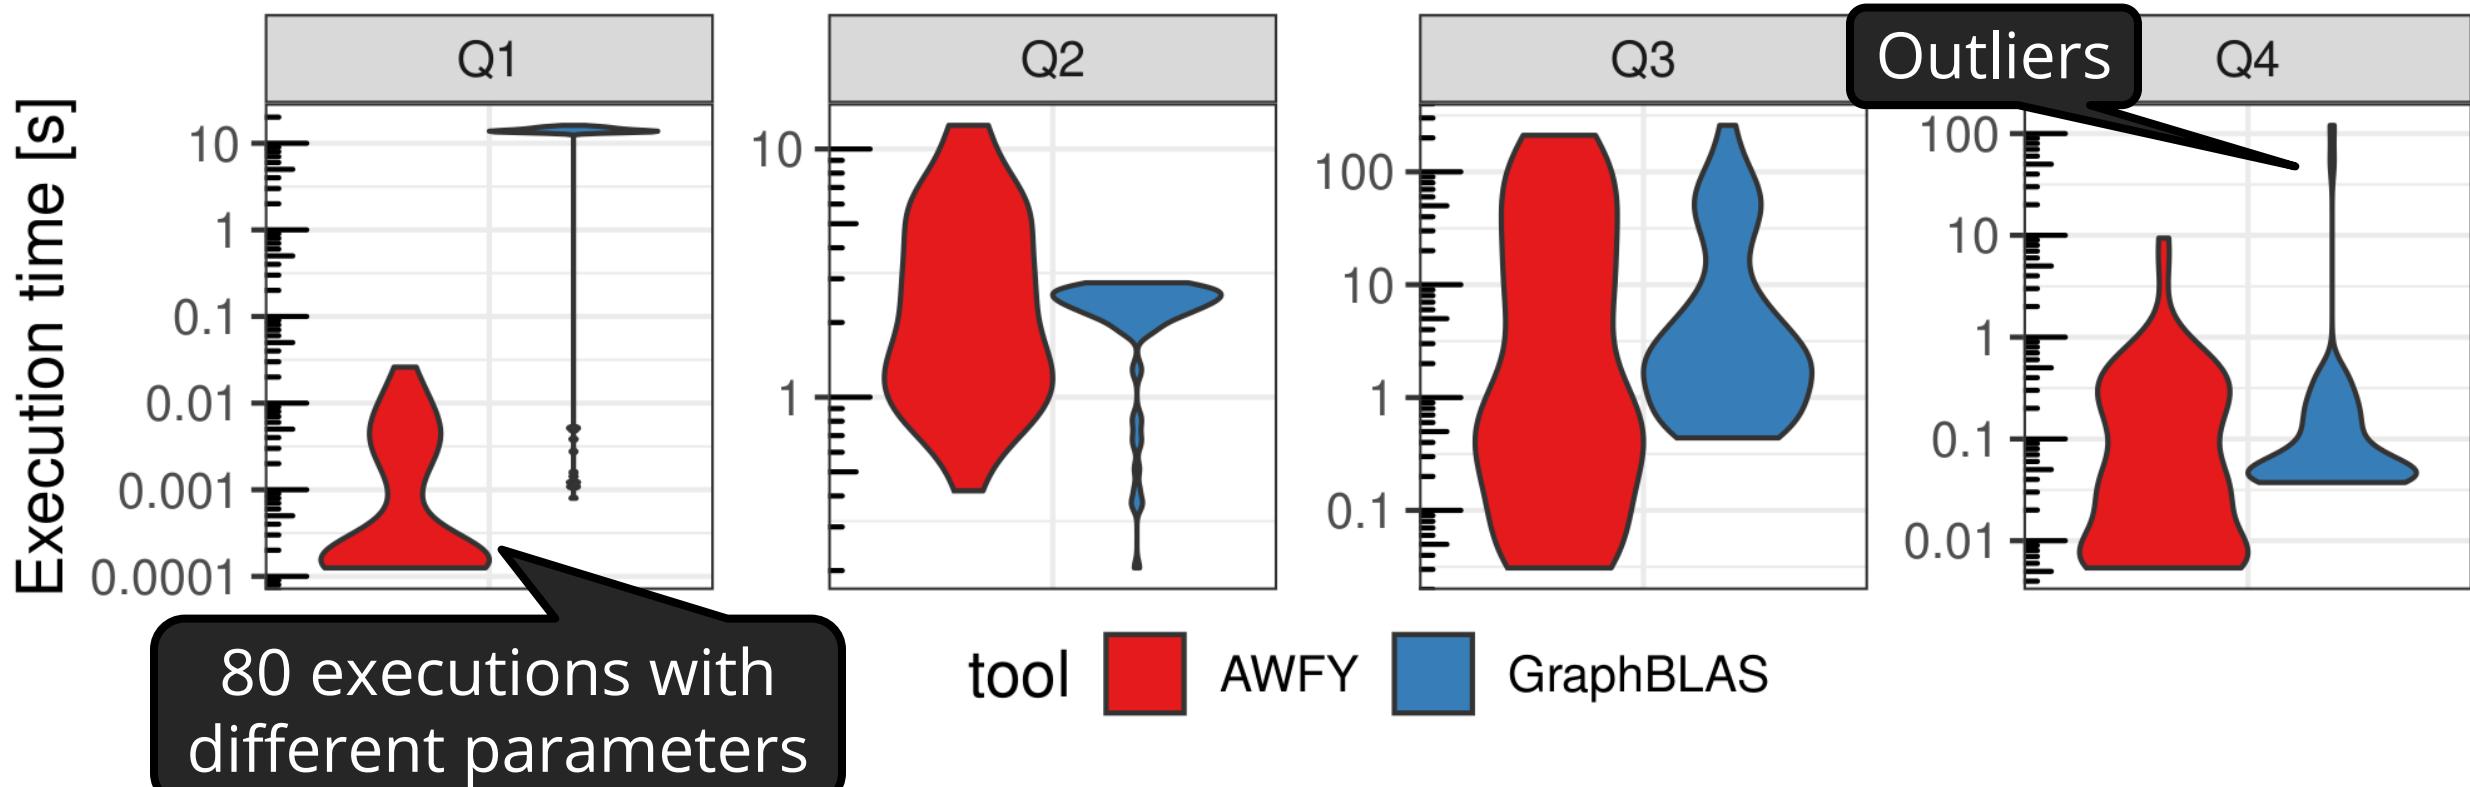
- from ① to ⑤
- from ⑤ to ⑥.

# Case study: SIGMOD 2014 Contest

## Results

# BENCHMARK RESULTS

- The top solution of AWFY vs. SuiteSparse:GraphBLAS v3.3.3
- AWFY's solution uses SIMD instructions → difficult to port
- GraphBLAS load times are slow (see details in paper)



# SUMMARY

- An interesting case study, see [technical report](#)
- GraphBLAS can capture mixed workloads
  - Induced subgraph computations are simple to express
  - Algorithms are concise, bitwise optimizations can be adopted
  - Performance is *sometimes* on par with specialized solutions
- Future optimizations
  - Q1: filter the induced subgraph on-the-fly
  - Q4: use more sophisticated unweighted shortest path algorithms



[sigmod2014-contest-graphblas](#)



M. Elekes et al., *A GraphBLAS solution to the SIGMOD 2014 Programming Contest using multi-source BFS*, HPEC 2020

# ACKNOWLEDGEMENTS

- Thanks for feedback on these slides to Tim Mattson, Tim Davis, Jeremy Kepner, John Gilbert, Aydin Buluc, Tamás Nyíri, Gyula Katona, Dávid Sándor, Attila Nagy, János Benjamin Antal, Márton Elekes, Florentin Dörre.
- Thanks to the LDBC Graphalytics task force for creating the benchmark and assisting in the measurements.
- Thanks to the Master's students at BME for exploring the area of GraphBLAS: Bálint Hegyi, Petra Várhegyi, Lehel Boér

# Notes

# ABOUT THIS PRESENTATION

- This presentation is intended to serve as an introduction to semiring-based graph processing and the GraphBLAS.
- Common graph algorithms (BFS, shortest path, PageRank, etc.) are used to demonstrate the features of GraphBLAS. Many of the algorithms presented are part of the LAGraph library.
- The presentation complements existing technical talks on GraphBLAS and can form a basis of anything from a short 20min overview to 2×90min lectures on the GraphBLAS.
- The slides contain numerous references to papers and talks.
- There are detailed textual explanations on some slides to make them usable as a learning material.

# TECHNICAL DETAILS

- The slides were created with PowerPoint 2016 using the Open Sans and DejaVu Sans Mono font families (embedded in the presentation).
- The mathematical expressions are typeset with PowerPoint's built-in Equation Editor.
- The circled numbers (denoting graph vertices) are rendered using the standard Wingdings font.
- The text is written in Oxford English.
- The icons for referenced papers and talks are clickable and will lead to an open-access / preprint / author's copy version of the referred work (if such copy exists). The icons depict the first page of the cited document.